

# knn

July 25, 2022

```
[40]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'cs231n/assignments/assignment1/'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.  
/content/drive/My Drive/cs231n/assignments/assignment1/cs231n/datasets  
/content/drive/My Drive/cs231n/assignments/assignment1

## 1 k-Nearest Neighbor (kNN) exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

The kNN classifier consists of two stages:

- During training, the classifier takes the training data and simply remembers it
- During testing, kNN classifies every test image by comparing to all training images and transferring the labels of the k most similar training examples

- The value of  $k$  is cross-validated

In this exercise you will implement these steps and understand the basic Image Classification pipeline, cross-validation, and gain proficiency in writing efficient, vectorized code.

```
[41]: # Run some setup code for this notebook.

import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

# This is a bit of magic to make matplotlib figures appear inline in the
↪ notebook
# rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/questions/1907993/
↪ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

```
[42]: # Load the raw CIFAR-10 data.
cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

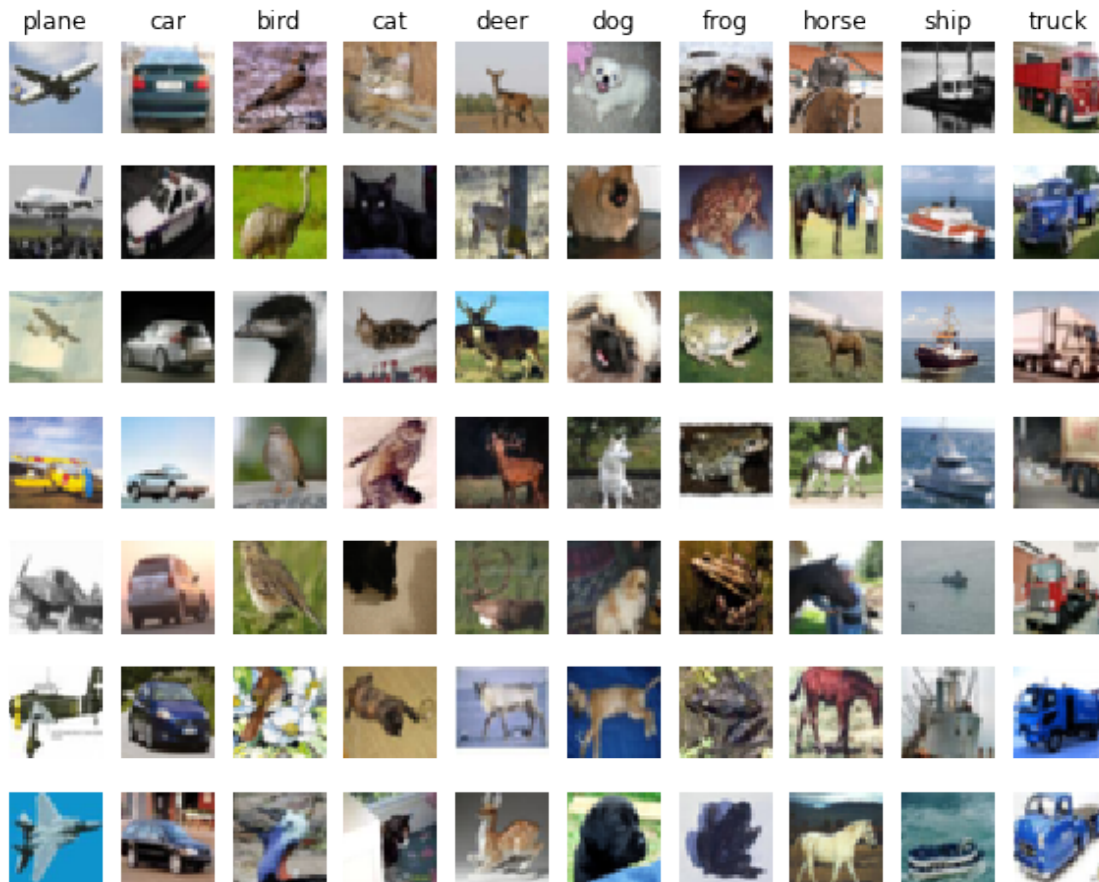
# Cleaning up variables to prevent loading data multiple times (which may cause
↪ memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# As a sanity check, we print out the size of the training and test data.
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Clear previously loaded data.
Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
```

```
[43]: # Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', '
↳ 'ship', 'truck']
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()
```



```
[44]: # Subsample the data for more efficient code execution in this exercise
num_training = 5000
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]

num_test = 500
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]

print(X_train.shape, X_test.shape)
print(X_train.shape[0])

# Reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
print(X_train.shape, X_test.shape)
```

```
(5000, 32, 32, 3) (500, 32, 32, 3)
5000
(5000, 3072) (500, 3072)
```

```
[45]: from cs231n.classifiers import KNearestNeighbor

# Create a kNN classifier instance.
# Remember that training a kNN classifier is a noop:
# the Classifier simply remembers the data and does no further processing
classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
```

We would now like to classify the test data with the kNN classifier. Recall that we can break down this process into two steps:

1. First we must compute the distances between all test examples and all train examples.
2. Given these distances, for each test example we find the k nearest examples and have them vote for the label

Lets begin with computing the distance matrix between all training and test examples. For example, if there are  $N_{tr}$  training examples and  $N_{te}$  test examples, this stage should result in a  $N_{te} \times N_{tr}$  matrix where each element  $(i,j)$  is the distance between the i-th test and j-th train example.

**Note:** For the three distance computations that we require you to implement in this notebook, you may not use the `np.linalg.norm()` function that numpy provides.


First, open `cs231n/classifiers/k_nearest_neighbor.py` and implement the function `compute_distances_two_loops` that uses a (very inefficient) double loop over all pairs of (test, train) examples and computes the distance matrix one element at a time.

```
[46]: # Open cs231n/classifiers/k_nearest_neighbor.py and implement
# compute_distances_two_loops.

# Test your implementation:
dists = classifier.compute_distances_two_loops(X_test)
print(dists.shape)
```

```
(500, 5000)
```

```
[47]: # We can visualize the distance matrix: each row is a single test example and
# its distances to training examples
plt.imshow(dists, interpolation='none')
plt.show()
```



### Inline Question 1

Notice the structured patterns in the distance matrix, where some rows or columns are visibly brighter. (Note that with the default color scheme black indicates low distances while white indicates high distances.)

- What in the data is the cause behind the distinctly bright rows?
- What causes the columns?

*Your Answer : fill this in.*

```
[48]: # Now implement the function predict_labels and run the code below:
# We use k = 1 (which is Nearest Neighbor).
y_test_pred = classifier.predict_labels(dists, k=1)

# Compute and print the fraction of correctly predicted examples
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
```

Got 137 / 500 correct => accuracy: 0.274000

You should expect to see approximately 27% accuracy. Now lets try out a larger k, say k = 5:

```
[49]: y_test_pred = classifier.predict_labels(dists, k=5)
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
```

Got 139 / 500 correct => accuracy: 0.278000

You should expect to see a slightly better performance than with k = 1.

### Inline Question 2

We can also use other distance metrics such as L1 distance. For pixel values  $p_{ij}^{(k)}$  at location  $(i, j)$  of some image  $I_k$ ,

the mean  $\mu$  across all pixels over all images is

$$\mu = \frac{1}{nhw} \sum_{k=1}^n \sum_{i=1}^h \sum_{j=1}^w p_{ij}^{(k)}$$

And the pixel-wise mean  $\mu_{ij}$  across all images is

$$\mu_{ij} = \frac{1}{n} \sum_{k=1}^n p_{ij}^{(k)}.$$

The general standard deviation  $\sigma$  and pixel-wise standard deviation  $\sigma_{ij}$  is defined similarly.

Which of the following preprocessing steps will not change the performance of a Nearest Neighbor classifier that uses L1 distance? Select all that apply. 1. Subtracting the mean  $\mu$  ( $\tilde{p}_{ij}^{(k)} = p_{ij}^{(k)} - \mu$ .) 2. Subtracting the per pixel mean  $\mu_{ij}$  ( $\tilde{p}_{ij}^{(k)} = p_{ij}^{(k)} - \mu_{ij}$ .) 3. Subtracting the mean  $\mu$  and dividing by the standard deviation  $\sigma$ . 4. Subtracting the pixel-wise mean  $\mu_{ij}$  and dividing by the pixel-wise standard deviation  $\sigma_{ij}$ . 5. Rotating the coordinate axes of the data.

*Your Answer* : 1,3

*Your Explanation* :

```
[50]: # Now lets speed up distance matrix computation by using partial vectorization
# with one loop. Implement the function compute_distances_one_loop and run the
# code below:
dists_one = classifier.compute_distances_one_loop(X_test)

# To ensure that our vectorized implementation is correct, we make sure that it
# agrees with the naive implementation. There are many ways to decide whether
# two matrices are similar; one of the simplest is the Frobenius norm. In case
# you haven't seen it before, the Frobenius norm of two matrices is the square
# root of the squared sum of differences of all elements; in other words,
# → reshape
# the matrices into vectors and compute the Euclidean distance between them.
difference = np.linalg.norm(dists - dists_one, ord='fro')
print('One loop difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')
```

One loop difference was: 0.000000

Good! The distance matrices are the same

```
[51]: # Now implement the fully vectorized version inside compute_distances_no_loops
# and run the code
dists_two = classifier.compute_distances_no_loops(X_test)

# check that the distance matrix agrees with the one we computed before:
difference = np.linalg.norm(dists - dists_two, ord='fro')
print('No loop difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')
```

No loop difference was: 0.000000

Good! The distance matrices are the same

```
[52]: # Let's compare how fast the implementations are
def time_function(f, *args):
    """
    Call a function f with args and return the time (in seconds) that it took
    →to execute.
    """
    import time
    tic = time.time()
    f(*args)
    toc = time.time()
    return toc - tic

two_loop_time = time_function(classifier.compute_distances_two_loops, X_test)
print('Two loop version took %f seconds' % two_loop_time)

one_loop_time = time_function(classifier.compute_distances_one_loop, X_test)
print('One loop version took %f seconds' % one_loop_time)

no_loop_time = time_function(classifier.compute_distances_no_loops, X_test)
print('No loop version took %f seconds' % no_loop_time)

# You should see significantly faster performance with the fully vectorized
→implementation!

# NOTE: depending on what machine you're using,
# you might not see a speedup when you go from two loops to one loop,
# and might even see a slow-down.
```

```
Two loop version took 39.096850 seconds
One loop version took 36.872584 seconds
No loop version took 0.509463 seconds
```

### 1.0.1 Cross-validation

We have implemented the k-Nearest Neighbor classifier but we set the value  $k = 5$  arbitrarily. We will now determine the best value of this hyperparameter with cross-validation.

```
[59]: num_folds = 5
k_choices = [1, 3, 5, 8, 10, 12, 15, 20, 50, 100]

X_train_folds = []
y_train_folds = []
#####
# TODO:
# Split up the training data into folds. After splitting, X_train_folds and
# y_train_folds should each be lists of length num_folds, where
# y_train_folds[i] is the label vector for the points in X_train_folds[i].
```



```

# Hint: Look up the numpy array_split function. #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
X_train_folds=np.array_split(X_train, num_folds)
y_train_folds=np.array_split(y_train, num_folds)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# A dictionary holding the accuracies for different values of k that we find
# when running cross-validation. After running cross-validation,
# k_to_accuracies[k] should be a list of length num_folds giving the different
# accuracy values that we found when using that value of k.
k_to_accuracies = {}

#####
# TODO: #
# Perform k-fold cross validation to find the best value of k. For each #
# possible value of k, run the k-nearest-neighbor algorithm num_folds times, #
# where in each case you use all but one of the folds as training data and the #
# last fold as a validation set. Store the accuracies for all fold and all #
# values of k in the k_to_accuracies dictionary. #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
for k in k_choices:
    k_to_accuracies[k]=[]
    for j in range(num_folds):
        X_train_fold=np.concatenate([x for num,x in enumerate(X_train_folds) if num!
        ==j])
        y_train_fold=np.concatenate([y for num,y in enumerate(y_train_folds) if num!
        ==j])

        classifier.train(X_train_fold, y_train_fold)
        y_fold_pred = classifier.predict(X_train_folds[j],k=k,num_loops=0)
        num_correct = np.sum(y_fold_pred == y_train_folds[j])

        accuracy = float(num_correct) / X_train_folds[j].shape[0]
        k_to_accuracies[k].append(accuracy)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out the computed accuracies
for k in sorted(k_to_accuracies):
    for accuracy in k_to_accuracies[k]:
        print('k = %d, accuracy = %f' % (k, accuracy))

```


k = 1, accuracy = 0.263000

k = 1, accuracy = 0.257000

k = 1, accuracy = 0.264000  
k = 1, accuracy = 0.278000  
k = 1, accuracy = 0.266000  
k = 3, accuracy = 0.239000  
k = 3, accuracy = 0.249000  
k = 3, accuracy = 0.240000  
k = 3, accuracy = 0.266000  
k = 3, accuracy = 0.254000  
k = 5, accuracy = 0.248000  
k = 5, accuracy = 0.266000  
k = 5, accuracy = 0.280000  
k = 5, accuracy = 0.292000  
k = 5, accuracy = 0.280000  
k = 8, accuracy = 0.262000  
k = 8, accuracy = 0.282000  
k = 8, accuracy = 0.273000  
k = 8, accuracy = 0.290000  
k = 8, accuracy = 0.273000  
k = 10, accuracy = 0.265000  
k = 10, accuracy = 0.296000  
k = 10, accuracy = 0.276000  
k = 10, accuracy = 0.284000  
k = 10, accuracy = 0.280000  
k = 12, accuracy = 0.260000  
k = 12, accuracy = 0.295000  
k = 12, accuracy = 0.279000  
k = 12, accuracy = 0.283000  
k = 12, accuracy = 0.280000  
k = 15, accuracy = 0.252000  
k = 15, accuracy = 0.289000  
k = 15, accuracy = 0.278000  
k = 15, accuracy = 0.282000  
k = 15, accuracy = 0.274000  
k = 20, accuracy = 0.270000  
k = 20, accuracy = 0.279000  
k = 20, accuracy = 0.279000  
k = 20, accuracy = 0.282000  
k = 20, accuracy = 0.285000  
k = 50, accuracy = 0.271000  
k = 50, accuracy = 0.288000  
k = 50, accuracy = 0.278000  
k = 50, accuracy = 0.269000  
k = 50, accuracy = 0.266000  
k = 100, accuracy = 0.256000  
k = 100, accuracy = 0.270000  
k = 100, accuracy = 0.263000  
k = 100, accuracy = 0.256000  
k = 100, accuracy = 0.263000

```
[60]: # plot the raw observations
for k in k_choices:
    accuracies = k_to_accuracies[k]
    plt.scatter([k] * len(accuracies), accuracies)

# plot the trend line with error bars that correspond to standard deviation
accuracies_mean = np.array([np.mean(v) for k,v in sorted(k_to_accuracies.
    ↳items())])
accuracies_std = np.array([np.std(v) for k,v in sorted(k_to_accuracies.
    ↳items())])
plt.errorbar(k_choices, accuracies_mean, yerr=accuracies_std)
plt.title('Cross-validation on k')
plt.xlabel('k')
plt.ylabel('Cross-validation accuracy')
plt.show()
```



```
[61]: # Based on the cross-validation results above, choose the best value for k,
# retrain the classifier using all the training data, and test it on the test
# data. You should be able to get above 28% accuracy on the test data.
```

```

best_k = 10

classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
y_test_pred = classifier.predict(X_test, k=best_k)

# Compute and display the accuracy
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))

```

Got 141 / 500 correct => accuracy: 0.282000

### Inline Question 3

Which of the following statements about  $k$ -Nearest Neighbor ( $k$ -NN) are true in a classification setting, and for all  $k$ ? Select all that apply. 1. The decision boundary of the  $k$ -NN classifier is linear. 2. The training error of a 1-NN will always be lower than or equal to that of 5-NN. 3. The test error of a 1-NN will always be lower than that of a 5-NN. 4. The time needed to classify a test example with the  $k$ -NN classifier grows with the size of the training set. 5. None of the above.

*Your Answer :* 2, 4

*Your Explanation :*

## SVM

July 25, 2022

```
[1]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'cs231n/assignments/assignment1/'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

```
Mounted at /content/drive
/content/drive/My Drive/cs231n/assignments/assignment1/cs231n/datasets
/content/drive/My Drive/cs231n/assignments/assignment1
```

## 1 Multiclass Support Vector Machine exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

In this exercise you will:

- implement a fully-vectorized **loss function** for the SVM
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** using numerical gradient
- use a validation set to **tune the learning rate and regularization strength**

- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
[3]: # Run some setup code for this notebook.
import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

# This is a bit of magic to make matplotlib figures appear inline in the
# notebook rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/questions/1907993/
# ↪ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

## 1.1 CIFAR-10 Data Loading and Preprocessing

```
[4]: # Load the raw CIFAR-10 data.
cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

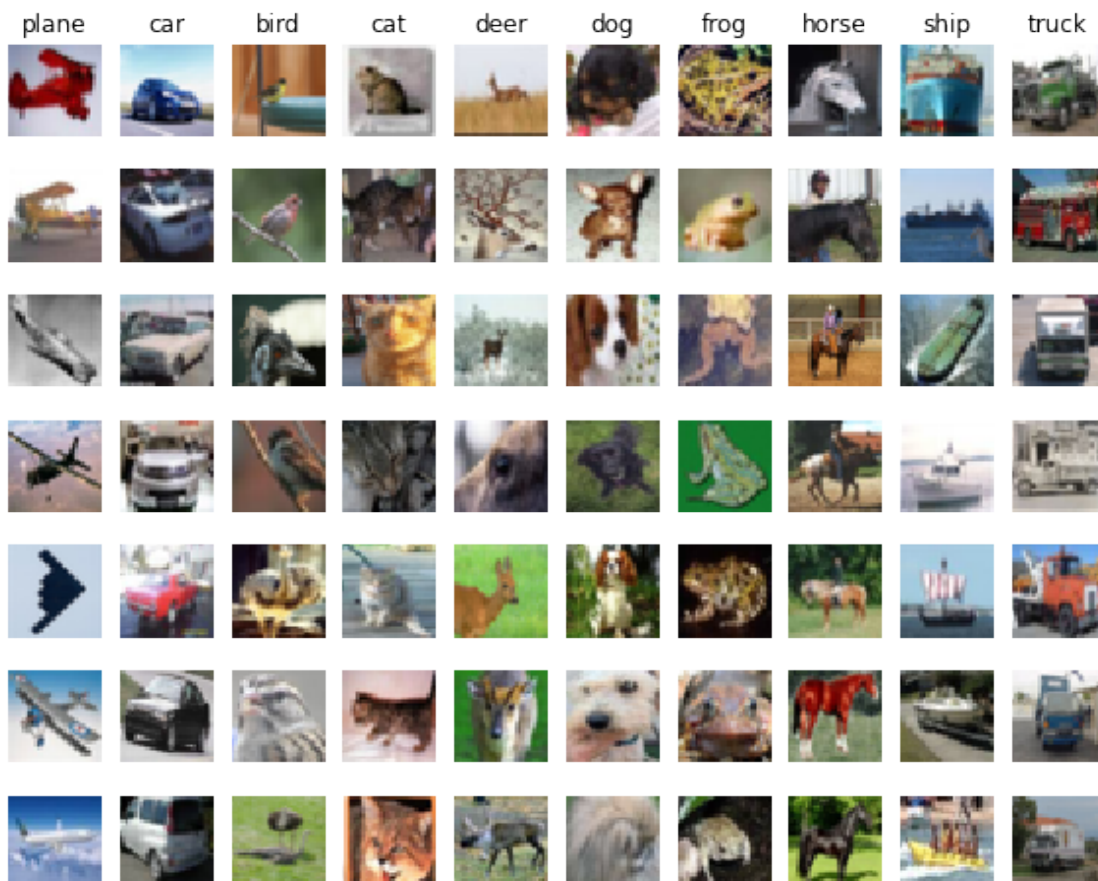
# Cleaning up variables to prevent loading data multiple times (which may cause
# ↪ memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# As a sanity check, we print out the size of the training and test data.
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

Training data shape: (50000, 32, 32, 3)  
 Training labels shape: (50000,)  
 Test data shape: (10000, 32, 32, 3)  
 Test labels shape: (10000,)

```
[5]: # Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
    if i == 0:
        plt.title(cls)
plt.show()
```



```
[6]: # Split the data into train, val, and test sets. In addition we will
# create a small development set as a subset of the training data;
# we can use this for development so our code runs faster.
num_training = 49000
num_validation = 1000
num_test = 1000
num_dev = 500

# Our validation set will be num_validation points from the original
# training set.
mask = range(num_training, num_training + num_validation)
X_val = X_train[mask]
y_val = y_train[mask]

# Our training set will be the first num_train points from the original
# training set.
mask = range(num_training)
X_train = X_train[mask]
y_train = y_train[mask]

# We will also make a development set, which is a small subset of
# the training set.
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

# We use the first num_test points of the original test set as our
# test set.
mask = range(num_test)
X_test = X_test[mask]
y_test = y_test[mask]

print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Train data shape: (49000, 32, 32, 3)
Train labels shape: (49000,)
Validation data shape: (1000, 32, 32, 3)
Validation labels shape: (1000,)
Test data shape: (1000, 32, 32, 3)
Test labels shape: (1000,)
```



```
[7]: # Preprocessing: reshape the image data into rows, 3 dimension to 1 dimension
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# As a sanity check, print out the shapes of the data
print('Training data shape: ', X_train.shape)
print('Validation data shape: ', X_val.shape)
print('Test data shape: ', X_test.shape)
print('dev data shape: ', X_dev.shape)
```

```
Training data shape: (49000, 3072)
Validation data shape: (1000, 3072)
Test data shape: (1000, 3072)
dev data shape: (500, 3072)
```


```
[ ]: # Preprocessing: subtract the mean image
# first: compute the image mean based on the training data
mean_image = np.mean(X_train, axis=0)
print(mean_image[:10]) # print a few of the elements
plt.figure(figsize=(4,4))
plt.imshow(mean_image.reshape((32,32,3)).astype('uint8')) # visualize the mean_
    ↪ image
plt.show()

# second: subtract the mean image from train and test data
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image

# third: append the bias dimension of ones (i.e. bias trick) so that our SVM
# only has to worry about optimizing a single weight matrix W.
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

print(X_train.shape, X_val.shape, X_test.shape, X_dev.shape)
```

```
[130.64189796 135.98173469 132.47391837 130.05569388 135.34804082
131.75402041 130.96055102 136.14328571 132.47636735 131.48467347]
```



(49000, 3073) (1000, 3073) (1000, 3073) (500, 3073)

## 1.2 SVM Classifier

Your code for this section will all be written inside `cs231n/classifiers/linear_svm.py`.

As you can see, we have prefilled the function `svm_loss_naive` which uses for loops to evaluate the multiclass SVM loss function.

```
[ ]: # Evaluate the naive implementation of the loss we provided for you:
      from cs231n.classifiers.linear_svm import svm_loss_naive
      import time

      # generate a random SVM weight matrix of small numbers
      W = np.random.randn(3073, 10) * 0.0001

      loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.000005)
      print('loss: %f' % (loss, ))
```

loss: 9.301101

The `grad` returned from the function above is right now all zero. Derive and implement the gradient for the SVM cost function and implement it inline inside the function `svm_loss_naive`. You will find it helpful to interleave your new code inside the existing function.

To check that you have correctly implemented the gradient, you can numerically estimate the gradient of the loss function and compare the numeric estimate to the gradient that you computed.

We have provided code that does this for you:

```
[ ]: # Once you've implemented the gradient, recompute it with the code below
# and gradient check it with the function we provided for you

# Compute the loss and its gradient at W.
loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.0)

# Numerically compute the gradient along several randomly chosen dimensions, and
# compare them with your analytically computed gradient. The numbers should
  ↳ match
# almost exactly along all dimensions.
from cs231n.gradient_check import grad_check_sparse
f = lambda w: svm_loss_naive(w, X_dev, y_dev, 0.0)[0] # no regularization
grad_numerical = grad_check_sparse(f, W, grad)

# do the gradient check once again with regularization turned on
# you didn't forget the regularization gradient did you?
loss, grad = svm_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: svm_loss_naive(w, X_dev, y_dev, 5e1)[0] #with regularization
grad_numerical = grad_check_sparse(f, W, grad)
```

```
numerical: 2.018202 analytic: 2.018202, relative error: 8.491378e-11
numerical: -28.776441 analytic: -28.776441, relative error: 1.151052e-11
numerical: -0.940961 analytic: -0.940961, relative error: 1.236439e-11
numerical: 9.891457 analytic: 9.891457, relative error: 1.295984e-11
numerical: 15.930939 analytic: 15.930939, relative error: 1.152546e-11
numerical: 8.977824 analytic: 8.977824, relative error: 7.242684e-12
numerical: -5.304467 analytic: -5.304467, relative error: 2.949141e-11
numerical: -2.566139 analytic: -2.566139, relative error: 5.327086e-11
numerical: 2.854758 analytic: 2.854758, relative error: 4.016735e-11
numerical: -22.276857 analytic: -22.276857, relative error: 1.663018e-11
numerical: -31.831138 analytic: -31.831138, relative error: 5.006324e-12
numerical: -11.202800 analytic: -11.202800, relative error: 3.586369e-11
numerical: -2.681966 analytic: -2.681966, relative error: 8.087692e-11
numerical: 25.047923 analytic: 25.047923, relative error: 1.563657e-11
numerical: 4.983970 analytic: 4.983970, relative error: 2.553056e-11
numerical: 8.284672 analytic: 8.284672, relative error: 3.416347e-11
numerical: -34.272393 analytic: -34.272393, relative error: 1.044645e-11
numerical: 1.311678 analytic: 1.311678, relative error: 2.635768e-11
numerical: -45.238797 analytic: -45.238797, relative error: 2.248698e-12
numerical: 4.485481 analytic: 4.485481, relative error: 2.767246e-11
```

### Inline Question 1

It is possible that once in a while a dimension in the gradcheck will not match exactly. What could such a discrepancy be caused by? Is it a reason for concern? What is a simple example in one dimension where a gradient check could fail? How would change the margin affect the frequency of this happening? *Hint: the SVM loss function is not strictly speaking differentiable*

*Your Answer* : if X value is between 0 and 1 and margin is 1, it is impossible to gradcheck

```
[ ]: # Next implement the function svm_loss_vectorized; for now only compute the
      ↪ loss;
      # we will implement the gradient in a moment.
      tic = time.time()
      loss_naive, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('Naive loss: %e computed in %fs' % (loss_naive, toc - tic))

      from cs231n.classifiers.linear_svm import svm_loss_vectorized
      tic = time.time()
      loss_vectorized, _ = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('Vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

      # The losses should match but your vectorized implementation should be much
      ↪ faster.
      print('difference: %f' % (loss_naive - loss_vectorized))
```

```
Naive loss: 9.301101e+00 computed in 0.137178s
Vectorized loss: 9.301101e+00 computed in 0.015659s
difference: -0.000000
```

```
[ ]: # Complete the implementation of svm_loss_vectorized, and compute the gradient
      # of the loss function in a vectorized way.

      # The naive implementation and the vectorized implementation should match, but
      # the vectorized version should still be much faster.
      tic = time.time()
      _, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('Naive loss and gradient: computed in %fs' % (toc - tic))

      tic = time.time()
      _, grad_vectorized = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('Vectorized loss and gradient: computed in %fs' % (toc - tic))

      # The loss is a single number, so it is easy to compare the values computed
      # by the two implementations. The gradient on the other hand is a matrix, so
      # we use the Frobenius norm to compare them.
      difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
      print('difference: %f' % difference)
```

```
Naive loss and gradient: computed in 0.111674s
Vectorized loss and gradient: computed in 0.012228s
difference: 0.000000
```


### 1.2.1 Stochastic Gradient Descent

We now have vectorized and efficient expressions for the loss, the gradient and our gradient matches the numerical gradient. We are therefore ready to do SGD to minimize the loss. Your code for this part will be written inside `cs231n/classifiers/linear_classifier.py`.

```
[ ]: # In the file linear_classifier.py, implement SGD in the function  
# LinearClassifier.train() and then run it with the code below.  
from cs231n.classifiers import LinearSVM  
svm = LinearSVM()  
tic = time.time()  
loss_hist = svm.train(X_train, y_train, learning_rate=1e-7, reg=2.5e4,  
                      num_iters=1500, verbose=True)  
toc = time.time()  
print('That took %fs' % (toc - tic))
```

```
iteration 0 / 1500: loss 791.174059  
iteration 100 / 1500: loss 286.889960  
iteration 200 / 1500: loss 107.923176  
iteration 300 / 1500: loss 42.361032  
iteration 400 / 1500: loss 18.678993  
iteration 500 / 1500: loss 10.026851  
iteration 600 / 1500: loss 6.856941  
iteration 700 / 1500: loss 5.599114  
iteration 800 / 1500: loss 5.519052  
iteration 900 / 1500: loss 5.737179  
iteration 1000 / 1500: loss 5.331880  
iteration 1100 / 1500: loss 5.330477  
iteration 1200 / 1500: loss 5.357079  
iteration 1300 / 1500: loss 5.502650  
iteration 1400 / 1500: loss 5.275219  
That took 12.869851s
```

```
[ ]: # A useful debugging strategy is to plot the loss as a function of  
# iteration number:  
plt.plot(loss_hist)  
plt.xlabel('Iteration number')  
plt.ylabel('Loss value')  
plt.show()
```



```
[ ]: # Write the LinearSVM.predict function and evaluate the performance on both the
# training and validation set
y_train_pred = svm.predict(X_train)
print('training accuracy: %f' % (np.mean(y_train == y_train_pred), ))
y_val_pred = svm.predict(X_val)
print('validation accuracy: %f' % (np.mean(y_val == y_val_pred), ))
```

```
training accuracy: 0.366878
validation accuracy: 0.384000
```

```
[ ]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of about 0.39 (> 0.385) on the validation set.

# Note: you may see runtime/overflow warnings during hyper-parameter search.
# This may be caused by extreme values, and is not a bug.

# results is dictionary mapping tuples of the form
# (learning_rate, regularization_strength) to tuples of the form
```

```

# (training_accuracy, validation_accuracy). The accuracy is simply the fraction
# of data points that are correctly classified.
results = {}
best_val = -1 # The highest validation accuracy that we have seen so far.
best_svm = None # The LinearSVM object that achieved the highest validation
    ↳rate.

#####
# TODO:
# Write code that chooses the best hyperparameters by tuning on the validation #
# set. For each combination of hyperparameters, train a linear SVM on the #
# training set, compute its accuracy on the training and validation sets, and #
# store these numbers in the results dictionary. In addition, store the best #
# validation accuracy in best_val and the LinearSVM object that achieves this #
# accuracy in best_svm.
#
# Hint: You should use a small value for num_iters as you develop your #
# validation code so that the SVMs don't take much time to train; once you are #
# confident that your validation code works, you should rerun the validation #
# code with a larger value for num_iters.
#####

# Provided as a reference. You may or may not want to change these
    ↳hyperparameters
learning_rates = [1e-7, 5e-5]
regularization_strengths = [2.5e4, 5e4]

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
for lr in learning_rates:
    for rs in regularization_strengths:
        svm = LinearSVM()
        svm.train(X_train, y_train, learning_rate=lr, reg=rs,
                  num_iters=1500, verbose=True)
        y_train_pred = svm.predict(X_train)
        train_accuracy= np.mean(y_train == y_train_pred)
        y_val_pred = svm.predict(X_val)
        validation_accuracy= np.mean(y_val == y_val_pred)
        results[(lr, rs)] = (train_accuracy, validation_accuracy)
        if validation_accuracy>best_val:
            best_val=validation_accuracy
            best_svm = svm

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):

```

```

train_accuracy, val_accuracy = results[(lr, reg)]
print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
    lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' %
      ↪best_val)

```

```

iteration 0 / 1500: loss 801.455447
iteration 100 / 1500: loss 292.376697
iteration 200 / 1500: loss 109.053176
iteration 300 / 1500: loss 42.844784
iteration 400 / 1500: loss 19.025903
iteration 500 / 1500: loss 10.064818
iteration 600 / 1500: loss 7.050020
iteration 700 / 1500: loss 5.801469
iteration 800 / 1500: loss 5.642600
iteration 900 / 1500: loss 5.294671
iteration 1000 / 1500: loss 5.613706
iteration 1100 / 1500: loss 5.248545
iteration 1200 / 1500: loss 5.360206
iteration 1300 / 1500: loss 5.366830
iteration 1400 / 1500: loss 5.512912
iteration 0 / 1500: loss 1545.335290
iteration 100 / 1500: loss 210.023252
iteration 200 / 1500: loss 32.505675
iteration 300 / 1500: loss 9.259946
iteration 400 / 1500: loss 6.198078
iteration 500 / 1500: loss 4.959885
iteration 600 / 1500: loss 5.442797
iteration 700 / 1500: loss 5.745977
iteration 800 / 1500: loss 5.697453
iteration 900 / 1500: loss 5.941142
iteration 1000 / 1500: loss 5.742636
iteration 1100 / 1500: loss 5.556872
iteration 1200 / 1500: loss 6.372953
iteration 1300 / 1500: loss 5.812734
iteration 1400 / 1500: loss 5.309607
iteration 0 / 1500: loss 795.031183
iteration 100 / 1500: loss 399986837018047105224874766387726254080.000000
iteration 200 / 1500: loss 66114620696033832252575901444739326472487341047685277
373847996760548442112.000000
iteration 300 / 1500: loss 10928217294268622138261895924185436716614004294162244
676543075668342579324135513835055384954274543164547137536.000000
iteration 400 / 1500: loss 18063467955116930485042438404427141173728138435673966
71914815983277240892655158288032543843898923394325010666981922215016146711810461
650430984192.000000
iteration 500 / 1500: loss 29857465840898010118650269295620268088706721978594533

```



```

06817177977976506957547380296942790781885676448949978846207018839507297985403204
66427945617105924008698012137839088019909378048.000000
iteration 600 / 1500: loss 49351999774099900890322982781166763525584265176433099
94625306644616530492276263077715703965504776591569726636096567632800944368812515
09511140930283500931787418092665258226898457412814074452325120328620142251852431
36.000000
iteration 700 / 1500: loss 81574903063826186167878845582100305488228195439022673
72533975739279403012534331666197417403885286490114730014662324825350312752959381
41753504908161099443407522374495805228719367286023184056231932184309729885212139
1854187882837626237155010953978314752.000000
iteration 800 / 1500: loss 13483678149481840333432729879422314204924613731828207
97831922690270875739899661148831131252735074422672343479374020795631994418560845
83056214165924834055849596487328314468925655730527656475196112626996667327859724
4827859856126682416992845282173898705582622990185156330095350337813610496.000000

/content/drive/My
Drive/cs231n/assignments/assignment1/cs231n/classifiers/linear_svm.py:93:
RuntimeWarning: overflow encountered in double_scalars
    loss += reg * np.sum(W * W)
/usr/local/lib/python3.7/dist-packages/numpy/core/fromnumeric.py:86:
RuntimeWarning: overflow encountered in reduce
    return ufunc.reduce(obj, axis, dtype, out, **passkwargs)
/content/drive/My
Drive/cs231n/assignments/assignment1/cs231n/classifiers/linear_svm.py:93:
RuntimeWarning: overflow encountered in multiply
    loss += reg * np.sum(W * W)

iteration 900 / 1500: loss inf
iteration 1000 / 1500: loss inf
iteration 1100 / 1500: loss inf
iteration 1200 / 1500: loss inf
iteration 1300 / 1500: loss inf
iteration 1400 / 1500: loss inf
iteration 0 / 1500: loss 1552.473676
iteration 100 / 1500: loss 42376358128369431068679955573141306540533375447852189
37413489496202200411463401288631885958870483869196587677791408881664000.000000
iteration 200 / 1500: loss 10942634561074914123571895876076447512441019495249007
24258146371783342597861137011821992623636884777090648327917144586977388373196775
10165748348643932968619948503107808829494942239106301842079949630049023718690575
12561403103089970159407857664000.000000
iteration 300 / 1500: loss inf
iteration 400 / 1500: loss inf
iteration 500 / 1500: loss inf

/content/drive/My
Drive/cs231n/assignments/assignment1/cs231n/classifiers/linear_svm.py:113:
RuntimeWarning: overflow encountered in multiply
    dW += reg*2*W

iteration 600 / 1500: loss nan

```

```

iteration 700 / 1500: loss nan
iteration 800 / 1500: loss nan
iteration 900 / 1500: loss nan
iteration 1000 / 1500: loss nan
iteration 1100 / 1500: loss nan
iteration 1200 / 1500: loss nan
iteration 1300 / 1500: loss nan
iteration 1400 / 1500: loss nan
lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.365020 val accuracy: 0.374000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.355816 val accuracy: 0.358000
lr 5.000000e-05 reg 2.500000e+04 train accuracy: 0.079857 val accuracy: 0.073000
lr 5.000000e-05 reg 5.000000e+04 train accuracy: 0.100265 val accuracy: 0.087000
best validation accuracy achieved during cross-validation: 0.374000

```

```

[ ]: # Visualize the cross-validation results
import math
import pdb


# pdb.set_trace()

x_scatter = [math.log10(x[0]) for x in results]
y_scatter = [math.log10(x[1]) for x in results]

# plot training accuracy
marker_size = 100
colors = [results[x][0] for x in results]
plt.subplot(2, 1, 1)
plt.tight_layout(pad=3)
plt.scatter(x_scatter, y_scatter, marker_size, c=colors, cmap=plt.cm.coolwarm)
plt.colorbar()
plt.xlabel('log learning rate')
plt.ylabel('log regularization strength')
plt.title('CIFAR-10 training accuracy')

# plot validation accuracy
colors = [results[x][1] for x in results] # default size of markers is 20
plt.subplot(2, 1, 2)
plt.scatter(x_scatter, y_scatter, marker_size, c=colors, cmap=plt.cm.coolwarm)
plt.colorbar()
plt.xlabel('log learning rate')
plt.ylabel('log regularization strength')
plt.title('CIFAR-10 validation accuracy')
plt.show()

```



```
[ ]: # Evaluate the best svm on test set
y_test_pred = best_svm.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('linear SVM on raw pixels final test set accuracy: %f' % test_accuracy)
```

linear SVM on raw pixels final test set accuracy: 0.354000

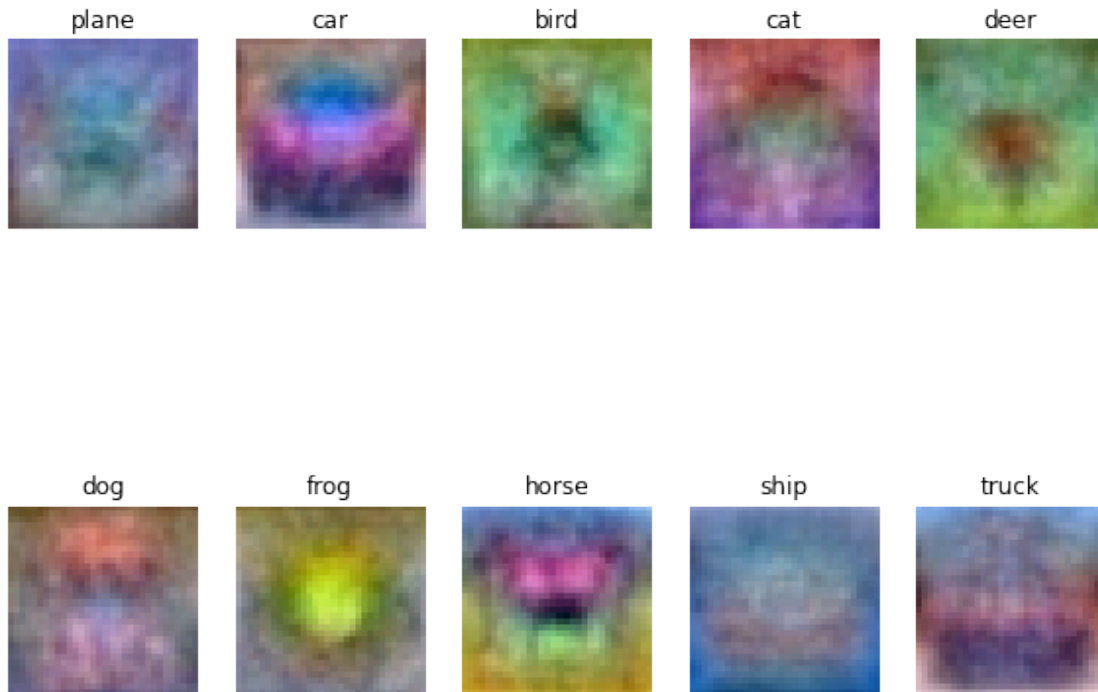
```
[ ]: # Visualize the learned weights for each class.
# Depending on your choice of learning rate and regularization strength, these
    ↪ may
# or may not be nice to look at.
w = best_svm.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)
w_min, w_max = np.min(w), np.max(w)
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse',
    ↪ 'ship', 'truck']
```

```

for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])

```



## Inline question 2

Describe what your visualized SVM weights look like, and offer a brief explanation for why they look the way they do.

*Your Answer :* we can briefly predict what weights explain by feature of SVM weights. ex) there are lots of green part in frog weight.

# softmax

July 25, 2022

```
[17]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'cs231n/assignments/assignment1/'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.  
/content/drive/My Drive/cs231n/assignments/assignment1/cs231n/datasets  
/content/drive/My Drive/cs231n/assignments/assignment1

## 1 Softmax exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

This exercise is analogous to the SVM exercise. You will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient

- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
[18]: import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
  ↳ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

```
[19]: def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000,
  ↳ num_dev=500):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the linear classifier. These are the same steps as we used for the
    SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

    # Cleaning up variables to prevent loading data multiple times (which may
    ↳ cause memory issue)
    try:
        del X_train, y_train
        del X_test, y_test
        print('Clear previously loaded data.')
    except:
        pass

    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
```

```

y_val = y_train[mask]
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

# Preprocessing: reshape the image data into rows: 3 dimension to 1
→ dimension
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# Normalize the data: subtract the mean image
mean_image = np.mean(X_train, axis = 0)
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image

# add bias dimension and transform into columns
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev =
→ get_CIFAR10_data()
print('Train data shape: ', X_train.shape) #49000 x 3073
print('Train labels shape: ', y_train.shape) #49000 x 1
print('Validation data shape: ', X_val.shape) #1000 x 3073
print('Validation labels shape: ', y_val.shape) #1000 x 1
print('Test data shape: ', X_test.shape) #1000 x 3073
print('Test labels shape: ', y_test.shape) #1000 x 1
print('dev data shape: ', X_dev.shape) #500 x 3073
print('dev labels shape: ', y_dev.shape) #500 x 1

```

Train data shape: (49000, 3073)

```
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)
```

## 1.1 Softmax Classifier

Your code for this section will all be written inside `cs231n/classifiers/softmax.py`.

```
[23]: # First implement the naive softmax loss function with nested loops.
      # Open the file cs231n/classifiers/softmax.py and implement the
      # softmax_loss_naive function.

      from cs231n.classifiers.softmax import softmax_loss_naive
      import time

      # Generate a random softmax weight matrix and use it to compute the loss.
      W = np.random.randn(3073, 10) * 0.0001
      loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

      # As a rough sanity check, our loss should be something close to -log(0.1).
      print('loss: %f' % loss)
      print('sanity check: %f' % (-np.log(0.1)))
```

```
loss: 2.368655
```

```
sanity check: 2.302585
```

### Inline Question 1

Why do we expect our loss to be close to  $-\log(0.1)$ ? Explain briefly.\*\*

*Your Answer* : there are 10 classes, so probability to detect correct answer is  $1/10$

```
[27]: # Complete the implementation of softmax_loss_naive and implement a (naive)
      # version of the gradient that uses nested loops.
      loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

      # As we did for the SVM, use numeric gradient checking as a debugging tool.
      # The numeric gradient should be close to the analytic gradient.
      from cs231n.gradient_check import grad_check_sparse
      f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 0.0)[0]
      grad_numerical = grad_check_sparse(f, W, grad, 10)

      # similar to SVM case, do another gradient check with regularization
      loss, grad = softmax_loss_naive(W, X_dev, y_dev, 5e1)
      f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 5e1)[0]
```



```
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

```
numerical: 1.133200 analytic: 1.133200, relative error: 2.979416e-08
numerical: 1.929890 analytic: 1.929890, relative error: 4.402701e-08
numerical: 1.789645 analytic: 1.789645, relative error: 1.829813e-08
numerical: 0.504461 analytic: 0.504461, relative error: 4.392747e-08
numerical: -2.590003 analytic: -2.590003, relative error: 2.046262e-08
numerical: 0.864664 analytic: 0.864664, relative error: 2.810858e-08
numerical: 1.724757 analytic: 1.724757, relative error: 2.893016e-08
numerical: -5.220883 analytic: -5.220883, relative error: 1.299076e-08
numerical: 1.472878 analytic: 1.472878, relative error: 2.100946e-09
numerical: 0.624058 analytic: 0.624058, relative error: 1.964306e-08
numerical: -2.311060 analytic: -2.311060, relative error: 7.892614e-09
numerical: 0.220874 analytic: 0.220874, relative error: 1.638380e-07
numerical: -1.642406 analytic: -1.642406, relative error: 1.624275e-08
numerical: 3.653459 analytic: 3.653459, relative error: 7.521597e-09
numerical: -4.855733 analytic: -4.855733, relative error: 1.342909e-09
numerical: -4.198332 analytic: -4.198331, relative error: 1.675519e-08
numerical: 3.866631 analytic: 3.866631, relative error: 4.099264e-09
numerical: -1.948127 analytic: -1.948127, relative error: 4.084492e-08
numerical: -1.904179 analytic: -1.904179, relative error: 2.629194e-09
numerical: 0.557432 analytic: 0.557432, relative error: 1.183513e-07
```

```
[37]: # Now that we have a naive implementation of the softmax loss function and its
      ↪ gradient,
      # implement a vectorized version in softmax_loss_vectorized.
      # The two versions should compute the same results, but the vectorized version
      ↪ should be
      # much faster.
      tic = time.time()
      loss_naive, grad_naive = softmax_loss_naive(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('naive loss: %e computed in %fs' % (loss_naive, toc - tic))

      from cs231n.classifiers.softmax import softmax_loss_vectorized
      tic = time.time()
      loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_dev, y_dev, 0.
      ↪ 000005)
      toc = time.time()
      print('vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

      # As we did for the SVM, we use the Frobenius norm to compare the two versions
      # of the gradient.
      grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
      print('Loss difference: %f' % np.abs(loss_naive - loss_vectorized))
      print('Gradient difference: %f' % grad_difference)
```

naive loss: 2.368655e+00 computed in 0.114715s  
vectorized loss: 2.368655e+00 computed in 0.023494s  
Loss difference: 0.000000  
Gradient difference: 0.000000

```
[39]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of over 0.35 on the validation set.

from cs231n.classifiers import Softmax
results = {}
best_val = -1
best_softmax = None

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength. #
# This should be identical to the validation that you did for the SVM; save #
# the best trained softmax classifier in best_softmax. #
#####

# Provided as a reference. You may or may not want to change these
# hyperparameters
learning_rates = [1e-7, 5e-7]
regularization_strengths = [2.5e4, 5e4]

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

for lr in learning_rates:
    for rs in regularization_strengths:
        sfx = Softmax()
        sfx.train(X_train, y_train, learning_rate=lr, reg=rs,
                  num_iters=1500, verbose=True)
        y_train_pred = sfx.predict(X_train)
        train_accuracy = np.mean(y_train == y_train_pred)
        y_val_pred = sfx.predict(X_val)
        validation_accuracy = np.mean(y_val == y_val_pred)
        results[(lr, rs)] = (train_accuracy, validation_accuracy)
        if validation_accuracy > best_val:
            best_val = validation_accuracy
            best_softmax = sfx

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):
```

```

train_accuracy, val_accuracy = results[(lr, reg)]
print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
    lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' %
      ↪best_val)

```

```

iteration 0 / 1500: loss 776.346208
iteration 100 / 1500: loss 284.915837
iteration 200 / 1500: loss 105.379845
iteration 300 / 1500: loss 39.982022
iteration 400 / 1500: loss 15.909922
iteration 500 / 1500: loss 7.153897
iteration 600 / 1500: loss 3.889405
iteration 700 / 1500: loss 2.815827
iteration 800 / 1500: loss 2.277305
iteration 900 / 1500: loss 2.189049
iteration 1000 / 1500: loss 2.154179
iteration 1100 / 1500: loss 2.103512
iteration 1200 / 1500: loss 2.027631
iteration 1300 / 1500: loss 2.061362
iteration 1400 / 1500: loss 2.118192
iteration 0 / 1500: loss 1538.723404
iteration 100 / 1500: loss 207.367120
iteration 200 / 1500: loss 29.580847
iteration 300 / 1500: loss 5.814450
iteration 400 / 1500: loss 2.569658
iteration 500 / 1500: loss 2.225496
iteration 600 / 1500: loss 2.147768
iteration 700 / 1500: loss 2.186226
iteration 800 / 1500: loss 2.123084
iteration 900 / 1500: loss 2.122774
iteration 1000 / 1500: loss 2.120958
iteration 1100 / 1500: loss 2.143974
iteration 1200 / 1500: loss 2.148913
iteration 1300 / 1500: loss 2.202296
iteration 1400 / 1500: loss 2.168277
iteration 0 / 1500: loss 767.287163
iteration 100 / 1500: loss 6.887006
iteration 200 / 1500: loss 2.102090
iteration 300 / 1500: loss 2.123372
iteration 400 / 1500: loss 2.136776
iteration 500 / 1500: loss 2.102515
iteration 600 / 1500: loss 2.134328
iteration 700 / 1500: loss 2.152824
iteration 800 / 1500: loss 2.054757
iteration 900 / 1500: loss 2.056627

```

```

iteration 1000 / 1500: loss 2.152816
iteration 1100 / 1500: loss 2.087911
iteration 1200 / 1500: loss 2.146382
iteration 1300 / 1500: loss 2.088835
iteration 1400 / 1500: loss 2.097501
iteration 0 / 1500: loss 1555.371411
iteration 100 / 1500: loss 2.258798
iteration 200 / 1500: loss 2.155087
iteration 300 / 1500: loss 2.175115
iteration 400 / 1500: loss 2.119690
iteration 500 / 1500: loss 2.088051
iteration 600 / 1500: loss 2.139761
iteration 700 / 1500: loss 2.178803
iteration 800 / 1500: loss 2.100830
iteration 900 / 1500: loss 2.121963
iteration 1000 / 1500: loss 2.151966
iteration 1100 / 1500: loss 2.115586
iteration 1200 / 1500: loss 2.123011
iteration 1300 / 1500: loss 2.186783
iteration 1400 / 1500: loss 2.170907
lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.329408 val accuracy: 0.343000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.309918 val accuracy: 0.322000
lr 5.000000e-07 reg 2.500000e+04 train accuracy: 0.329531 val accuracy: 0.339000
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.300837 val accuracy: 0.316000
best validation accuracy achieved during cross-validation: 0.343000

```

```

[40]: # evaluate on test set
# Evaluate the best softmax on test set
y_test_pred = best_softmax.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))

```

softmax on raw pixels final test set accuracy: 0.334000

### Inline Question 2 - True or False

Suppose the overall training loss is defined as the sum of the per-datapoint loss over all training examples. It is possible to add a new datapoint to a training set that would leave the SVM loss unchanged, but this is not the case with the Softmax classifier loss.

*Your Answer* : True

*Your Explanation* : unlike SVM, Softmax classifier compute sum of all values of datapoint.

```

[41]: # Visualize the learned weights for each class
w = best_softmax.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)

w_min, w_max = np.min(w), np.max(w)

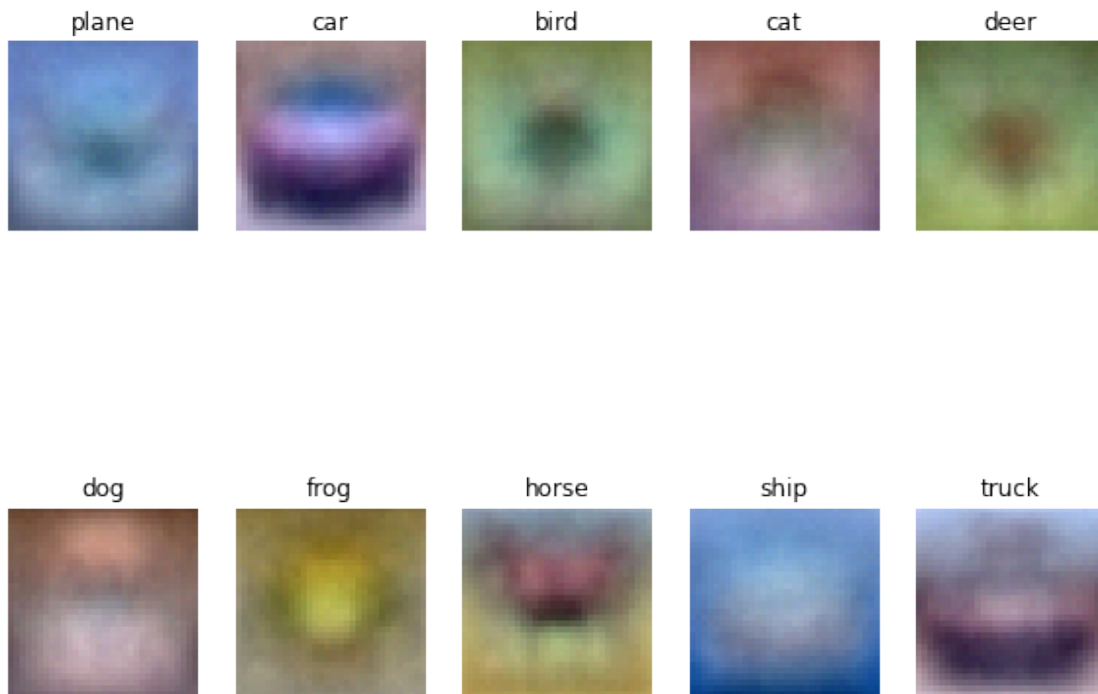
```

```

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])

```



[ ]:

# two\_layer\_net

July 25, 2022

```
[ ]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'cs231n/assignments/assignment1/'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.  
/content/drive/My Drive/cs231n/assignments/assignment1/cs231n/datasets  
/content/drive/My Drive/cs231n/assignments/assignment1

## 1 Fully-Connected Neural Nets

In this exercise we will implement fully-connected networks using a modular approach. For each layer we will implement a `forward` and a `backward` function. The `forward` function will receive inputs, weights, and other parameters and will return both an output and a `cache` object storing data needed for the backward pass, like this:

```
def layer_forward(x, w):
    """ Receive inputs x and weights w """
    # Do some computations ...
    z = # ... some intermediate value
```

```

# Do some more computations ...
out = # the output

cache = (x, w, z, out) # Values we need to compute gradients

return out, cache

```

The backward pass will receive upstream derivatives and the `cache` object, and will return gradients with respect to the inputs and weights, like this:

```

def layer_backward(dout, cache):
    """
    Receive dout (derivative of loss with respect to outputs) and cache,
    and compute derivative with respect to inputs.
    """
    # Unpack cache values
    x, w, z, out = cache

    # Use values in cache to compute derivatives
    dx = # Derivative of loss with respect to x
    dw = # Derivative of loss with respect to w

    return dx, dw

```

After implementing a bunch of layers this way, we will be able to easily combine them to build classifiers with different architectures.

```

[ ]: # As usual, a bit of setup
from __future__ import print_function
import time
import numpy as np
import matplotlib.pyplot as plt
from cs231n.classifiers.fc_net import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
→ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

```

```
def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

```
[ ]: # Load the (preprocessed) CIFAR10 data.
```

```
data = get_CIFAR10_data()
for k, v in list(data.items()):
    print('%s: ' % k, v.shape)
```

```
('X_train: ', (49000, 3, 32, 32))
('y_train: ', (49000,))
('X_val: ', (1000, 3, 32, 32))
('y_val: ', (1000,))
('X_test: ', (1000, 3, 32, 32))
('y_test: ', (1000,))
```

## 2 Affine layer: forward

Open the file `cs231n/layers.py` and implement the `affine_forward` function.

Once you are done you can test your implementation by running the following:

```
[ ]: # Test the affine_forward function
```

```
num_inputs = 2
input_shape = (4, 5, 6)
output_dim = 3

input_size = num_inputs * np.prod(input_shape)
weight_size = output_dim * np.prod(input_shape)

x = np.linspace(-0.1, 0.5, num=input_size).reshape(num_inputs, *input_shape)
w = np.linspace(-0.2, 0.3, num=weight_size).reshape(np.prod(input_shape),
    ↪output_dim)
b = np.linspace(-0.3, 0.1, num=output_dim)

out, _ = affine_forward(x, w, b)
correct_out = np.array([[ 1.49834967,  1.70660132,  1.91485297],
                        [ 3.25553199,  3.5141327,   3.77273342]])

# Compare your output with ours. The error should be around e-9 or less.
print('Testing affine_forward function:')
print('difference: ', rel_error(out, correct_out))
```



Testing affine\_forward function:  
difference: 9.769849468192957e-10

### 3 Affine layer: backward

Now implement the `affine_backward` function and test your implementation using numeric gradient checking.

```
[ ]: # Test the affine_backward function
np.random.seed(231)
x = np.random.randn(10, 2, 3)
w = np.random.randn(6, 5)
b = np.random.randn(5)
dout = np.random.randn(10, 5)

dx_num = eval_numerical_gradient_array(lambda x: affine_forward(x, w, b)[0], x,
    ↪dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_forward(x, w, b)[0], w,
    ↪dout)
db_num = eval_numerical_gradient_array(lambda b: affine_forward(x, w, b)[0], b,
    ↪dout)

_, cache = affine_forward(x, w, b)
dx, dw, db = affine_backward(dout, cache)

# The error should be around e-10 or less
print('Testing affine_backward function:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

Testing affine\_backward function:  
dx error: 5.399100368651805e-11  
dw error: 9.904211865398145e-11  
db error: 2.4122867568119087e-11

### 4 ReLU activation: forward

Implement the forward pass for the ReLU activation function in the `relu_forward` function and test your implementation using the following:

```
[ ]: # Test the relu_forward function

x = np.linspace(-0.5, 0.5, num=12).reshape(3, 4)
```

```

out, _ = relu_forward(x)
correct_out = np.array([[ 0.,          0.,          0.,          0.,          ],
                        [ 0.,          0.,          0.04545455, 0.13636364,],
                        [ 0.22727273, 0.31818182, 0.40909091, 0.5,          ]])

# Compare your output with ours. The error should be on the order of e-8
print('Testing relu_forward function:')
print('difference: ', rel_error(out, correct_out))

```

Testing relu\_forward function:  
 difference: 4.999999798022158e-08

## 5 ReLU activation: backward

Now implement the backward pass for the ReLU activation function in the `relu_backward` function and test your implementation using numeric gradient checking:

```

[ ]: np.random.seed(231)
x = np.random.randn(10, 10)
dout = np.random.randn(*x.shape)

dx_num = eval_numerical_gradient_array(lambda x: relu_forward(x)[0], x, dout)

_, cache = relu_forward(x)
dx = relu_backward(dout, cache)

# The error should be on the order of e-12
print('Testing relu_backward function:')
print('dx error: ', rel_error(dx_num, dx))

```

Testing relu\_backward function:  
 dx error: 3.2756349136310288e-12

### 5.1 Inline Question 1:

We've only asked you to implement ReLU, but there are a number of different activation functions that one could use in neural networks, each with its pros and cons. In particular, an issue commonly seen with activation functions is getting zero (or close to zero) gradient flow during backpropagation. Which of the following activation functions have this problem? If you consider these functions in the one dimensional case, what types of input would lead to this behaviour? 1. Sigmoid 2. ReLU 3. Leaky ReLU

### 5.2 Answer:

if input of sigmoid function is getting bigger, gradient is getting close to zero

## 6 "Sandwich" layers

There are some common patterns of layers that are frequently used in neural nets. For example, affine layers are frequently followed by a ReLU nonlinearity. To make these common patterns easy, we define several convenience layers in the file `cs231n/layer_utils.py`.

For now take a look at the `affine_relu_forward` and `affine_relu_backward` functions, and run the following to numerically gradient check the backward pass:

```
[ ]: from cs231n.layer_utils import affine_relu_forward, affine_relu_backward
np.random.seed(231)
x = np.random.randn(2, 3, 4)
w = np.random.randn(12, 10)
b = np.random.randn(10)
dout = np.random.randn(2, 10)

out, cache = affine_relu_forward(x, w, b)
dx, dw, db = affine_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: affine_relu_forward(x, w,
    ↪b)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_relu_forward(x, w,
    ↪b)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: affine_relu_forward(x, w,
    ↪b)[0], b, dout)

# Relative error should be around e-10 or less
print('Testing affine_relu_forward and affine_relu_backward:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

```
Testing affine_relu_forward and affine_relu_backward:
dx error:  2.299579177309368e-11
dw error:  8.162011105764925e-11
db error:  7.826724021458994e-12
```

## 7 Loss layers: Softmax and SVM

Now implement the loss and gradient for softmax and SVM in the `softmax_loss` and `svm_loss` function in `cs231n/layers.py`. These should be similar to what you implemented in `cs231n/classifiers/softmax.py` and `cs231n/classifiers/linear_svm.py`.

You can make sure that the implementations are correct by running the following:

```
[ ]: np.random.seed(231)
num_classes, num_inputs = 10, 50
```

```

x = 0.001 * np.random.randn(num_inputs, num_classes)
y = np.random.randint(num_classes, size=num_inputs)

dx_num = eval_numerical_gradient(lambda x: svm_loss(x, y)[0], x, verbose=False)
loss, dx = svm_loss(x, y)

# Test svm_loss function. Loss should be around 9 and dx error should be around
→ the order of e-9
print('Testing svm_loss:')
print('loss: ', loss)
print('dx error: ', rel_error(dx_num, dx))

dx_num = eval_numerical_gradient(lambda x: softmax_loss(x, y)[0], x,
→ verbose=False)
loss, dx = softmax_loss(x, y)

# Test softmax_loss function. Loss should be close to 2.3 and dx error should
→ be around e-8
print('\nTesting softmax_loss:')
print('loss: ', loss)
print('dx error: ', rel_error(dx_num, dx))

```

```

Testing svm_loss:
loss: 8.999602749096233
dx error: 1.4021566006651672e-09

```

```

Testing softmax_loss:
loss: 2.3025458445007376
dx error: 0.3333335566823834

```

## 8 Two-layer network

Open the file `cs231n/classifiers/fc_net.py` and complete the implementation of the `TwoLayerNet` class. Read through it to make sure you understand the API. You can run the cell below to test your implementation.

```

[ ]: np.random.seed(231)
N, D, H, C = 3, 5, 50, 7
X = np.random.randn(N, D)
y = np.random.randint(C, size=N)

std = 1e-3
model = TwoLayerNet(input_dim=D, hidden_dim=H, num_classes=C, weight_scale=std)

print('Testing initialization ... ')
W1_std = abs(model.params['W1'].std() - std)

```

```

b1 = model.params['b1']
W2_std = abs(model.params['W2'].std() - std)
b2 = model.params['b2']
assert W1_std < std / 10, 'First layer weights do not seem right'
assert np.all(b1 == 0), 'First layer biases do not seem right'
assert W2_std < std / 10, 'Second layer weights do not seem right'
assert np.all(b2 == 0), 'Second layer biases do not seem right'

print('Testing test-time forward pass ... ')
model.params['W1'] = np.linspace(-0.7, 0.3, num=D*H).reshape(D, H)
model.params['b1'] = np.linspace(-0.1, 0.9, num=H)
model.params['W2'] = np.linspace(-0.3, 0.4, num=H*C).reshape(H, C)
model.params['b2'] = np.linspace(-0.9, 0.1, num=C)
X = np.linspace(-5.5, 4.5, num=N*D).reshape(D, N).T
scores = model.loss(X)
correct_scores = np.asarray(
    [[11.53165108, 12.2917344, 13.05181771, 13.81190102, 14.57198434, 15.
    ↪33206765, 16.09215096],
    [12.05769098, 12.74614105, 13.43459113, 14.1230412, 14.81149128, 15.
    ↪49994135, 16.18839143],
    [12.58373087, 13.20054771, 13.81736455, 14.43418138, 15.05099822, 15.
    ↪66781506, 16.2846319 ]])
scores_diff = np.abs(scores - correct_scores).sum()
assert scores_diff < 1e-6, 'Problem with test-time forward pass'

print('Testing training loss (no regularization)')
y = np.asarray([0, 5, 1])
loss, grads = model.loss(X, y)
correct_loss = 3.4702243556
assert abs(loss - correct_loss) < 1e-10, 'Problem with training-time loss'

model.reg = 1.0
loss, grads = model.loss(X, y)
correct_loss = 26.5948426952
assert abs(loss - correct_loss) < 1e-10, 'Problem with regularization loss'

# Errors should be around e-7 or less
for reg in [0.0, 0.7]:
    print('Running numeric gradient check with reg = ', reg)
    model.reg = reg
    loss, grads = model.loss(X, y)

    for name in sorted(grads):
        f = lambda _: model.loss(X, y)[0]
        grad_num = eval_numerical_gradient(f, model.params[name], verbose=False)
        print('%s relative error: %.2e' % (name, rel_error(grad_num, grads[name])))

```

```

Testing initialization ...
Testing test-time forward pass ...
Testing training loss (no regularization)
Running numeric gradient check with reg = 0.0
W1 relative error: 1.83e-08
W2 relative error: 3.20e-10
b1 relative error: 9.83e-09
b2 relative error: 4.33e-10
Running numeric gradient check with reg = 0.7
W1 relative error: 2.53e-07
W2 relative error: 7.98e-08
b1 relative error: 1.35e-08
b2 relative error: 7.76e-10

```

## 9 Solver

Open the file `cs231n/solver.py` and read through it to familiarize yourself with the API. You also need to implement the `sgd` function in `cs231n/optim.py`. After doing so, use a `Solver` instance to train a `TwoLayerNet` that achieves about 36% accuracy on the validation set.

```

[ ]: input_size = 32 * 32 * 3
hidden_size = 50
num_classes = 10
model = TwoLayerNet(input_size, hidden_size, num_classes)
solver = None

#####
# TODO: Use a Solver instance to train a TwoLayerNet that achieves about 36% #
# accuracy on the validation set.                                           #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

solver = Solver(model, data,
                 update_rule='sgd',
                 optim_config={
                     'learning_rate': 1e-4,
                 },
                 lr_decay=0.95,
                 num_epochs=5, batch_size=200,
                 print_every=100)

solver.train()

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#####
#                               END OF YOUR CODE                               #
#####

```

```

(Iteration 1 / 1225) loss: 2.303735
(Epoch 0 / 5) train acc: 0.108000; val_acc: 0.092000
(Iteration 101 / 1225) loss: 2.263315
(Iteration 201 / 1225) loss: 2.183178
(Epoch 1 / 5) train acc: 0.227000; val_acc: 0.263000
(Iteration 301 / 1225) loss: 2.073311
(Iteration 401 / 1225) loss: 2.035833
(Epoch 2 / 5) train acc: 0.312000; val_acc: 0.290000
(Iteration 501 / 1225) loss: 1.988650
(Iteration 601 / 1225) loss: 1.833697
(Iteration 701 / 1225) loss: 1.894654
(Epoch 3 / 5) train acc: 0.310000; val_acc: 0.325000
(Iteration 801 / 1225) loss: 1.815753
(Iteration 901 / 1225) loss: 1.953207
(Epoch 4 / 5) train acc: 0.337000; val_acc: 0.353000
(Iteration 1001 / 1225) loss: 1.813236
(Iteration 1101 / 1225) loss: 1.885266
(Iteration 1201 / 1225) loss: 1.909371
(Epoch 5 / 5) train acc: 0.369000; val_acc: 0.372000

```

## 10 Debug the training

With the default parameters we provided above, you should get a validation accuracy of about 0.36 on the validation set. This isn't very good.

One strategy for getting insight into what's wrong is to plot the loss function and the accuracies on the training and validation sets during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

```
[ ]: # Run this cell to visualize training loss and train / val accuracy
```


```

plt.subplot(2, 1, 1)
plt.title('Training loss')
plt.plot(solver.loss_history, 'o')
plt.xlabel('Iteration')

plt.subplot(2, 1, 2)
plt.title('Accuracy')
plt.plot(solver.train_acc_history, '-o', label='train')
plt.plot(solver.val_acc_history, '-o', label='val')
plt.plot([0.5] * len(solver.val_acc_history), 'k--')
plt.xlabel('Epoch')
plt.legend(loc='lower right')
plt.gcf().set_size_inches(15, 12)

```

```
plt.show()
```




```
[ ]: from cs231n.vis_utils import visualize_grid

# Visualize the weights of the network

def show_net_weights(net):
    W1 = net.params['W1']
    W1 = W1.reshape(3, 32, 32, -1).transpose(3, 1, 2, 0)
    plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()

show_net_weights(model)
```





## 11 Tune your hyperparameters

**What's wrong?.** Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

**Tuning.** Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer

size, learning rate, number of training epochs, and regularization strength. You might also consider tuning the learning rate decay, but you should be able to get good performance using the default value.

**Approximate results.** You should be able to achieve a classification accuracy of greater than 48% on the validation set. Our best network gets over 52% on the validation set.

**Experiment:** Your goal in this exercise is to get as good of a result on CIFAR-10 as you can (52% could serve as a reference), with a fully-connected Neural Network. Feel free to implement your own techniques (e.g. PCA to reduce dimensionality, or adding dropout, or adding features to the solver, etc.).

```
[ ]: best_model = None

#####
# TODO: Tune hyperparameters using the validation set. Store your best trained
#
# model in best_model.
#
#
# To help debug your network, it may help to use visualizations similar to the
#
# ones we used above; these visualizations will have significant qualitative
#
# differences from the ones we saw above for the poorly tuned network.
#
#
# Tweaking hyperparameters by hand can be fun, but you might find it useful to
#
# write code to sweep through possible combinations of hyperparameters
#
# automatically like we did on the previous exercises.
#
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
best_val_acc=0
best_params={}
params={}
for learning_rate in (1e-4,1e-3):
    for lr_decay in (0.95, 0.90):
        for num_epochs in (5,10):
            for batch_size in (200, 100):
                solver = Solver(model, data,
                                update_rule='sgd',
                                optim_config={
```

```

        'learning_rate': learning_rater,
    },
    lr_decay=lr_decayr,
    num_epochs=num_epochsr, batch_size=batch_sizer,
    print_every=100)
params["learning_rate"]=learning_rater
params["lr_decay"]=lr_decayr
params["num_epochs"]=num_epochsr
params["batch_size"]=batch_sizer
print(params)
solver.train()
if solver.best_val_acc > best_val_acc:
    best_val_acc=solver.best_val_acc
    best_params["learning_rate"]=learning_rater
    best_params["lr_decay"]=lr_decayr
    best_params["num_epochs"]=num_epochsr
    best_params["batch_size"]=batch_sizer

print(best_val_acc)
print(best_params)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#####
#                                     END OF YOUR CODE                                     #
#####

```

```

{'learning_rate': 0.0001, 'lr_decay': 0.95, 'num_epochs': 5, 'batch_size': 200}
(Iteration 1 / 1225) loss: 1.334293
(Epoch 0 / 5) train acc: 0.555000; val_acc: 0.514000
(Iteration 101 / 1225) loss: 1.215951
(Iteration 201 / 1225) loss: 1.274708
(Epoch 1 / 5) train acc: 0.583000; val_acc: 0.509000
(Iteration 301 / 1225) loss: 1.193970
(Iteration 401 / 1225) loss: 1.263497
(Epoch 2 / 5) train acc: 0.580000; val_acc: 0.504000
(Iteration 501 / 1225) loss: 1.535640
(Iteration 601 / 1225) loss: 1.312291
(Iteration 701 / 1225) loss: 1.124291
(Epoch 3 / 5) train acc: 0.598000; val_acc: 0.504000
(Iteration 801 / 1225) loss: 1.236232
(Iteration 901 / 1225) loss: 1.161867
(Epoch 4 / 5) train acc: 0.604000; val_acc: 0.507000
(Iteration 1001 / 1225) loss: 1.303681
(Iteration 1101 / 1225) loss: 1.185320
(Iteration 1201 / 1225) loss: 1.164725
(Epoch 5 / 5) train acc: 0.596000; val_acc: 0.504000

```

```

{'learning rate': 0.0001, 'lr_decay': 0.95, 'num_epochs': 5, 'batch_size': 100}
(Iteration 1 / 2450) loss: 1.236851
(Epoch 0 / 5) train acc: 0.536000; val_acc: 0.511000
(Iteration 101 / 2450) loss: 1.169309
(Iteration 201 / 2450) loss: 1.151894
(Iteration 301 / 2450) loss: 1.178712
(Iteration 401 / 2450) loss: 1.239360
(Epoch 1 / 5) train acc: 0.579000; val_acc: 0.494000
(Iteration 501 / 2450) loss: 1.575943
(Iteration 601 / 2450) loss: 1.287569
(Iteration 701 / 2450) loss: 1.421999
(Iteration 801 / 2450) loss: 1.299235
(Iteration 901 / 2450) loss: 1.182849
(Epoch 2 / 5) train acc: 0.565000; val_acc: 0.517000
(Iteration 1001 / 2450) loss: 1.471185
(Iteration 1101 / 2450) loss: 1.224873
(Iteration 1201 / 2450) loss: 1.050676
(Iteration 1301 / 2450) loss: 1.091264
(Iteration 1401 / 2450) loss: 1.318064
(Epoch 3 / 5) train acc: 0.616000; val_acc: 0.503000
(Iteration 1501 / 2450) loss: 1.315003
(Iteration 1601 / 2450) loss: 1.151905
(Iteration 1701 / 2450) loss: 1.243062
(Iteration 1801 / 2450) loss: 1.192332
(Iteration 1901 / 2450) loss: 1.300316
(Epoch 4 / 5) train acc: 0.586000; val_acc: 0.508000
(Iteration 2001 / 2450) loss: 1.283836
(Iteration 2101 / 2450) loss: 1.298036
(Iteration 2201 / 2450) loss: 1.253789
(Iteration 2301 / 2450) loss: 1.093579
(Iteration 2401 / 2450) loss: 1.399402
(Epoch 5 / 5) train acc: 0.589000; val_acc: 0.504000
{'learning rate': 0.0001, 'lr_decay': 0.95, 'num_epochs': 10, 'batch_size': 200}
(Iteration 1 / 2450) loss: 1.340776
(Epoch 0 / 10) train acc: 0.581000; val_acc: 0.511000
(Iteration 101 / 2450) loss: 1.152492
(Iteration 201 / 2450) loss: 1.211404
(Epoch 1 / 10) train acc: 0.603000; val_acc: 0.505000
(Iteration 301 / 2450) loss: 1.459455
(Iteration 401 / 2450) loss: 1.316210
(Epoch 2 / 10) train acc: 0.587000; val_acc: 0.508000
(Iteration 501 / 2450) loss: 1.210939
(Iteration 601 / 2450) loss: 1.192460
(Iteration 701 / 2450) loss: 1.140021
(Epoch 3 / 10) train acc: 0.568000; val_acc: 0.518000
(Iteration 801 / 2450) loss: 1.218555
(Iteration 901 / 2450) loss: 1.272382
(Epoch 4 / 10) train acc: 0.559000; val_acc: 0.508000

```

```

(Iteration 1001 / 2450) loss: 1.164954
(Iteration 1101 / 2450) loss: 1.243192
(Iteration 1201 / 2450) loss: 1.296231
(Epoch 5 / 10) train acc: 0.556000; val_acc: 0.515000
(Iteration 1301 / 2450) loss: 1.299272
(Iteration 1401 / 2450) loss: 1.084854
(Epoch 6 / 10) train acc: 0.601000; val_acc: 0.516000
(Iteration 1501 / 2450) loss: 1.220441
(Iteration 1601 / 2450) loss: 1.304873
(Iteration 1701 / 2450) loss: 1.260851
(Epoch 7 / 10) train acc: 0.594000; val_acc: 0.515000
(Iteration 1801 / 2450) loss: 1.203636
(Iteration 1901 / 2450) loss: 1.257709
(Epoch 8 / 10) train acc: 0.593000; val_acc: 0.511000
(Iteration 2001 / 2450) loss: 1.193002
(Iteration 2101 / 2450) loss: 1.272920
(Iteration 2201 / 2450) loss: 1.333933
(Epoch 9 / 10) train acc: 0.590000; val_acc: 0.505000
(Iteration 2301 / 2450) loss: 1.101427
(Iteration 2401 / 2450) loss: 1.170644
(Epoch 10 / 10) train acc: 0.599000; val_acc: 0.509000
{'learning rate': 0.0001, 'lr_decay': 0.95, 'num_epochs': 10, 'batch_size': 100}
(Iteration 1 / 4900) loss: 1.174249
(Epoch 0 / 10) train acc: 0.587000; val_acc: 0.515000
(Iteration 101 / 4900) loss: 1.159400
(Iteration 201 / 4900) loss: 1.197879
(Iteration 301 / 4900) loss: 1.225508
(Iteration 401 / 4900) loss: 1.314446
(Epoch 1 / 10) train acc: 0.589000; val_acc: 0.505000
(Iteration 501 / 4900) loss: 1.173523
(Iteration 601 / 4900) loss: 1.241319
(Iteration 701 / 4900) loss: 1.314660
(Iteration 801 / 4900) loss: 1.314505
(Iteration 901 / 4900) loss: 1.177624
(Epoch 2 / 10) train acc: 0.581000; val_acc: 0.502000
(Iteration 1001 / 4900) loss: 1.218231
(Iteration 1101 / 4900) loss: 1.264991
(Iteration 1201 / 4900) loss: 1.235472
(Iteration 1301 / 4900) loss: 1.374937
(Iteration 1401 / 4900) loss: 1.209611
(Epoch 3 / 10) train acc: 0.594000; val_acc: 0.507000
(Iteration 1501 / 4900) loss: 1.196669
(Iteration 1601 / 4900) loss: 1.225460
(Iteration 1701 / 4900) loss: 1.321993
(Iteration 1801 / 4900) loss: 1.245697
(Iteration 1901 / 4900) loss: 1.124853
(Epoch 4 / 10) train acc: 0.589000; val_acc: 0.520000
(Iteration 2001 / 4900) loss: 1.172097

```

```

(Iteration 2101 / 4900) loss: 1.302930
(Iteration 2201 / 4900) loss: 1.071661
(Iteration 2301 / 4900) loss: 1.076187
(Iteration 2401 / 4900) loss: 1.197772
(Epoch 5 / 10) train acc: 0.595000; val_acc: 0.514000
(Iteration 2501 / 4900) loss: 1.271050
(Iteration 2601 / 4900) loss: 0.966551
(Iteration 2701 / 4900) loss: 1.024883
(Iteration 2801 / 4900) loss: 1.478797
(Iteration 2901 / 4900) loss: 1.341238
(Epoch 6 / 10) train acc: 0.587000; val_acc: 0.516000
(Iteration 3001 / 4900) loss: 1.106458
(Iteration 3101 / 4900) loss: 1.071064
(Iteration 3201 / 4900) loss: 1.196401
(Iteration 3301 / 4900) loss: 1.091961
(Iteration 3401 / 4900) loss: 1.162516
(Epoch 7 / 10) train acc: 0.590000; val_acc: 0.507000
(Iteration 3501 / 4900) loss: 1.161735
(Iteration 3601 / 4900) loss: 1.151226
(Iteration 3701 / 4900) loss: 1.132763
(Iteration 3801 / 4900) loss: 1.205242
(Iteration 3901 / 4900) loss: 1.308913
(Epoch 8 / 10) train acc: 0.572000; val_acc: 0.510000
(Iteration 4001 / 4900) loss: 1.228737
(Iteration 4101 / 4900) loss: 1.128851
(Iteration 4201 / 4900) loss: 1.125424
(Iteration 4301 / 4900) loss: 1.073133
(Iteration 4401 / 4900) loss: 1.219704
(Epoch 9 / 10) train acc: 0.608000; val_acc: 0.503000
(Iteration 4501 / 4900) loss: 1.277812
(Iteration 4601 / 4900) loss: 1.106107
(Iteration 4701 / 4900) loss: 1.296202
(Iteration 4801 / 4900) loss: 1.144922
(Epoch 10 / 10) train acc: 0.590000; val_acc: 0.515000
{'learning rate': 0.0001, 'lr_decay': 0.9, 'num_epochs': 5, 'batch_size': 200}
(Iteration 1 / 1225) loss: 1.240464
(Epoch 0 / 5) train acc: 0.596000; val_acc: 0.518000
(Iteration 101 / 1225) loss: 1.203417
(Iteration 201 / 1225) loss: 1.222274
(Epoch 1 / 5) train acc: 0.597000; val_acc: 0.517000
(Iteration 301 / 1225) loss: 1.071790
(Iteration 401 / 1225) loss: 1.253716
(Epoch 2 / 5) train acc: 0.587000; val_acc: 0.511000
(Iteration 501 / 1225) loss: 1.133081
(Iteration 601 / 1225) loss: 1.226793
(Iteration 701 / 1225) loss: 1.236190
(Epoch 3 / 5) train acc: 0.595000; val_acc: 0.519000
(Iteration 801 / 1225) loss: 1.250712

```

```

(Iteration 901 / 1225) loss: 1.188699
(Epoch 4 / 5) train acc: 0.590000; val_acc: 0.519000
(Iteration 1001 / 1225) loss: 1.292413
(Iteration 1101 / 1225) loss: 1.357767
(Iteration 1201 / 1225) loss: 1.175122
(Epoch 5 / 5) train acc: 0.583000; val_acc: 0.517000
{'learning rate': 0.0001, 'lr_decay': 0.9, 'num_epochs': 5, 'batch_size': 100}
(Iteration 1 / 2450) loss: 1.246927
(Epoch 0 / 5) train acc: 0.576000; val_acc: 0.520000
(Iteration 101 / 2450) loss: 1.238294
(Iteration 201 / 2450) loss: 1.299699
(Iteration 301 / 2450) loss: 1.057388
(Iteration 401 / 2450) loss: 1.065496
(Epoch 1 / 5) train acc: 0.563000; val_acc: 0.516000
(Iteration 501 / 2450) loss: 1.066503
(Iteration 601 / 2450) loss: 1.373125
(Iteration 701 / 2450) loss: 1.265260
(Iteration 801 / 2450) loss: 1.181214
(Iteration 901 / 2450) loss: 1.173887
(Epoch 2 / 5) train acc: 0.595000; val_acc: 0.509000
(Iteration 1001 / 2450) loss: 1.045874
(Iteration 1101 / 2450) loss: 1.404552
(Iteration 1201 / 2450) loss: 1.179401
(Iteration 1301 / 2450) loss: 1.183274
(Iteration 1401 / 2450) loss: 1.033464
(Epoch 3 / 5) train acc: 0.596000; val_acc: 0.499000
(Iteration 1501 / 2450) loss: 1.254620
(Iteration 1601 / 2450) loss: 1.297570
(Iteration 1701 / 2450) loss: 1.227940
(Iteration 1801 / 2450) loss: 1.253173
(Iteration 1901 / 2450) loss: 1.069068
(Epoch 4 / 5) train acc: 0.621000; val_acc: 0.511000
(Iteration 2001 / 2450) loss: 1.127980
(Iteration 2101 / 2450) loss: 1.544804
(Iteration 2201 / 2450) loss: 1.195122
(Iteration 2301 / 2450) loss: 1.052843
(Iteration 2401 / 2450) loss: 1.264248
(Epoch 5 / 5) train acc: 0.586000; val_acc: 0.509000
{'learning rate': 0.0001, 'lr_decay': 0.9, 'num_epochs': 10, 'batch_size': 200}
(Iteration 1 / 2450) loss: 1.257871
(Epoch 0 / 10) train acc: 0.598000; val_acc: 0.518000
(Iteration 101 / 2450) loss: 1.209629
(Iteration 201 / 2450) loss: 1.183085
(Epoch 1 / 10) train acc: 0.588000; val_acc: 0.525000
(Iteration 301 / 2450) loss: 1.246515
(Iteration 401 / 2450) loss: 1.161239
(Epoch 2 / 10) train acc: 0.619000; val_acc: 0.512000
(Iteration 501 / 2450) loss: 1.234444

```

```

(Iteration 601 / 2450) loss: 1.186589
(Iteration 701 / 2450) loss: 1.243026
(Epoch 3 / 10) train acc: 0.610000; val_acc: 0.523000
(Iteration 801 / 2450) loss: 1.232217
(Iteration 901 / 2450) loss: 0.972824
(Epoch 4 / 10) train acc: 0.576000; val_acc: 0.524000
(Iteration 1001 / 2450) loss: 1.214154
(Iteration 1101 / 2450) loss: 1.162776
(Iteration 1201 / 2450) loss: 1.207865
(Epoch 5 / 10) train acc: 0.608000; val_acc: 0.518000
(Iteration 1301 / 2450) loss: 1.152043
(Iteration 1401 / 2450) loss: 1.101142
(Epoch 6 / 10) train acc: 0.592000; val_acc: 0.510000
(Iteration 1501 / 2450) loss: 1.210411
(Iteration 1601 / 2450) loss: 1.161065
(Iteration 1701 / 2450) loss: 1.146600
(Epoch 7 / 10) train acc: 0.644000; val_acc: 0.511000
(Iteration 1801 / 2450) loss: 1.168704
(Iteration 1901 / 2450) loss: 1.146179
(Epoch 8 / 10) train acc: 0.590000; val_acc: 0.510000
(Iteration 2001 / 2450) loss: 1.222747
(Iteration 2101 / 2450) loss: 1.209122
(Iteration 2201 / 2450) loss: 1.166366
(Epoch 9 / 10) train acc: 0.655000; val_acc: 0.516000
(Iteration 2301 / 2450) loss: 1.038828
(Iteration 2401 / 2450) loss: 1.103141
(Epoch 10 / 10) train acc: 0.590000; val_acc: 0.511000
{'learning rate': 0.0001, 'lr_decay': 0.9, 'num_epochs': 10, 'batch_size': 100}
(Iteration 1 / 4900) loss: 1.148802
(Epoch 0 / 10) train acc: 0.589000; val_acc: 0.524000
(Iteration 101 / 4900) loss: 1.086203
(Iteration 201 / 4900) loss: 1.096573
(Iteration 301 / 4900) loss: 1.206868
(Iteration 401 / 4900) loss: 1.096110
(Epoch 1 / 10) train acc: 0.554000; val_acc: 0.506000
(Iteration 501 / 4900) loss: 1.069904
(Iteration 601 / 4900) loss: 1.245022
(Iteration 701 / 4900) loss: 1.231844
(Iteration 801 / 4900) loss: 1.027262
(Iteration 901 / 4900) loss: 1.057421
(Epoch 2 / 10) train acc: 0.592000; val_acc: 0.506000
(Iteration 1001 / 4900) loss: 1.080580
(Iteration 1101 / 4900) loss: 1.190796
(Iteration 1201 / 4900) loss: 1.099369
(Iteration 1301 / 4900) loss: 1.332302
(Iteration 1401 / 4900) loss: 1.303097
(Epoch 3 / 10) train acc: 0.592000; val_acc: 0.513000
(Iteration 1501 / 4900) loss: 1.274668

```



```

(Iteration 1601 / 4900) loss: 1.206060
(Iteration 1701 / 4900) loss: 1.153838
(Iteration 1801 / 4900) loss: 1.247895
(Iteration 1901 / 4900) loss: 1.309677
(Epoch 4 / 10) train acc: 0.593000; val_acc: 0.497000
(Iteration 2001 / 4900) loss: 1.031616
(Iteration 2101 / 4900) loss: 1.237734
(Iteration 2201 / 4900) loss: 1.150693
(Iteration 2301 / 4900) loss: 1.149603
(Iteration 2401 / 4900) loss: 1.392014
(Epoch 5 / 10) train acc: 0.611000; val_acc: 0.498000
(Iteration 2501 / 4900) loss: 1.160109
(Iteration 2601 / 4900) loss: 1.269934
(Iteration 2701 / 4900) loss: 1.007172
(Iteration 2801 / 4900) loss: 1.198622
(Iteration 2901 / 4900) loss: 1.353890
(Epoch 6 / 10) train acc: 0.622000; val_acc: 0.518000
(Iteration 3001 / 4900) loss: 1.175625
(Iteration 3101 / 4900) loss: 1.302547
(Iteration 3201 / 4900) loss: 1.167640
(Iteration 3301 / 4900) loss: 1.032510
(Iteration 3401 / 4900) loss: 1.229017
(Epoch 7 / 10) train acc: 0.597000; val_acc: 0.515000
(Iteration 3501 / 4900) loss: 1.176262
(Iteration 3601 / 4900) loss: 1.175052
(Iteration 3701 / 4900) loss: 1.018686
(Iteration 3801 / 4900) loss: 1.252146
(Iteration 3901 / 4900) loss: 1.037884
(Epoch 8 / 10) train acc: 0.609000; val_acc: 0.502000
(Iteration 4001 / 4900) loss: 1.152067
(Iteration 4101 / 4900) loss: 1.073804
(Iteration 4201 / 4900) loss: 1.081678
(Iteration 4301 / 4900) loss: 1.184315
(Iteration 4401 / 4900) loss: 1.126315
(Epoch 9 / 10) train acc: 0.615000; val_acc: 0.513000
(Iteration 4501 / 4900) loss: 1.176094
(Iteration 4601 / 4900) loss: 1.131342
(Iteration 4701 / 4900) loss: 1.170025
(Iteration 4801 / 4900) loss: 0.932521
(Epoch 10 / 10) train acc: 0.637000; val_acc: 0.501000
{'learning rate': 0.001, 'lr_decay': 0.95, 'num_epochs': 5, 'batch_size': 200}
(Iteration 1 / 1225) loss: 1.102216
(Epoch 0 / 5) train acc: 0.552000; val_acc: 0.503000
(Iteration 101 / 1225) loss: 1.263104
(Iteration 201 / 1225) loss: 1.283665
(Epoch 1 / 5) train acc: 0.546000; val_acc: 0.483000
(Iteration 301 / 1225) loss: 1.318610
(Iteration 401 / 1225) loss: 1.337281

```

```

(Epoch 2 / 5) train acc: 0.534000; val_acc: 0.465000
(Iteration 501 / 1225) loss: 1.389820
(Iteration 601 / 1225) loss: 1.294946
(Iteration 701 / 1225) loss: 1.290945
(Epoch 3 / 5) train acc: 0.546000; val_acc: 0.479000
(Iteration 801 / 1225) loss: 1.243757
(Iteration 901 / 1225) loss: 1.237142
(Epoch 4 / 5) train acc: 0.551000; val_acc: 0.478000
(Iteration 1001 / 1225) loss: 1.303622
(Iteration 1101 / 1225) loss: 1.334273
(Iteration 1201 / 1225) loss: 1.247191
(Epoch 5 / 5) train acc: 0.574000; val_acc: 0.514000
{'learning rate': 0.001, 'lr_decay': 0.95, 'num_epochs': 5, 'batch_size': 100}
(Iteration 1 / 2450) loss: 1.192671
(Epoch 0 / 5) train acc: 0.558000; val_acc: 0.469000
(Iteration 101 / 2450) loss: 1.563985
(Iteration 201 / 2450) loss: 1.424677
(Iteration 301 / 2450) loss: 1.182432
(Iteration 401 / 2450) loss: 1.100741
(Epoch 1 / 5) train acc: 0.531000; val_acc: 0.485000
(Iteration 501 / 2450) loss: 1.421617
(Iteration 601 / 2450) loss: 1.410256
(Iteration 701 / 2450) loss: 1.270836
(Iteration 801 / 2450) loss: 1.507237
(Iteration 901 / 2450) loss: 1.181630
(Epoch 2 / 5) train acc: 0.551000; val_acc: 0.492000
(Iteration 1001 / 2450) loss: 1.137249
(Iteration 1101 / 2450) loss: 1.246342
(Iteration 1201 / 2450) loss: 1.337724
(Iteration 1301 / 2450) loss: 1.256458
(Iteration 1401 / 2450) loss: 1.275648
(Epoch 3 / 5) train acc: 0.551000; val_acc: 0.478000
(Iteration 1501 / 2450) loss: 1.165692
(Iteration 1601 / 2450) loss: 1.320590
(Iteration 1701 / 2450) loss: 1.415335
(Iteration 1801 / 2450) loss: 1.178841
(Iteration 1901 / 2450) loss: 1.183951
(Epoch 4 / 5) train acc: 0.581000; val_acc: 0.503000
(Iteration 2001 / 2450) loss: 1.516323
(Iteration 2101 / 2450) loss: 1.192316
(Iteration 2201 / 2450) loss: 1.253786
(Iteration 2301 / 2450) loss: 1.077134
(Iteration 2401 / 2450) loss: 1.111350
(Epoch 5 / 5) train acc: 0.572000; val_acc: 0.498000
{'learning rate': 0.001, 'lr_decay': 0.95, 'num_epochs': 10, 'batch_size': 200}
(Iteration 1 / 2450) loss: 1.211421
(Epoch 0 / 10) train acc: 0.570000; val_acc: 0.489000
(Iteration 101 / 2450) loss: 1.245639

```

```

(Iteration 201 / 2450) loss: 1.174411
(Epoch 1 / 10) train acc: 0.585000; val_acc: 0.482000
(Iteration 301 / 2450) loss: 1.165967
(Iteration 401 / 2450) loss: 1.169161
(Epoch 2 / 10) train acc: 0.602000; val_acc: 0.499000
(Iteration 501 / 2450) loss: 1.149993
(Iteration 601 / 2450) loss: 1.224108
(Iteration 701 / 2450) loss: 1.115634
(Epoch 3 / 10) train acc: 0.595000; val_acc: 0.490000
(Iteration 801 / 2450) loss: 1.086061
(Iteration 901 / 2450) loss: 1.150472
(Epoch 4 / 10) train acc: 0.628000; val_acc: 0.492000
(Iteration 1001 / 2450) loss: 1.151101
(Iteration 1101 / 2450) loss: 1.162405
(Iteration 1201 / 2450) loss: 0.971081
(Epoch 5 / 10) train acc: 0.629000; val_acc: 0.504000
(Iteration 1301 / 2450) loss: 1.037766
(Iteration 1401 / 2450) loss: 1.063797
(Epoch 6 / 10) train acc: 0.616000; val_acc: 0.512000
(Iteration 1501 / 2450) loss: 1.220252
(Iteration 1601 / 2450) loss: 1.093340
(Iteration 1701 / 2450) loss: 1.112789
(Epoch 7 / 10) train acc: 0.642000; val_acc: 0.512000
(Iteration 1801 / 2450) loss: 1.210731
(Iteration 1901 / 2450) loss: 0.991425
(Epoch 8 / 10) train acc: 0.606000; val_acc: 0.494000
(Iteration 2001 / 2450) loss: 1.081700
(Iteration 2101 / 2450) loss: 1.098585
(Iteration 2201 / 2450) loss: 0.894657
(Epoch 9 / 10) train acc: 0.638000; val_acc: 0.494000
(Iteration 2301 / 2450) loss: 1.107582
(Iteration 2401 / 2450) loss: 0.945513
(Epoch 10 / 10) train acc: 0.624000; val_acc: 0.506000
{'learning rate': 0.001, 'lr_decay': 0.95, 'num_epochs': 10, 'batch_size': 100}
(Iteration 1 / 4900) loss: 1.045242
(Epoch 0 / 10) train acc: 0.610000; val_acc: 0.498000
(Iteration 101 / 4900) loss: 1.242157
(Iteration 201 / 4900) loss: 1.240966
(Iteration 301 / 4900) loss: 1.146744
(Iteration 401 / 4900) loss: 1.624467
(Epoch 1 / 10) train acc: 0.538000; val_acc: 0.501000
(Iteration 501 / 4900) loss: 1.232376
(Iteration 601 / 4900) loss: 1.174558
(Iteration 701 / 4900) loss: 1.232762
(Iteration 801 / 4900) loss: 1.438685
(Iteration 901 / 4900) loss: 1.199499
(Epoch 2 / 10) train acc: 0.578000; val_acc: 0.494000
(Iteration 1001 / 4900) loss: 1.123911

```

```

(Iteration 1101 / 4900) loss: 1.130369
(Iteration 1201 / 4900) loss: 1.112972
(Iteration 1301 / 4900) loss: 1.110786
(Iteration 1401 / 4900) loss: 1.070576
(Epoch 3 / 10) train acc: 0.568000; val_acc: 0.476000
(Iteration 1501 / 4900) loss: 1.092014
(Iteration 1601 / 4900) loss: 1.075848
(Iteration 1701 / 4900) loss: 1.137584
(Iteration 1801 / 4900) loss: 1.294126
(Iteration 1901 / 4900) loss: 1.323767
(Epoch 4 / 10) train acc: 0.592000; val_acc: 0.495000
(Iteration 2001 / 4900) loss: 1.206453
(Iteration 2101 / 4900) loss: 1.141266
(Iteration 2201 / 4900) loss: 0.943430
(Iteration 2301 / 4900) loss: 1.157373
(Iteration 2401 / 4900) loss: 1.114343
(Epoch 5 / 10) train acc: 0.614000; val_acc: 0.494000
(Iteration 2501 / 4900) loss: 1.461067
(Iteration 2601 / 4900) loss: 0.961333
(Iteration 2701 / 4900) loss: 1.054090
(Iteration 2801 / 4900) loss: 1.131885
(Iteration 2901 / 4900) loss: 1.169057
(Epoch 6 / 10) train acc: 0.597000; val_acc: 0.472000
(Iteration 3001 / 4900) loss: 1.315391
(Iteration 3101 / 4900) loss: 1.064121
(Iteration 3201 / 4900) loss: 1.272993
(Iteration 3301 / 4900) loss: 1.115863
(Iteration 3401 / 4900) loss: 0.991262
(Epoch 7 / 10) train acc: 0.627000; val_acc: 0.489000
(Iteration 3501 / 4900) loss: 1.048818
(Iteration 3601 / 4900) loss: 1.198624
(Iteration 3701 / 4900) loss: 1.036087
(Iteration 3801 / 4900) loss: 1.108384
(Iteration 3901 / 4900) loss: 1.103451
(Epoch 8 / 10) train acc: 0.606000; val_acc: 0.473000
(Iteration 4001 / 4900) loss: 1.016211
(Iteration 4101 / 4900) loss: 1.204607
(Iteration 4201 / 4900) loss: 0.881028
(Iteration 4301 / 4900) loss: 1.057467
(Iteration 4401 / 4900) loss: 1.051947
(Epoch 9 / 10) train acc: 0.609000; val_acc: 0.493000
(Iteration 4501 / 4900) loss: 1.093004
(Iteration 4601 / 4900) loss: 0.927405
(Iteration 4701 / 4900) loss: 1.028503
(Iteration 4801 / 4900) loss: 1.260495
(Epoch 10 / 10) train acc: 0.605000; val_acc: 0.510000
{'learning rate': 0.001, 'lr_decay': 0.9, 'num_epochs': 5, 'batch_size': 200}
(Iteration 1 / 1225) loss: 1.257363

```

```

(Epoch 0 / 5) train acc: 0.611000; val_acc: 0.462000
(Iteration 101 / 1225) loss: 0.994359
(Iteration 201 / 1225) loss: 1.014956
(Epoch 1 / 5) train acc: 0.661000; val_acc: 0.499000
(Iteration 301 / 1225) loss: 1.023271
(Iteration 401 / 1225) loss: 0.946393
(Epoch 2 / 5) train acc: 0.632000; val_acc: 0.491000
(Iteration 501 / 1225) loss: 1.121263
(Iteration 601 / 1225) loss: 1.009656
(Iteration 701 / 1225) loss: 1.129153
(Epoch 3 / 5) train acc: 0.641000; val_acc: 0.499000
(Iteration 801 / 1225) loss: 1.047171
(Iteration 901 / 1225) loss: 0.882001
(Epoch 4 / 5) train acc: 0.642000; val_acc: 0.488000
(Iteration 1001 / 1225) loss: 1.020817
(Iteration 1101 / 1225) loss: 0.925783
(Iteration 1201 / 1225) loss: 0.890988
(Epoch 5 / 5) train acc: 0.652000; val_acc: 0.485000
{'learning_rate': 0.001, 'lr_decay': 0.9, 'num_epochs': 5, 'batch_size': 100}
(Iteration 1 / 2450) loss: 0.985955
(Epoch 0 / 5) train acc: 0.621000; val_acc: 0.459000
(Iteration 101 / 2450) loss: 1.141552
(Iteration 201 / 2450) loss: 1.198852
(Iteration 301 / 2450) loss: 0.870221
(Iteration 401 / 2450) loss: 1.166746
(Epoch 1 / 5) train acc: 0.628000; val_acc: 0.463000
(Iteration 501 / 2450) loss: 1.161744
(Iteration 601 / 2450) loss: 1.078657
(Iteration 701 / 2450) loss: 1.200801
(Iteration 801 / 2450) loss: 1.074184
(Iteration 901 / 2450) loss: 0.935765
(Epoch 2 / 5) train acc: 0.591000; val_acc: 0.475000
(Iteration 1001 / 2450) loss: 1.093714
(Iteration 1101 / 2450) loss: 1.233433
(Iteration 1201 / 2450) loss: 1.174666
(Iteration 1301 / 2450) loss: 1.071223
(Iteration 1401 / 2450) loss: 0.817991
(Epoch 3 / 5) train acc: 0.622000; val_acc: 0.488000
(Iteration 1501 / 2450) loss: 1.068485
(Iteration 1601 / 2450) loss: 1.228715
(Iteration 1701 / 2450) loss: 0.964206
(Iteration 1801 / 2450) loss: 1.044713
(Iteration 1901 / 2450) loss: 0.898048
(Epoch 4 / 5) train acc: 0.626000; val_acc: 0.490000
(Iteration 2001 / 2450) loss: 1.149295
(Iteration 2101 / 2450) loss: 1.114429
(Iteration 2201 / 2450) loss: 1.160378
(Iteration 2301 / 2450) loss: 0.958404

```

```

(Iteration 2401 / 2450) loss: 1.116809
(Epoch 5 / 5) train acc: 0.662000; val_acc: 0.488000
{'learning rate': 0.001, 'lr_decay': 0.9, 'num_epochs': 10, 'batch_size': 200}
(Iteration 1 / 2450) loss: 1.090109
(Epoch 0 / 10) train acc: 0.651000; val_acc: 0.466000
(Iteration 101 / 2450) loss: 0.957648
(Iteration 201 / 2450) loss: 0.992259
(Epoch 1 / 10) train acc: 0.636000; val_acc: 0.482000
(Iteration 301 / 2450) loss: 0.969113
(Iteration 401 / 2450) loss: 0.996230
(Epoch 2 / 10) train acc: 0.634000; val_acc: 0.498000
(Iteration 501 / 2450) loss: 1.009661
(Iteration 601 / 2450) loss: 1.078417
(Iteration 701 / 2450) loss: 0.970738
(Epoch 3 / 10) train acc: 0.625000; val_acc: 0.488000
(Iteration 801 / 2450) loss: 0.831877
(Iteration 901 / 2450) loss: 1.009152
(Epoch 4 / 10) train acc: 0.659000; val_acc: 0.496000
(Iteration 1001 / 2450) loss: 0.837540
(Iteration 1101 / 2450) loss: 1.015009
(Iteration 1201 / 2450) loss: 1.038237
(Epoch 5 / 10) train acc: 0.651000; val_acc: 0.498000
(Iteration 1301 / 2450) loss: 0.952840
(Iteration 1401 / 2450) loss: 0.927106
(Epoch 6 / 10) train acc: 0.680000; val_acc: 0.498000
(Iteration 1501 / 2450) loss: 0.809438
(Iteration 1601 / 2450) loss: 0.970926
(Iteration 1701 / 2450) loss: 0.934069
(Epoch 7 / 10) train acc: 0.682000; val_acc: 0.496000
(Iteration 1801 / 2450) loss: 1.017739
(Iteration 1901 / 2450) loss: 0.887228
(Epoch 8 / 10) train acc: 0.686000; val_acc: 0.494000
(Iteration 2001 / 2450) loss: 0.867831
(Iteration 2101 / 2450) loss: 0.896811
(Iteration 2201 / 2450) loss: 0.869605
(Epoch 9 / 10) train acc: 0.688000; val_acc: 0.493000
(Iteration 2301 / 2450) loss: 0.912281
(Iteration 2401 / 2450) loss: 0.911624
(Epoch 10 / 10) train acc: 0.707000; val_acc: 0.510000
{'learning rate': 0.001, 'lr_decay': 0.9, 'num_epochs': 10, 'batch_size': 100}
(Iteration 1 / 4900) loss: 0.829227
(Epoch 0 / 10) train acc: 0.623000; val_acc: 0.451000
(Iteration 101 / 4900) loss: 1.037059
(Iteration 201 / 4900) loss: 1.383311
(Iteration 301 / 4900) loss: 1.356798
(Iteration 401 / 4900) loss: 1.100087
(Epoch 1 / 10) train acc: 0.566000; val_acc: 0.443000
(Iteration 501 / 4900) loss: 1.083080

```

(Iteration 601 / 4900) loss: 1.157631  
(Iteration 701 / 4900) loss: 1.204532  
(Iteration 801 / 4900) loss: 1.035617  
(Iteration 901 / 4900) loss: 1.093349  
(Epoch 2 / 10) train acc: 0.626000; val\_acc: 0.466000  
(Iteration 1001 / 4900) loss: 0.857518  
(Iteration 1101 / 4900) loss: 0.872058  
(Iteration 1201 / 4900) loss: 1.071382  
(Iteration 1301 / 4900) loss: 1.037549  
(Iteration 1401 / 4900) loss: 0.863025  
(Epoch 3 / 10) train acc: 0.622000; val\_acc: 0.476000  
(Iteration 1501 / 4900) loss: 0.996405  
(Iteration 1601 / 4900) loss: 1.012951  
(Iteration 1701 / 4900) loss: 1.036150  
(Iteration 1801 / 4900) loss: 0.781415  
(Iteration 1901 / 4900) loss: 1.093766  
(Epoch 4 / 10) train acc: 0.645000; val\_acc: 0.492000  
(Iteration 2001 / 4900) loss: 1.115827  
(Iteration 2101 / 4900) loss: 0.930011  
(Iteration 2201 / 4900) loss: 0.989916  
(Iteration 2301 / 4900) loss: 0.976613  
(Iteration 2401 / 4900) loss: 1.135477  
(Epoch 5 / 10) train acc: 0.641000; val\_acc: 0.471000  
(Iteration 2501 / 4900) loss: 0.879866  
(Iteration 2601 / 4900) loss: 0.914753  
(Iteration 2701 / 4900) loss: 0.934667  
(Iteration 2801 / 4900) loss: 1.004229  
(Iteration 2901 / 4900) loss: 0.985295  
(Epoch 6 / 10) train acc: 0.666000; val\_acc: 0.471000  
(Iteration 3001 / 4900) loss: 1.123575  
(Iteration 3101 / 4900) loss: 0.902014  
(Iteration 3201 / 4900) loss: 0.915457  
(Iteration 3301 / 4900) loss: 0.807424  
(Iteration 3401 / 4900) loss: 1.048530  
(Epoch 7 / 10) train acc: 0.672000; val\_acc: 0.476000  
(Iteration 3501 / 4900) loss: 1.058749  
(Iteration 3601 / 4900) loss: 1.148586  
(Iteration 3701 / 4900) loss: 0.774565  
(Iteration 3801 / 4900) loss: 0.807034  
(Iteration 3901 / 4900) loss: 1.120696  
(Epoch 8 / 10) train acc: 0.692000; val\_acc: 0.484000  
(Iteration 4001 / 4900) loss: 0.966395  
(Iteration 4101 / 4900) loss: 0.956864  
(Iteration 4201 / 4900) loss: 0.881146  
(Iteration 4301 / 4900) loss: 0.986604  
(Iteration 4401 / 4900) loss: 0.803794  
(Epoch 9 / 10) train acc: 0.694000; val\_acc: 0.477000  
(Iteration 4501 / 4900) loss: 0.783058

```
(Iteration 4601 / 4900) loss: 0.819460
(Iteration 4701 / 4900) loss: 0.927896
(Iteration 4801 / 4900) loss: 1.125727
(Epoch 10 / 10) train acc: 0.683000; val_acc: 0.468000
0.525
{'learning_rate': 0.0001, 'lr_decay': 0.9, 'num_epochs': 10, 'batch_size': 200}
```

```
[ ]: solver = Solver(model, data,
                    update_rule='sgd',
                    optim_config={
                        'learning_rate': 1e-4,
                    },
                    lr_decay=0.9,
                    num_epochs=10, batch_size=200,
                    print_every=100)

solver.train()
best_model=model
```

```
(Iteration 1 / 2450) loss: 0.910035
(Epoch 0 / 10) train acc: 0.685000; val_acc: 0.510000
(Iteration 101 / 2450) loss: 0.930505
(Iteration 201 / 2450) loss: 0.877415
(Epoch 1 / 10) train acc: 0.705000; val_acc: 0.505000
(Iteration 301 / 2450) loss: 0.905470
(Iteration 401 / 2450) loss: 0.881467
(Epoch 2 / 10) train acc: 0.694000; val_acc: 0.502000
(Iteration 501 / 2450) loss: 0.818054
(Iteration 601 / 2450) loss: 0.846530
(Iteration 701 / 2450) loss: 0.905243
(Epoch 3 / 10) train acc: 0.695000; val_acc: 0.496000
(Iteration 801 / 2450) loss: 0.993924
(Iteration 901 / 2450) loss: 0.980080
(Epoch 4 / 10) train acc: 0.698000; val_acc: 0.497000
(Iteration 1001 / 2450) loss: 0.917439
(Iteration 1101 / 2450) loss: 0.959833
(Iteration 1201 / 2450) loss: 0.923533
(Epoch 5 / 10) train acc: 0.724000; val_acc: 0.496000
(Iteration 1301 / 2450) loss: 1.030214
(Iteration 1401 / 2450) loss: 0.967664
(Epoch 6 / 10) train acc: 0.686000; val_acc: 0.499000
(Iteration 1501 / 2450) loss: 0.824528
(Iteration 1601 / 2450) loss: 0.878561
(Iteration 1701 / 2450) loss: 0.894332
(Epoch 7 / 10) train acc: 0.699000; val_acc: 0.494000
(Iteration 1801 / 2450) loss: 0.955539
(Iteration 1901 / 2450) loss: 0.856374
(Epoch 8 / 10) train acc: 0.697000; val_acc: 0.498000
(Iteration 2001 / 2450) loss: 0.891288
```




```
(Iteration 2101 / 2450) loss: 0.765385
(Iteration 2201 / 2450) loss: 0.830216
(Epoch 9 / 10) train acc: 0.716000; val_acc: 0.498000
(Iteration 2301 / 2450) loss: 0.912148
(Iteration 2401 / 2450) loss: 0.736636
(Epoch 10 / 10) train acc: 0.714000; val_acc: 0.497000
```

```
[ ]: from cs231n.vis_utils import visualize_grid

# Visualize the weights of the network

def show_net_weights(net):
    W1 = net.params['W1']
    W1 = W1.reshape(3, 32, 32, -1).transpose(3, 1, 2, 0)
    plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()

show_net_weights(model)
```



## 12 Test your model!

Run your best model on the validation and test sets. You should achieve above 48% accuracy on the validation set and the test set.

```
[ ]: y_val_pred = np.argmax(best_model.loss(data['X_val']), axis=1)
      print('Validation set accuracy: ', (y_val_pred == data['y_val']).mean())
```

Validation set accuracy: 0.51

```
[ ]: y_test_pred = np.argmax(best_model.loss(data['X_test']), axis=1)
      print('Test set accuracy: ', (y_test_pred == data['y_test']).mean())
```

Test set accuracy: 0.505

# features

July 25, 2022

```
[2]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'cs231n/assignments/assignment1/'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the CIFAR-10 dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.  
/content/drive/My Drive/cs231n/assignments/assignment1/cs231n/datasets  
/content/drive/My Drive/cs231n/assignments/assignment1

## 1 Image features exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

We have seen that we can achieve reasonable performance on an image classification task by training a linear classifier on the pixels of the input image. In this exercise we will show that we can improve our classification performance by training linear classifiers not on raw pixels but on features that are computed from the raw pixels.

All of your work for this exercise will be done in this notebook.

```
[4]: import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# → autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

## 1.1 Load data

Similar to previous exercises, we will load CIFAR-10 data from disk.

```
[5]: from cs231n.features import color_histogram_hsv, hog_feature

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

    # Cleaning up variables to prevent loading data multiple times (which may
    # → cause memory issue)
    try:
        del X_train, y_train
        del X_test, y_test
        print('Clear previously loaded data.')
    except:
        pass

    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
```

```

mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]

return X_train, y_train, X_val, y_val, X_test, y_test

X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()

```

## 1.2 Extract Features

For each image we will compute a Histogram of Oriented Gradients (HOG) as well as a color histogram using the hue channel in HSV color space. We form our final feature vector for each image by concatenating the HOG and color histogram feature vectors.

Roughly speaking, HOG should capture the texture of the image while ignoring color information, and the color histogram represents the color of the input image while ignoring texture. As a result, we expect that using both together ought to work better than using either alone. Verifying this assumption would be a good thing to try for your own interest.

The `hog_feature` and `color_histogram_hsv` functions both operate on a single image and return a feature vector for that image. The `extract_features` function takes a set of images and a list of feature functions and evaluates each feature function on each image, storing the results in a matrix where each column is the concatenation of all feature vectors for a single image.

```

[6]: from cs231n.features import *

num_color_bins = 10 # Number of bins in the color histogram
feature_fns = [hog_feature, lambda img: color_histogram_hsv(img,
    ↪nbin=num_color_bins)]
X_train_feats = extract_features(X_train, feature_fns, verbose=True)
X_val_feats = extract_features(X_val, feature_fns)
X_test_feats = extract_features(X_test, feature_fns)

# Preprocessing: Subtract the mean feature
mean_feat = np.mean(X_train_feats, axis=0, keepdims=True)
X_train_feats -= mean_feat
X_val_feats -= mean_feat
X_test_feats -= mean_feat

# Preprocessing: Divide by standard deviation. This ensures that each feature
# has roughly the same scale.
std_feat = np.std(X_train_feats, axis=0, keepdims=True)
X_train_feats /= std_feat
X_val_feats /= std_feat

```

```
X_test_feats /= std_feat
```

```
# Preprocessing: Add a bias dimension
```

```
X_train_feats = np.hstack([X_train_feats, np.ones((X_train_feats.shape[0], 1))])
```

```
X_val_feats = np.hstack([X_val_feats, np.ones((X_val_feats.shape[0], 1))])
```

```
X_test_feats = np.hstack([X_test_feats, np.ones((X_test_feats.shape[0], 1))])
```

```
Done extracting features for 1000 / 49000 images
Done extracting features for 2000 / 49000 images
Done extracting features for 3000 / 49000 images
Done extracting features for 4000 / 49000 images
Done extracting features for 5000 / 49000 images
Done extracting features for 6000 / 49000 images
Done extracting features for 7000 / 49000 images
Done extracting features for 8000 / 49000 images
Done extracting features for 9000 / 49000 images
Done extracting features for 10000 / 49000 images
Done extracting features for 11000 / 49000 images
Done extracting features for 12000 / 49000 images
Done extracting features for 13000 / 49000 images
Done extracting features for 14000 / 49000 images
Done extracting features for 15000 / 49000 images
Done extracting features for 16000 / 49000 images
Done extracting features for 17000 / 49000 images
Done extracting features for 18000 / 49000 images
Done extracting features for 19000 / 49000 images
Done extracting features for 20000 / 49000 images
Done extracting features for 21000 / 49000 images
Done extracting features for 22000 / 49000 images
Done extracting features for 23000 / 49000 images
Done extracting features for 24000 / 49000 images
Done extracting features for 25000 / 49000 images
Done extracting features for 26000 / 49000 images
Done extracting features for 27000 / 49000 images
Done extracting features for 28000 / 49000 images
Done extracting features for 29000 / 49000 images
Done extracting features for 30000 / 49000 images
Done extracting features for 31000 / 49000 images
Done extracting features for 32000 / 49000 images
Done extracting features for 33000 / 49000 images
Done extracting features for 34000 / 49000 images
Done extracting features for 35000 / 49000 images
Done extracting features for 36000 / 49000 images
Done extracting features for 37000 / 49000 images
Done extracting features for 38000 / 49000 images
Done extracting features for 39000 / 49000 images
Done extracting features for 40000 / 49000 images
```

```

Done extracting features for 41000 / 49000 images
Done extracting features for 42000 / 49000 images
Done extracting features for 43000 / 49000 images
Done extracting features for 44000 / 49000 images
Done extracting features for 45000 / 49000 images
Done extracting features for 46000 / 49000 images
Done extracting features for 47000 / 49000 images
Done extracting features for 48000 / 49000 images
Done extracting features for 49000 / 49000 images

```

### 1.3 Train SVM on features

Using the multiclass SVM code developed earlier in the assignment, train SVMs on top of the features extracted above; this should achieve better results than training SVMs directly on top of raw pixels.

```

[7]: # Use the validation set to tune the learning rate and regularization strength

from cs231n.classifiers.linear_classifier import LinearSVM

learning_rates = [1e-9, 1e-8, 1e-7]
regularization_strengths = [5e4, 5e5, 5e6]

results = {}
best_val = -1
best_svm = None

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength. #
# This should be identical to the validation that you did for the SVM; save #
# the best trained classifier in best_svm. You might also want to play #
# with different numbers of bins in the color histogram. If you are careful #
# you should be able to get accuracy of near 0.44 on the validation set. #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

for lr in learning_rates:
    for rs in regularization_strengths:
        svm = LinearSVM()
        svm.train(X_train_feats, y_train, learning_rate=lr, reg=rs,
                  num_iters=1500, verbose=True)
        y_train_pred = svm.predict(X_train_feats)
        train_accuracy= np.mean(y_train == y_train_pred)
        y_val_pred = svm.predict(X_val_feats)
        validation_accuracy= np.mean(y_val == y_val_pred)
        results[(lr, rs)] = (train_accuracy, validation_accuracy)

```



```

    if validation_accuracy>best_val:
        best_val=validation_accuracy
        best_svm = svm

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved: %f' % best_val)

```

```

iteration 0 / 1500: loss 84.994982
iteration 100 / 1500: loss 83.478411
iteration 200 / 1500: loss 82.010806
iteration 300 / 1500: loss 80.574769
iteration 400 / 1500: loss 79.148311
iteration 500 / 1500: loss 77.752788
iteration 600 / 1500: loss 76.396646
iteration 700 / 1500: loss 75.051871
iteration 800 / 1500: loss 73.736667
iteration 900 / 1500: loss 72.461444
iteration 1000 / 1500: loss 71.210670
iteration 1100 / 1500: loss 69.974092
iteration 1200 / 1500: loss 68.776618
iteration 1300 / 1500: loss 67.579058
iteration 1400 / 1500: loss 66.412083
iteration 0 / 1500: loss 851.575606
iteration 100 / 1500: loss 698.765416
iteration 200 / 1500: loss 573.684546
iteration 300 / 1500: loss 471.286140
iteration 400 / 1500: loss 387.444504
iteration 500 / 1500: loss 318.816575
iteration 600 / 1500: loss 262.628420
iteration 700 / 1500: loss 216.637879
iteration 800 / 1500: loss 178.982539
iteration 900 / 1500: loss 148.149557
iteration 1000 / 1500: loss 122.915263
iteration 1100 / 1500: loss 102.257416
iteration 1200 / 1500: loss 85.342222
iteration 1300 / 1500: loss 71.500591
iteration 1400 / 1500: loss 60.165982
iteration 0 / 1500: loss 7460.906679
iteration 100 / 1500: loss 1007.408351
iteration 200 / 1500: loss 142.764340

```

iteration 300 / 1500: loss 26.921708  
iteration 400 / 1500: loss 11.401231  
iteration 500 / 1500: loss 9.321751  
iteration 600 / 1500: loss 9.043087  
iteration 700 / 1500: loss 9.005767  
iteration 800 / 1500: loss 9.000767  
iteration 900 / 1500: loss 9.000099  
iteration 1000 / 1500: loss 9.000011  
iteration 1100 / 1500: loss 8.999999  
iteration 1200 / 1500: loss 8.999997  
iteration 1300 / 1500: loss 8.999997  
iteration 1400 / 1500: loss 8.999996  
iteration 0 / 1500: loss 85.490349  
iteration 100 / 1500: loss 71.631637  
iteration 200 / 1500: loss 60.275785  
iteration 300 / 1500: loss 50.985872  
iteration 400 / 1500: loss 43.360973  
iteration 500 / 1500: loss 37.131303  
iteration 600 / 1500: loss 32.028071  
iteration 700 / 1500: loss 27.858206  
iteration 800 / 1500: loss 24.437240  
iteration 900 / 1500: loss 21.630859  
iteration 1000 / 1500: loss 19.340107  
iteration 1100 / 1500: loss 17.462118  
iteration 1200 / 1500: loss 15.932119  
iteration 1300 / 1500: loss 14.675586  
iteration 1400 / 1500: loss 13.644937  
iteration 0 / 1500: loss 791.956456  
iteration 100 / 1500: loss 113.900367  
iteration 200 / 1500: loss 23.054552  
iteration 300 / 1500: loss 10.883075  
iteration 400 / 1500: loss 9.252086  
iteration 500 / 1500: loss 9.033754  
iteration 600 / 1500: loss 9.004495  
iteration 700 / 1500: loss 9.000568  
iteration 800 / 1500: loss 9.000046  
iteration 900 / 1500: loss 8.999972  
iteration 1000 / 1500: loss 8.999967  
iteration 1100 / 1500: loss 8.999961  
iteration 1200 / 1500: loss 8.999959  
iteration 1300 / 1500: loss 8.999969  
iteration 1400 / 1500: loss 8.999967  
iteration 0 / 1500: loss 7977.777592  
iteration 100 / 1500: loss 9.000003  
iteration 200 / 1500: loss 8.999996  
iteration 300 / 1500: loss 8.999997  
iteration 400 / 1500: loss 8.999996  
iteration 500 / 1500: loss 8.999997

iteration 600 / 1500: loss 8.999997  
iteration 700 / 1500: loss 8.999997  
iteration 800 / 1500: loss 8.999997  
iteration 900 / 1500: loss 8.999997  
iteration 1000 / 1500: loss 8.999997  
iteration 1100 / 1500: loss 8.999996  
iteration 1200 / 1500: loss 8.999997  
iteration 1300 / 1500: loss 8.999996  
iteration 1400 / 1500: loss 8.999996  
iteration 0 / 1500: loss 81.803552  
iteration 100 / 1500: loss 18.751479  
iteration 200 / 1500: loss 10.305304  
iteration 300 / 1500: loss 9.174871  
iteration 400 / 1500: loss 9.022779  
iteration 500 / 1500: loss 9.002765  
iteration 600 / 1500: loss 9.000093  
iteration 700 / 1500: loss 8.999686  
iteration 800 / 1500: loss 8.999651  
iteration 900 / 1500: loss 8.999629  
iteration 1000 / 1500: loss 8.999776  
iteration 1100 / 1500: loss 8.999660  
iteration 1200 / 1500: loss 8.999689  
iteration 1300 / 1500: loss 8.999630  
iteration 1400 / 1500: loss 8.999698  
iteration 0 / 1500: loss 785.828069  
iteration 100 / 1500: loss 8.999963  
iteration 200 / 1500: loss 8.999972  
iteration 300 / 1500: loss 8.999964  
iteration 400 / 1500: loss 8.999966  
iteration 500 / 1500: loss 8.999957  
iteration 600 / 1500: loss 8.999965  
iteration 700 / 1500: loss 8.999964  
iteration 800 / 1500: loss 8.999972  
iteration 900 / 1500: loss 8.999972  
iteration 1000 / 1500: loss 8.999966  
iteration 1100 / 1500: loss 8.999969  
iteration 1200 / 1500: loss 8.999977  
iteration 1300 / 1500: loss 8.999965  
iteration 1400 / 1500: loss 8.999973  
iteration 0 / 1500: loss 7713.563991  
iteration 100 / 1500: loss 9.000000  
iteration 200 / 1500: loss 9.000000  
iteration 300 / 1500: loss 8.999999  
iteration 400 / 1500: loss 9.000001  
iteration 500 / 1500: loss 9.000000  
iteration 600 / 1500: loss 9.000000  
iteration 700 / 1500: loss 8.999999  
iteration 800 / 1500: loss 8.999999

```

iteration 900 / 1500: loss 9.000000
iteration 1000 / 1500: loss 9.000000
iteration 1100 / 1500: loss 9.000000
iteration 1200 / 1500: loss 9.000000
iteration 1300 / 1500: loss 9.000000
iteration 1400 / 1500: loss 9.000000
lr 1.000000e-09 reg 5.000000e+04 train accuracy: 0.083388 val accuracy: 0.082000
lr 1.000000e-09 reg 5.000000e+05 train accuracy: 0.094122 val accuracy: 0.078000
lr 1.000000e-09 reg 5.000000e+06 train accuracy: 0.415367 val accuracy: 0.412000
lr 1.000000e-08 reg 5.000000e+04 train accuracy: 0.126163 val accuracy: 0.141000
lr 1.000000e-08 reg 5.000000e+05 train accuracy: 0.412551 val accuracy: 0.409000
lr 1.000000e-08 reg 5.000000e+06 train accuracy: 0.412041 val accuracy: 0.408000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.417306 val accuracy: 0.416000
lr 1.000000e-07 reg 5.000000e+05 train accuracy: 0.401673 val accuracy: 0.393000
lr 1.000000e-07 reg 5.000000e+06 train accuracy: 0.294694 val accuracy: 0.301000
best validation accuracy achieved: 0.416000

```

```

[8]: # Evaluate your trained SVM on the test set: you should be able to get at least
      ↪ 0.40
y_test_pred = best_svm.predict(X_test_feats)
test_accuracy = np.mean(y_test == y_test_pred)
print(test_accuracy)

```

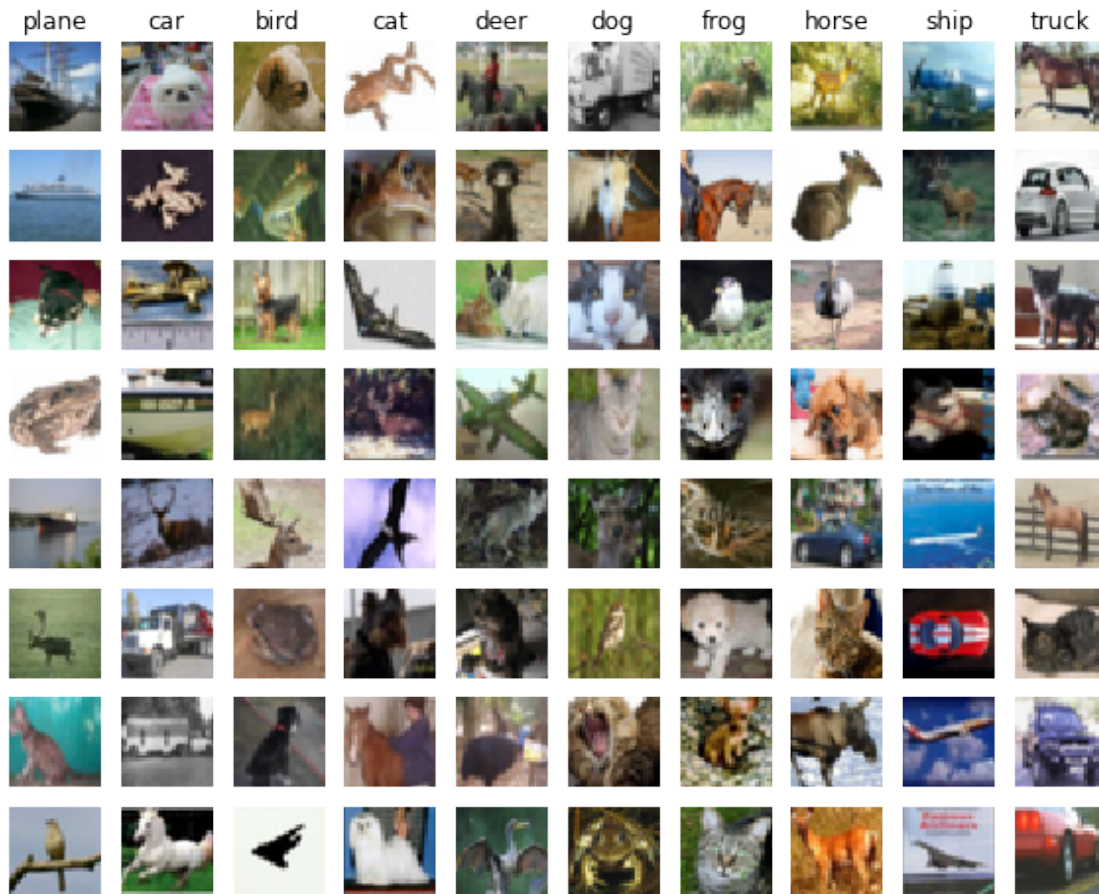
0.421

```

[9]: # An important way to gain intuition about how an algorithm works is to
      # visualize the mistakes that it makes. In this visualization, we show examples
      # of images that are misclassified by our current system. The first column
      # shows images that our system labeled as "plane" but whose true label is
      # something other than "plane".

examples_per_class = 8
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse',
          ↪ 'ship', 'truck']
for cls, cls_name in enumerate(classes):
    idxs = np.where((y_test != cls) & (y_test_pred == cls))[0]
    idxs = np.random.choice(idxs, examples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt.subplot(examples_per_class, len(classes), i * len(classes) + cls +
          ↪ 1)
        plt.imshow(X_test[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls_name)
plt.show()

```



### 1.3.1 Inline question 1:

Describe the misclassification results that you see. Do they make sense?

Most of classification is wrong. It doesn't make sense

## 1.4 Neural Network on image features

Earlier in this assignment we saw that training a two-layer neural network on raw pixels achieved better classification performance than linear classifiers on raw pixels. In this notebook we have seen that linear classifiers on image features outperform linear classifiers on raw pixels.

For completeness, we should also try training a neural network on image features. This approach should outperform all previous approaches: you should easily be able to achieve over 55% classification accuracy on the test set; our best model achieves about 60% classification accuracy.

```
[12]: # Preprocessing: Remove the bias dimension
      # Make sure to run this cell only ONCE
```

```

print(X_train_feats.shape)
X_train_feats = X_train_feats[:, :-1]
X_val_feats = X_val_feats[:, :-1]
X_test_feats = X_test_feats[:, :-1]

print(X_train_feats.shape)

```

(49000, 154)

(49000, 153)

```

[17]: from numpy.random.mtrand import noncentral_f
      from cs231n.classifiers.fc_net import TwoLayerNet
      from cs231n.solver import Solver

input_dim = X_train_feats.shape[1]
hidden_dim = 500
num_classes = 10

data = {
    'X_train': X_train_feats,
    'y_train': y_train,
    'X_val': X_val_feats,
    'y_val': y_val,
    'X_test': X_test_feats,
    'y_test': y_test,
}

net = TwoLayerNet(input_dim, hidden_dim, num_classes)
best_net = None
best_solver= None

#####
# TODO: Train a two-layer neural network on image features. You may want to      #
# cross-validate various parameters as in previous sections. Store your best    #
# model in the best_net variable.                                              #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
for _learning_rate in (0.01,0.02,0.03,0.04,0.05):
    for _lr_decay in (0.95, 0.90):
        for _reg in range(3):
            _reg= np.random.uniform(0, 0.01)
            model = TwoLayerNet(input_dim, hidden_dim, num_classes, reg = _reg)
            solver = Solver(model, data,
                            update_rule='sgd',
                            optim_config={
                                'learning_rate': _learning_rate,
                            },

```

```

        lr_decay=lr_decay,
        num_epochs=10,
        batch_size=100,
        print_every=100)

    print("reg: {}, lr: {}, lr_decay: {}".format(_reg, _learning_rate,
→_lr_decay))

    solver.train()

    if best_net == None:
        best_net = model
        best_solver = solver

    else:
        if solver.best_val_acc > best_solver.best_val_acc:
            best_net = model
            best_solver= solver
            print("Best model updated! Current best_val_acc = {}".
→format(best_solver.best_val_acc))
        else:
            print("Best model not updated! Current best_val_acc = {}".
→format(best_solver.best_val_acc))

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

```

```

reg: 0.0082856780016254, lr: 0.01, lr_decay: 0.95
(Iteration 1 / 4900) loss: 2.302952
(Epoch 0 / 10) train acc: 0.081000; val_acc: 0.098000
(Iteration 101 / 4900) loss: 2.303068
(Iteration 201 / 4900) loss: 2.302702
(Iteration 301 / 4900) loss: 2.302438
(Iteration 401 / 4900) loss: 2.301651
(Epoch 1 / 10) train acc: 0.147000; val_acc: 0.163000
(Iteration 501 / 4900) loss: 2.300065
(Iteration 601 / 4900) loss: 2.299623
(Iteration 701 / 4900) loss: 2.294666
(Iteration 801 / 4900) loss: 2.286044
(Iteration 901 / 4900) loss: 2.278483
(Epoch 2 / 10) train acc: 0.222000; val_acc: 0.237000
(Iteration 1001 / 4900) loss: 2.237441
(Iteration 1101 / 4900) loss: 2.221596
(Iteration 1201 / 4900) loss: 2.195081
(Iteration 1301 / 4900) loss: 2.096610
(Iteration 1401 / 4900) loss: 2.174302
(Epoch 3 / 10) train acc: 0.296000; val_acc: 0.281000

```

(Iteration 1501 / 4900) loss: 2.063345  
(Iteration 1601 / 4900) loss: 1.937163  
(Iteration 1701 / 4900) loss: 1.948185  
(Iteration 1801 / 4900) loss: 1.927292  
(Iteration 1901 / 4900) loss: 1.889842  
(Epoch 4 / 10) train acc: 0.292000; val\_acc: 0.317000  
(Iteration 2001 / 4900) loss: 1.914234  
(Iteration 2101 / 4900) loss: 1.914068  
(Iteration 2201 / 4900) loss: 1.773900  
(Iteration 2301 / 4900) loss: 1.761124  
(Iteration 2401 / 4900) loss: 1.754307  
(Epoch 5 / 10) train acc: 0.360000; val\_acc: 0.366000  
(Iteration 2501 / 4900) loss: 1.698999  
(Iteration 2601 / 4900) loss: 1.738604  
(Iteration 2701 / 4900) loss: 1.736538  
(Iteration 2801 / 4900) loss: 1.708529  
(Iteration 2901 / 4900) loss: 1.661625  
(Epoch 6 / 10) train acc: 0.431000; val\_acc: 0.397000  
(Iteration 3001 / 4900) loss: 1.866120  
(Iteration 3101 / 4900) loss: 1.846813  
(Iteration 3201 / 4900) loss: 1.586486  
(Iteration 3301 / 4900) loss: 1.594109  
(Iteration 3401 / 4900) loss: 1.599612  
(Epoch 7 / 10) train acc: 0.452000; val\_acc: 0.418000  
(Iteration 3501 / 4900) loss: 1.604730  
(Iteration 3601 / 4900) loss: 1.687003  
(Iteration 3701 / 4900) loss: 1.628311  
(Iteration 3801 / 4900) loss: 1.673763  
(Iteration 3901 / 4900) loss: 1.474054  
(Epoch 8 / 10) train acc: 0.452000; val\_acc: 0.441000  
(Iteration 4001 / 4900) loss: 1.486433  
(Iteration 4101 / 4900) loss: 1.352856  
(Iteration 4201 / 4900) loss: 1.643070  
(Iteration 4301 / 4900) loss: 1.433257  
(Iteration 4401 / 4900) loss: 1.608780  
(Epoch 9 / 10) train acc: 0.479000; val\_acc: 0.459000  
(Iteration 4501 / 4900) loss: 1.582180  
(Iteration 4601 / 4900) loss: 1.607665  
(Iteration 4701 / 4900) loss: 1.542218  
(Iteration 4801 / 4900) loss: 1.509256  
(Epoch 10 / 10) train acc: 0.460000; val\_acc: 0.458000  
reg: 0.004422818187092686, lr: 0.01, lr\_decay: 0.95  
(Iteration 1 / 4900) loss: 2.302764  
(Epoch 0 / 10) train acc: 0.089000; val\_acc: 0.102000  
(Iteration 101 / 4900) loss: 2.302363  
(Iteration 201 / 4900) loss: 2.301851  
(Iteration 301 / 4900) loss: 2.302038  
(Iteration 401 / 4900) loss: 2.301744



(Epoch 1 / 10) train acc: 0.108000; val\_acc: 0.093000  
(Iteration 501 / 4900) loss: 2.300212  
(Iteration 601 / 4900) loss: 2.298061  
(Iteration 701 / 4900) loss: 2.292461  
(Iteration 801 / 4900) loss: 2.287209  
(Iteration 901 / 4900) loss: 2.284690  
(Epoch 2 / 10) train acc: 0.246000; val\_acc: 0.261000  
(Iteration 1001 / 4900) loss: 2.253825  
(Iteration 1101 / 4900) loss: 2.208090  
(Iteration 1201 / 4900) loss: 2.183988  
(Iteration 1301 / 4900) loss: 2.129307  
(Iteration 1401 / 4900) loss: 2.128592  
(Epoch 3 / 10) train acc: 0.238000; val\_acc: 0.266000  
(Iteration 1501 / 4900) loss: 2.013006  
(Iteration 1601 / 4900) loss: 1.887733  
(Iteration 1701 / 4900) loss: 1.913889  
(Iteration 1801 / 4900) loss: 1.858685  
(Iteration 1901 / 4900) loss: 1.955646  
(Epoch 4 / 10) train acc: 0.345000; val\_acc: 0.327000  
(Iteration 2001 / 4900) loss: 1.766364  
(Iteration 2101 / 4900) loss: 1.825416  
(Iteration 2201 / 4900) loss: 1.904902  
(Iteration 2301 / 4900) loss: 1.821696  
(Iteration 2401 / 4900) loss: 1.787852  
(Epoch 5 / 10) train acc: 0.370000; val\_acc: 0.381000  
(Iteration 2501 / 4900) loss: 1.711158  
(Iteration 2601 / 4900) loss: 1.782857  
(Iteration 2701 / 4900) loss: 1.803401  
(Iteration 2801 / 4900) loss: 1.673248  
(Iteration 2901 / 4900) loss: 1.641201  
(Epoch 6 / 10) train acc: 0.408000; val\_acc: 0.402000  
(Iteration 3001 / 4900) loss: 1.697081  
(Iteration 3101 / 4900) loss: 1.503331  
(Iteration 3201 / 4900) loss: 1.705221  
(Iteration 3301 / 4900) loss: 1.558441  
(Iteration 3401 / 4900) loss: 1.624144  
(Epoch 7 / 10) train acc: 0.421000; val\_acc: 0.422000  
(Iteration 3501 / 4900) loss: 1.437250  
(Iteration 3601 / 4900) loss: 1.586389  
(Iteration 3701 / 4900) loss: 1.659224  
(Iteration 3801 / 4900) loss: 1.613325  
(Iteration 3901 / 4900) loss: 1.510576  
(Epoch 8 / 10) train acc: 0.463000; val\_acc: 0.441000  
(Iteration 4001 / 4900) loss: 1.728436  
(Iteration 4101 / 4900) loss: 1.462304  
(Iteration 4201 / 4900) loss: 1.629648  
(Iteration 4301 / 4900) loss: 1.522504  
(Iteration 4401 / 4900) loss: 1.559979

(Epoch 9 / 10) train acc: 0.481000; val\_acc: 0.452000  
(Iteration 4501 / 4900) loss: 1.427136  
(Iteration 4601 / 4900) loss: 1.442288  
(Iteration 4701 / 4900) loss: 1.462266  
(Iteration 4801 / 4900) loss: 1.432488  
(Epoch 10 / 10) train acc: 0.472000; val\_acc: 0.471000  
Best model updated! Current best\_val\_acc = 0.471  
reg: 0.004042896197502762, lr: 0.01, lr\_decay: 0.95  
(Iteration 1 / 4900) loss: 2.302748  
(Epoch 0 / 10) train acc: 0.105000; val\_acc: 0.119000  
(Iteration 101 / 4900) loss: 2.302444  
(Iteration 201 / 4900) loss: 2.303289  
(Iteration 301 / 4900) loss: 2.302276  
(Iteration 401 / 4900) loss: 2.301569  
(Epoch 1 / 10) train acc: 0.214000; val\_acc: 0.226000  
(Iteration 501 / 4900) loss: 2.300446  
(Iteration 601 / 4900) loss: 2.299462  
(Iteration 701 / 4900) loss: 2.295401  
(Iteration 801 / 4900) loss: 2.291823  
(Iteration 901 / 4900) loss: 2.269293  
(Epoch 2 / 10) train acc: 0.259000; val\_acc: 0.263000  
(Iteration 1001 / 4900) loss: 2.268683  
(Iteration 1101 / 4900) loss: 2.247209  
(Iteration 1201 / 4900) loss: 2.190293  
(Iteration 1301 / 4900) loss: 2.121814  
(Iteration 1401 / 4900) loss: 2.066860  
(Epoch 3 / 10) train acc: 0.260000; val\_acc: 0.282000  
(Iteration 1501 / 4900) loss: 2.052916  
(Iteration 1601 / 4900) loss: 2.017226  
(Iteration 1701 / 4900) loss: 2.008096  
(Iteration 1801 / 4900) loss: 1.940882  
(Iteration 1901 / 4900) loss: 1.970122  
(Epoch 4 / 10) train acc: 0.334000; val\_acc: 0.338000  
(Iteration 2001 / 4900) loss: 1.797432  
(Iteration 2101 / 4900) loss: 1.836631  
(Iteration 2201 / 4900) loss: 1.739533  
(Iteration 2301 / 4900) loss: 1.770812  
(Iteration 2401 / 4900) loss: 1.707155  
(Epoch 5 / 10) train acc: 0.391000; val\_acc: 0.373000  
(Iteration 2501 / 4900) loss: 1.756946  
(Iteration 2601 / 4900) loss: 1.806035  
(Iteration 2701 / 4900) loss: 1.687534  
(Iteration 2801 / 4900) loss: 1.706342  
(Iteration 2901 / 4900) loss: 1.598212  
(Epoch 6 / 10) train acc: 0.414000; val\_acc: 0.398000  
(Iteration 3001 / 4900) loss: 1.728750  
(Iteration 3101 / 4900) loss: 1.624681  
(Iteration 3201 / 4900) loss: 1.751387

(Iteration 3301 / 4900) loss: 1.688346  
(Iteration 3401 / 4900) loss: 1.640967  
(Epoch 7 / 10) train acc: 0.422000; val\_acc: 0.420000  
(Iteration 3501 / 4900) loss: 1.609378  
(Iteration 3601 / 4900) loss: 1.479075  
(Iteration 3701 / 4900) loss: 1.556336  
(Iteration 3801 / 4900) loss: 1.559605  
(Iteration 3901 / 4900) loss: 1.413187  
(Epoch 8 / 10) train acc: 0.465000; val\_acc: 0.442000  
(Iteration 4001 / 4900) loss: 1.555094  
(Iteration 4101 / 4900) loss: 1.534199  
(Iteration 4201 / 4900) loss: 1.431372  
(Iteration 4301 / 4900) loss: 1.517576  
(Iteration 4401 / 4900) loss: 1.501845  
(Epoch 9 / 10) train acc: 0.478000; val\_acc: 0.461000  
(Iteration 4501 / 4900) loss: 1.681407  
(Iteration 4601 / 4900) loss: 1.447981  
(Iteration 4701 / 4900) loss: 1.518066  
(Iteration 4801 / 4900) loss: 1.555409  
(Epoch 10 / 10) train acc: 0.481000; val\_acc: 0.475000  
Best model updated! Current best\_val\_acc = 0.475  
reg: 0.0013879040383895781, lr: 0.01, lr\_decay: 0.9  
(Iteration 1 / 4900) loss: 2.302633  
(Epoch 0 / 10) train acc: 0.130000; val\_acc: 0.142000  
(Iteration 101 / 4900) loss: 2.302458  
(Iteration 201 / 4900) loss: 2.302240  
(Iteration 301 / 4900) loss: 2.302584  
(Iteration 401 / 4900) loss: 2.300332  
(Epoch 1 / 10) train acc: 0.191000; val\_acc: 0.230000  
(Iteration 501 / 4900) loss: 2.300075  
(Iteration 601 / 4900) loss: 2.296833  
(Iteration 701 / 4900) loss: 2.293020  
(Iteration 801 / 4900) loss: 2.285628  
(Iteration 901 / 4900) loss: 2.276788  
(Epoch 2 / 10) train acc: 0.290000; val\_acc: 0.264000  
(Iteration 1001 / 4900) loss: 2.256149  
(Iteration 1101 / 4900) loss: 2.225298  
(Iteration 1201 / 4900) loss: 2.189077  
(Iteration 1301 / 4900) loss: 2.153887  
(Iteration 1401 / 4900) loss: 2.162795  
(Epoch 3 / 10) train acc: 0.269000; val\_acc: 0.265000  
(Iteration 1501 / 4900) loss: 2.081289  
(Iteration 1601 / 4900) loss: 2.070573  
(Iteration 1701 / 4900) loss: 1.927056  
(Iteration 1801 / 4900) loss: 1.920348  
(Iteration 1901 / 4900) loss: 1.979611  
(Epoch 4 / 10) train acc: 0.326000; val\_acc: 0.307000  
(Iteration 2001 / 4900) loss: 1.840644

(Iteration 2101 / 4900) loss: 1.987865  
(Iteration 2201 / 4900) loss: 1.785827  
(Iteration 2301 / 4900) loss: 1.797904  
(Iteration 2401 / 4900) loss: 1.726353  
(Epoch 5 / 10) train acc: 0.351000; val\_acc: 0.354000  
(Iteration 2501 / 4900) loss: 1.847005  
(Iteration 2601 / 4900) loss: 1.671093  
(Iteration 2701 / 4900) loss: 1.819464  
(Iteration 2801 / 4900) loss: 1.679568  
(Iteration 2901 / 4900) loss: 1.659404  
(Epoch 6 / 10) train acc: 0.374000; val\_acc: 0.378000  
(Iteration 3001 / 4900) loss: 1.724590  
(Iteration 3101 / 4900) loss: 1.559886  
(Iteration 3201 / 4900) loss: 1.587863  
(Iteration 3301 / 4900) loss: 1.701093  
(Iteration 3401 / 4900) loss: 1.577733  
(Epoch 7 / 10) train acc: 0.415000; val\_acc: 0.402000  
(Iteration 3501 / 4900) loss: 1.666840  
(Iteration 3601 / 4900) loss: 1.746563  
(Iteration 3701 / 4900) loss: 1.623177  
(Iteration 3801 / 4900) loss: 1.673230  
(Iteration 3901 / 4900) loss: 1.642476  
(Epoch 8 / 10) train acc: 0.426000; val\_acc: 0.414000  
(Iteration 4001 / 4900) loss: 1.593319  
(Iteration 4101 / 4900) loss: 1.703651  
(Iteration 4201 / 4900) loss: 1.487295  
(Iteration 4301 / 4900) loss: 1.511337  
(Iteration 4401 / 4900) loss: 1.526084  
(Epoch 9 / 10) train acc: 0.429000; val\_acc: 0.423000  
(Iteration 4501 / 4900) loss: 1.599901  
(Iteration 4601 / 4900) loss: 1.522028  
(Iteration 4701 / 4900) loss: 1.558974  
(Iteration 4801 / 4900) loss: 1.509707  
(Epoch 10 / 10) train acc: 0.461000; val\_acc: 0.435000  
Best model not updated! Current best\_val\_acc = 0.475  
reg: 0.004541117088053265, lr: 0.01, lr\_decay: 0.9  
(Iteration 1 / 4900) loss: 2.302773  
(Epoch 0 / 10) train acc: 0.082000; val\_acc: 0.120000  
(Iteration 101 / 4900) loss: 2.302613  
(Iteration 201 / 4900) loss: 2.302002  
(Iteration 301 / 4900) loss: 2.302301  
(Iteration 401 / 4900) loss: 2.300412  
(Epoch 1 / 10) train acc: 0.178000; val\_acc: 0.140000  
(Iteration 501 / 4900) loss: 2.300260  
(Iteration 601 / 4900) loss: 2.299101  
(Iteration 701 / 4900) loss: 2.295326  
(Iteration 801 / 4900) loss: 2.289698  
(Iteration 901 / 4900) loss: 2.277526

(Epoch 2 / 10) train acc: 0.279000; val\_acc: 0.273000  
(Iteration 1001 / 4900) loss: 2.258337  
(Iteration 1101 / 4900) loss: 2.245316  
(Iteration 1201 / 4900) loss: 2.170239  
(Iteration 1301 / 4900) loss: 2.173962  
(Iteration 1401 / 4900) loss: 2.122456  
(Epoch 3 / 10) train acc: 0.274000; val\_acc: 0.281000  
(Iteration 1501 / 4900) loss: 2.087637  
(Iteration 1601 / 4900) loss: 2.104578  
(Iteration 1701 / 4900) loss: 2.010335  
(Iteration 1801 / 4900) loss: 2.030183  
(Iteration 1901 / 4900) loss: 1.926631  
(Epoch 4 / 10) train acc: 0.312000; val\_acc: 0.331000  
(Iteration 2001 / 4900) loss: 1.978199  
(Iteration 2101 / 4900) loss: 1.926773  
(Iteration 2201 / 4900) loss: 1.933867  
(Iteration 2301 / 4900) loss: 1.909266  
(Iteration 2401 / 4900) loss: 1.739471  
(Epoch 5 / 10) train acc: 0.357000; val\_acc: 0.355000  
(Iteration 2501 / 4900) loss: 1.879235  
(Iteration 2601 / 4900) loss: 1.775321  
(Iteration 2701 / 4900) loss: 1.738343  
(Iteration 2801 / 4900) loss: 1.805363  
(Iteration 2901 / 4900) loss: 1.687363  
(Epoch 6 / 10) train acc: 0.408000; val\_acc: 0.389000  
(Iteration 3001 / 4900) loss: 1.715237  
(Iteration 3101 / 4900) loss: 1.760099  
(Iteration 3201 / 4900) loss: 1.657720  
(Iteration 3301 / 4900) loss: 1.612720  
(Iteration 3401 / 4900) loss: 1.702443  
(Epoch 7 / 10) train acc: 0.406000; val\_acc: 0.395000  
(Iteration 3501 / 4900) loss: 1.519267  
(Iteration 3601 / 4900) loss: 1.608164  
(Iteration 3701 / 4900) loss: 1.703807  
(Iteration 3801 / 4900) loss: 1.604700  
(Iteration 3901 / 4900) loss: 1.602502  
(Epoch 8 / 10) train acc: 0.433000; val\_acc: 0.413000  
(Iteration 4001 / 4900) loss: 1.706013  
(Iteration 4101 / 4900) loss: 1.555866  
(Iteration 4201 / 4900) loss: 1.616726  
(Iteration 4301 / 4900) loss: 1.498019  
(Iteration 4401 / 4900) loss: 1.575365  
(Epoch 9 / 10) train acc: 0.408000; val\_acc: 0.422000  
(Iteration 4501 / 4900) loss: 1.600104  
(Iteration 4601 / 4900) loss: 1.530999  
(Iteration 4701 / 4900) loss: 1.563508  
(Iteration 4801 / 4900) loss: 1.456248  
(Epoch 10 / 10) train acc: 0.442000; val\_acc: 0.432000

Best model not updated! Current best\_val\_acc = 0.475  
reg: 0.006352794645402055, lr: 0.01, lr\_decay: 0.9  
(Iteration 1 / 4900) loss: 2.302879  
(Epoch 0 / 10) train acc: 0.114000; val\_acc: 0.080000  
(Iteration 101 / 4900) loss: 2.302296  
(Iteration 201 / 4900) loss: 2.302739  
(Iteration 301 / 4900) loss: 2.302533  
(Iteration 401 / 4900) loss: 2.301967  
(Epoch 1 / 10) train acc: 0.218000; val\_acc: 0.203000  
(Iteration 501 / 4900) loss: 2.299993  
(Iteration 601 / 4900) loss: 2.297253  
(Iteration 701 / 4900) loss: 2.293996  
(Iteration 801 / 4900) loss: 2.287245  
(Iteration 901 / 4900) loss: 2.276151  
(Epoch 2 / 10) train acc: 0.278000; val\_acc: 0.267000  
(Iteration 1001 / 4900) loss: 2.244938  
(Iteration 1101 / 4900) loss: 2.252318  
(Iteration 1201 / 4900) loss: 2.204819  
(Iteration 1301 / 4900) loss: 2.193363  
(Iteration 1401 / 4900) loss: 2.147631  
(Epoch 3 / 10) train acc: 0.280000; val\_acc: 0.279000  
(Iteration 1501 / 4900) loss: 2.091074  
(Iteration 1601 / 4900) loss: 2.074451  
(Iteration 1701 / 4900) loss: 1.995389  
(Iteration 1801 / 4900) loss: 2.043981  
(Iteration 1901 / 4900) loss: 1.965283  
(Epoch 4 / 10) train acc: 0.304000; val\_acc: 0.311000  
(Iteration 2001 / 4900) loss: 1.871330  
(Iteration 2101 / 4900) loss: 1.872956  
(Iteration 2201 / 4900) loss: 1.942090  
(Iteration 2301 / 4900) loss: 1.953274  
(Iteration 2401 / 4900) loss: 1.829346  
(Epoch 5 / 10) train acc: 0.363000; val\_acc: 0.358000  
(Iteration 2501 / 4900) loss: 1.796759  
(Iteration 2601 / 4900) loss: 1.848936  
(Iteration 2701 / 4900) loss: 1.753758  
(Iteration 2801 / 4900) loss: 1.806299  
(Iteration 2901 / 4900) loss: 1.741331  
(Epoch 6 / 10) train acc: 0.386000; val\_acc: 0.376000  
(Iteration 3001 / 4900) loss: 1.690842  
(Iteration 3101 / 4900) loss: 1.599711  
(Iteration 3201 / 4900) loss: 1.765151  
(Iteration 3301 / 4900) loss: 1.744950  
(Iteration 3401 / 4900) loss: 1.678391  
(Epoch 7 / 10) train acc: 0.408000; val\_acc: 0.391000  
(Iteration 3501 / 4900) loss: 1.638719  
(Iteration 3601 / 4900) loss: 1.723927  
(Iteration 3701 / 4900) loss: 1.744239

```

(Iteration 3801 / 4900) loss: 1.836110
(Iteration 3901 / 4900) loss: 1.755307
(Epoch 8 / 10) train acc: 0.410000; val_acc: 0.410000
(Iteration 4001 / 4900) loss: 1.555389
(Iteration 4101 / 4900) loss: 1.659318
(Iteration 4201 / 4900) loss: 1.550655
(Iteration 4301 / 4900) loss: 1.602639
(Iteration 4401 / 4900) loss: 1.682250
(Epoch 9 / 10) train acc: 0.415000; val_acc: 0.419000
(Iteration 4501 / 4900) loss: 1.486189
(Iteration 4601 / 4900) loss: 1.602723
(Iteration 4701 / 4900) loss: 1.575334
(Iteration 4801 / 4900) loss: 1.461706
(Epoch 10 / 10) train acc: 0.432000; val_acc: 0.428000
Best model not updated! Current best_val_acc = 0.475
reg: 0.0010600803465194918, lr: 0.02, lr_decay: 0.95
(Iteration 1 / 4900) loss: 2.302609
(Epoch 0 / 10) train acc: 0.105000; val_acc: 0.087000
(Iteration 101 / 4900) loss: 2.302531
(Iteration 201 / 4900) loss: 2.301460
(Iteration 301 / 4900) loss: 2.296980
(Iteration 401 / 4900) loss: 2.283699
(Epoch 1 / 10) train acc: 0.255000; val_acc: 0.256000
(Iteration 501 / 4900) loss: 2.239671
(Iteration 601 / 4900) loss: 2.163001
(Iteration 701 / 4900) loss: 2.025441
(Iteration 801 / 4900) loss: 1.852932
(Iteration 901 / 4900) loss: 1.887833
(Epoch 2 / 10) train acc: 0.325000; val_acc: 0.347000
(Iteration 1001 / 4900) loss: 1.892246
(Iteration 1101 / 4900) loss: 1.815852
(Iteration 1201 / 4900) loss: 1.612884
(Iteration 1301 / 4900) loss: 1.769701
(Iteration 1401 / 4900) loss: 1.527408
(Epoch 3 / 10) train acc: 0.447000; val_acc: 0.418000
(Iteration 1501 / 4900) loss: 1.524845
(Iteration 1601 / 4900) loss: 1.554797
(Iteration 1701 / 4900) loss: 1.479290
(Iteration 1801 / 4900) loss: 1.554796
(Iteration 1901 / 4900) loss: 1.510538
(Epoch 4 / 10) train acc: 0.458000; val_acc: 0.462000
(Iteration 2001 / 4900) loss: 1.382540
(Iteration 2101 / 4900) loss: 1.496296
(Iteration 2201 / 4900) loss: 1.439898
(Iteration 2301 / 4900) loss: 1.496867
(Iteration 2401 / 4900) loss: 1.521149
(Epoch 5 / 10) train acc: 0.508000; val_acc: 0.498000
(Iteration 2501 / 4900) loss: 1.445834

```

(Iteration 2601 / 4900) loss: 1.614661  
(Iteration 2701 / 4900) loss: 1.642059  
(Iteration 2801 / 4900) loss: 1.381352  
(Iteration 2901 / 4900) loss: 1.413971  
(Epoch 6 / 10) train acc: 0.498000; val\_acc: 0.495000  
(Iteration 3001 / 4900) loss: 1.373956  
(Iteration 3101 / 4900) loss: 1.450300  
(Iteration 3201 / 4900) loss: 1.433239  
(Iteration 3301 / 4900) loss: 1.341114  
(Iteration 3401 / 4900) loss: 1.479868  
(Epoch 7 / 10) train acc: 0.502000; val\_acc: 0.503000  
(Iteration 3501 / 4900) loss: 1.383100  
(Iteration 3601 / 4900) loss: 1.366725  
(Iteration 3701 / 4900) loss: 1.466683  
(Iteration 3801 / 4900) loss: 1.404898  
(Iteration 3901 / 4900) loss: 1.532525  
(Epoch 8 / 10) train acc: 0.528000; val\_acc: 0.508000  
(Iteration 4001 / 4900) loss: 1.318628  
(Iteration 4101 / 4900) loss: 1.437999  
(Iteration 4201 / 4900) loss: 1.399376  
(Iteration 4301 / 4900) loss: 1.404009  
(Iteration 4401 / 4900) loss: 1.325312  
(Epoch 9 / 10) train acc: 0.508000; val\_acc: 0.515000  
(Iteration 4501 / 4900) loss: 1.222564  
(Iteration 4601 / 4900) loss: 1.323065  
(Iteration 4701 / 4900) loss: 1.451059  
(Iteration 4801 / 4900) loss: 1.485891  
(Epoch 10 / 10) train acc: 0.541000; val\_acc: 0.514000  
Best model updated! Current best\_val\_acc = 0.515  
reg: 0.007019815455129428, lr: 0.02, lr\_decay: 0.95  
(Iteration 1 / 4900) loss: 2.302849  
(Epoch 0 / 10) train acc: 0.085000; val\_acc: 0.113000  
(Iteration 101 / 4900) loss: 2.303050  
(Iteration 201 / 4900) loss: 2.301033  
(Iteration 301 / 4900) loss: 2.297904  
(Iteration 401 / 4900) loss: 2.285494  
(Epoch 1 / 10) train acc: 0.284000; val\_acc: 0.290000  
(Iteration 501 / 4900) loss: 2.255026  
(Iteration 601 / 4900) loss: 2.172744  
(Iteration 701 / 4900) loss: 2.052262  
(Iteration 801 / 4900) loss: 1.958689  
(Iteration 901 / 4900) loss: 1.906558  
(Epoch 2 / 10) train acc: 0.364000; val\_acc: 0.341000  
(Iteration 1001 / 4900) loss: 1.806598  
(Iteration 1101 / 4900) loss: 1.712235  
(Iteration 1201 / 4900) loss: 1.756357  
(Iteration 1301 / 4900) loss: 1.588448  
(Iteration 1401 / 4900) loss: 1.643616



```

(Epoch 3 / 10) train acc: 0.429000; val_acc: 0.413000
(Iteration 1501 / 4900) loss: 1.688299
(Iteration 1601 / 4900) loss: 1.802886
(Iteration 1701 / 4900) loss: 1.516247
(Iteration 1801 / 4900) loss: 1.613960
(Iteration 1901 / 4900) loss: 1.516186
(Epoch 4 / 10) train acc: 0.465000; val_acc: 0.460000
(Iteration 2001 / 4900) loss: 1.647518
(Iteration 2101 / 4900) loss: 1.473746
(Iteration 2201 / 4900) loss: 1.439705
(Iteration 2301 / 4900) loss: 1.564485
(Iteration 2401 / 4900) loss: 1.495032
(Epoch 5 / 10) train acc: 0.488000; val_acc: 0.486000
(Iteration 2501 / 4900) loss: 1.591337
(Iteration 2601 / 4900) loss: 1.574301
(Iteration 2701 / 4900) loss: 1.510283
(Iteration 2801 / 4900) loss: 1.523980
(Iteration 2901 / 4900) loss: 1.543734
(Epoch 6 / 10) train acc: 0.489000; val_acc: 0.500000
(Iteration 3001 / 4900) loss: 1.331654
(Iteration 3101 / 4900) loss: 1.489216
(Iteration 3201 / 4900) loss: 1.582673
(Iteration 3301 / 4900) loss: 1.392089
(Iteration 3401 / 4900) loss: 1.323347
(Epoch 7 / 10) train acc: 0.516000; val_acc: 0.504000
(Iteration 3501 / 4900) loss: 1.550671
(Iteration 3601 / 4900) loss: 1.612497
(Iteration 3701 / 4900) loss: 1.482155
(Iteration 3801 / 4900) loss: 1.352495
(Iteration 3901 / 4900) loss: 1.252472
(Epoch 8 / 10) train acc: 0.496000; val_acc: 0.519000
(Iteration 4001 / 4900) loss: 1.436499
(Iteration 4101 / 4900) loss: 1.420174
(Iteration 4201 / 4900) loss: 1.542864
(Iteration 4301 / 4900) loss: 1.356605
(Iteration 4401 / 4900) loss: 1.260831
(Epoch 9 / 10) train acc: 0.498000; val_acc: 0.509000
(Iteration 4501 / 4900) loss: 1.309612
(Iteration 4601 / 4900) loss: 1.442143
(Iteration 4701 / 4900) loss: 1.250552
(Iteration 4801 / 4900) loss: 1.546444
(Epoch 10 / 10) train acc: 0.554000; val_acc: 0.521000
Best model updated! Current best_val_acc = 0.521
reg: 0.005151684585375265, lr: 0.02, lr_decay: 0.95
(Iteration 1 / 4900) loss: 2.302786
(Epoch 0 / 10) train acc: 0.105000; val_acc: 0.107000
(Iteration 101 / 4900) loss: 2.301872
(Iteration 201 / 4900) loss: 2.301544

```

(Iteration 301 / 4900) loss: 2.297373  
(Iteration 401 / 4900) loss: 2.276950  
(Epoch 1 / 10) train acc: 0.198000; val\_acc: 0.232000  
(Iteration 501 / 4900) loss: 2.232660  
(Iteration 601 / 4900) loss: 2.144628  
(Iteration 701 / 4900) loss: 2.020608  
(Iteration 801 / 4900) loss: 1.946932  
(Iteration 901 / 4900) loss: 1.857572  
(Epoch 2 / 10) train acc: 0.349000; val\_acc: 0.360000  
(Iteration 1001 / 4900) loss: 1.863931  
(Iteration 1101 / 4900) loss: 1.770289  
(Iteration 1201 / 4900) loss: 1.743853  
(Iteration 1301 / 4900) loss: 1.703551  
(Iteration 1401 / 4900) loss: 1.792097  
(Epoch 3 / 10) train acc: 0.455000; val\_acc: 0.417000  
(Iteration 1501 / 4900) loss: 1.545361  
(Iteration 1601 / 4900) loss: 1.682266  
(Iteration 1701 / 4900) loss: 1.397306  
(Iteration 1801 / 4900) loss: 1.531347  
(Iteration 1901 / 4900) loss: 1.616176  
(Epoch 4 / 10) train acc: 0.465000; val\_acc: 0.457000  
(Iteration 2001 / 4900) loss: 1.634437  
(Iteration 2101 / 4900) loss: 1.382785  
(Iteration 2201 / 4900) loss: 1.597081  
(Iteration 2301 / 4900) loss: 1.435523  
(Iteration 2401 / 4900) loss: 1.425663  
(Epoch 5 / 10) train acc: 0.476000; val\_acc: 0.479000  
(Iteration 2501 / 4900) loss: 1.415965  
(Iteration 2601 / 4900) loss: 1.308795  
(Iteration 2701 / 4900) loss: 1.382718  
(Iteration 2801 / 4900) loss: 1.557069  
(Iteration 2901 / 4900) loss: 1.555524  
(Epoch 6 / 10) train acc: 0.516000; val\_acc: 0.504000  
(Iteration 3001 / 4900) loss: 1.352778  
(Iteration 3101 / 4900) loss: 1.544827  
(Iteration 3201 / 4900) loss: 1.430120  
(Iteration 3301 / 4900) loss: 1.527941  
(Iteration 3401 / 4900) loss: 1.576892  
(Epoch 7 / 10) train acc: 0.509000; val\_acc: 0.508000  
(Iteration 3501 / 4900) loss: 1.477605  
(Iteration 3601 / 4900) loss: 1.365561  
(Iteration 3701 / 4900) loss: 1.386682  
(Iteration 3801 / 4900) loss: 1.342671  
(Iteration 3901 / 4900) loss: 1.581286  
(Epoch 8 / 10) train acc: 0.536000; val\_acc: 0.515000  
(Iteration 4001 / 4900) loss: 1.310970  
(Iteration 4101 / 4900) loss: 1.237487  
(Iteration 4201 / 4900) loss: 1.317435

(Iteration 4301 / 4900) loss: 1.320935  
(Iteration 4401 / 4900) loss: 1.407751  
(Epoch 9 / 10) train acc: 0.566000; val\_acc: 0.513000  
(Iteration 4501 / 4900) loss: 1.255265  
(Iteration 4601 / 4900) loss: 1.458556  
(Iteration 4701 / 4900) loss: 1.301191  
(Iteration 4801 / 4900) loss: 1.407941  
(Epoch 10 / 10) train acc: 0.515000; val\_acc: 0.517000  
Best model not updated! Current best\_val\_acc = 0.521  
reg: 0.0019317508692079754, lr: 0.02, lr\_decay: 0.9  
(Iteration 1 / 4900) loss: 2.302615  
(Epoch 0 / 10) train acc: 0.114000; val\_acc: 0.110000  
(Iteration 101 / 4900) loss: 2.303442  
(Iteration 201 / 4900) loss: 2.301251  
(Iteration 301 / 4900) loss: 2.297069  
(Iteration 401 / 4900) loss: 2.286893  
(Epoch 1 / 10) train acc: 0.242000; val\_acc: 0.256000  
(Iteration 501 / 4900) loss: 2.234655  
(Iteration 601 / 4900) loss: 2.161380  
(Iteration 701 / 4900) loss: 1.989328  
(Iteration 801 / 4900) loss: 1.906629  
(Iteration 901 / 4900) loss: 2.008846  
(Epoch 2 / 10) train acc: 0.321000; val\_acc: 0.333000  
(Iteration 1001 / 4900) loss: 1.760331  
(Iteration 1101 / 4900) loss: 1.687476  
(Iteration 1201 / 4900) loss: 1.691303  
(Iteration 1301 / 4900) loss: 1.696342  
(Iteration 1401 / 4900) loss: 1.693222  
(Epoch 3 / 10) train acc: 0.418000; val\_acc: 0.404000  
(Iteration 1501 / 4900) loss: 1.597743  
(Iteration 1601 / 4900) loss: 1.540403  
(Iteration 1701 / 4900) loss: 1.510420  
(Iteration 1801 / 4900) loss: 1.525759  
(Iteration 1901 / 4900) loss: 1.491598  
(Epoch 4 / 10) train acc: 0.465000; val\_acc: 0.446000  
(Iteration 2001 / 4900) loss: 1.469113  
(Iteration 2101 / 4900) loss: 1.456815  
(Iteration 2201 / 4900) loss: 1.456389  
(Iteration 2301 / 4900) loss: 1.415455  
(Iteration 2401 / 4900) loss: 1.536049  
(Epoch 5 / 10) train acc: 0.483000; val\_acc: 0.477000  
(Iteration 2501 / 4900) loss: 1.413906  
(Iteration 2601 / 4900) loss: 1.423220  
(Iteration 2701 / 4900) loss: 1.456312  
(Iteration 2801 / 4900) loss: 1.449138  
(Iteration 2901 / 4900) loss: 1.622345  
(Epoch 6 / 10) train acc: 0.480000; val\_acc: 0.498000  
(Iteration 3001 / 4900) loss: 1.642100

(Iteration 3101 / 4900) loss: 1.408091  
(Iteration 3201 / 4900) loss: 1.463604  
(Iteration 3301 / 4900) loss: 1.467074  
(Iteration 3401 / 4900) loss: 1.427519  
(Epoch 7 / 10) train acc: 0.503000; val\_acc: 0.501000  
(Iteration 3501 / 4900) loss: 1.386555  
(Iteration 3601 / 4900) loss: 1.310342  
(Iteration 3701 / 4900) loss: 1.299070  
(Iteration 3801 / 4900) loss: 1.387008  
(Iteration 3901 / 4900) loss: 1.287588  
(Epoch 8 / 10) train acc: 0.514000; val\_acc: 0.511000  
(Iteration 4001 / 4900) loss: 1.339905  
(Iteration 4101 / 4900) loss: 1.391509  
(Iteration 4201 / 4900) loss: 1.270407  
(Iteration 4301 / 4900) loss: 1.450360  
(Iteration 4401 / 4900) loss: 1.580512  
(Epoch 9 / 10) train acc: 0.501000; val\_acc: 0.518000  
(Iteration 4501 / 4900) loss: 1.412303  
(Iteration 4601 / 4900) loss: 1.451452  
(Iteration 4701 / 4900) loss: 1.346848  
(Iteration 4801 / 4900) loss: 1.194317  
(Epoch 10 / 10) train acc: 0.476000; val\_acc: 0.515000  
Best model not updated! Current best\_val\_acc = 0.521  
reg: 0.002141341985272848, lr: 0.02, lr\_decay: 0.9  
(Iteration 1 / 4900) loss: 2.302723  
(Epoch 0 / 10) train acc: 0.093000; val\_acc: 0.117000  
(Iteration 101 / 4900) loss: 2.300360  
(Iteration 201 / 4900) loss: 2.302300  
(Iteration 301 / 4900) loss: 2.296404  
(Iteration 401 / 4900) loss: 2.283622  
(Epoch 1 / 10) train acc: 0.234000; val\_acc: 0.213000  
(Iteration 501 / 4900) loss: 2.246961  
(Iteration 601 / 4900) loss: 2.154027  
(Iteration 701 / 4900) loss: 2.092982  
(Iteration 801 / 4900) loss: 2.009880  
(Iteration 901 / 4900) loss: 1.909013  
(Epoch 2 / 10) train acc: 0.351000; val\_acc: 0.354000  
(Iteration 1001 / 4900) loss: 1.821060  
(Iteration 1101 / 4900) loss: 1.833776  
(Iteration 1201 / 4900) loss: 1.747348  
(Iteration 1301 / 4900) loss: 1.785415  
(Iteration 1401 / 4900) loss: 1.649440  
(Epoch 3 / 10) train acc: 0.436000; val\_acc: 0.415000  
(Iteration 1501 / 4900) loss: 1.615880  
(Iteration 1601 / 4900) loss: 1.590933  
(Iteration 1701 / 4900) loss: 1.471040  
(Iteration 1801 / 4900) loss: 1.621667  
(Iteration 1901 / 4900) loss: 1.537475

(Epoch 4 / 10) train acc: 0.437000; val\_acc: 0.444000  
(Iteration 2001 / 4900) loss: 1.467800  
(Iteration 2101 / 4900) loss: 1.555494  
(Iteration 2201 / 4900) loss: 1.462082  
(Iteration 2301 / 4900) loss: 1.398239  
(Iteration 2401 / 4900) loss: 1.274356  
(Epoch 5 / 10) train acc: 0.476000; val\_acc: 0.478000  
(Iteration 2501 / 4900) loss: 1.470804  
(Iteration 2601 / 4900) loss: 1.400513  
(Iteration 2701 / 4900) loss: 1.349935  
(Iteration 2801 / 4900) loss: 1.412017  
(Iteration 2901 / 4900) loss: 1.474202  
(Epoch 6 / 10) train acc: 0.514000; val\_acc: 0.481000  
(Iteration 3001 / 4900) loss: 1.526689  
(Iteration 3101 / 4900) loss: 1.383219  
(Iteration 3201 / 4900) loss: 1.474176  
(Iteration 3301 / 4900) loss: 1.354879  
(Iteration 3401 / 4900) loss: 1.465474  
(Epoch 7 / 10) train acc: 0.495000; val\_acc: 0.498000  
(Iteration 3501 / 4900) loss: 1.416281  
(Iteration 3601 / 4900) loss: 1.414478  
(Iteration 3701 / 4900) loss: 1.478006  
(Iteration 3801 / 4900) loss: 1.444896  
(Iteration 3901 / 4900) loss: 1.459130  
(Epoch 8 / 10) train acc: 0.524000; val\_acc: 0.508000  
(Iteration 4001 / 4900) loss: 1.510932  
(Iteration 4101 / 4900) loss: 1.409774  
(Iteration 4201 / 4900) loss: 1.302575  
(Iteration 4301 / 4900) loss: 1.444532  
(Iteration 4401 / 4900) loss: 1.519370  
(Epoch 9 / 10) train acc: 0.506000; val\_acc: 0.518000  
(Iteration 4501 / 4900) loss: 1.487605  
(Iteration 4601 / 4900) loss: 1.448664  
(Iteration 4701 / 4900) loss: 1.323189  
(Iteration 4801 / 4900) loss: 1.373579  
(Epoch 10 / 10) train acc: 0.523000; val\_acc: 0.510000  
Best model not updated! Current best\_val\_acc = 0.521  
reg: 0.00256242250532733, lr: 0.02, lr\_decay: 0.9  
(Iteration 1 / 4900) loss: 2.302696  
(Epoch 0 / 10) train acc: 0.095000; val\_acc: 0.102000  
(Iteration 101 / 4900) loss: 2.302159  
(Iteration 201 / 4900) loss: 2.301258  
(Iteration 301 / 4900) loss: 2.299632  
(Iteration 401 / 4900) loss: 2.282574  
(Epoch 1 / 10) train acc: 0.242000; val\_acc: 0.251000  
(Iteration 501 / 4900) loss: 2.228827  
(Iteration 601 / 4900) loss: 2.176412  
(Iteration 701 / 4900) loss: 2.070379

(Iteration 801 / 4900) loss: 2.029730  
(Iteration 901 / 4900) loss: 1.973998  
(Epoch 2 / 10) train acc: 0.301000; val\_acc: 0.345000  
(Iteration 1001 / 4900) loss: 1.858579  
(Iteration 1101 / 4900) loss: 1.839002  
(Iteration 1201 / 4900) loss: 1.714864  
(Iteration 1301 / 4900) loss: 1.715657  
(Iteration 1401 / 4900) loss: 1.710699  
(Epoch 3 / 10) train acc: 0.415000; val\_acc: 0.407000  
(Iteration 1501 / 4900) loss: 1.717364  
(Iteration 1601 / 4900) loss: 1.486253  
(Iteration 1701 / 4900) loss: 1.446220  
(Iteration 1801 / 4900) loss: 1.691938  
(Iteration 1901 / 4900) loss: 1.530626  
(Epoch 4 / 10) train acc: 0.477000; val\_acc: 0.446000  
(Iteration 2001 / 4900) loss: 1.541425  
(Iteration 2101 / 4900) loss: 1.488888  
(Iteration 2201 / 4900) loss: 1.494762  
(Iteration 2301 / 4900) loss: 1.623792  
(Iteration 2401 / 4900) loss: 1.415034  
(Epoch 5 / 10) train acc: 0.500000; val\_acc: 0.480000  
(Iteration 2501 / 4900) loss: 1.332568  
(Iteration 2601 / 4900) loss: 1.439741  
(Iteration 2701 / 4900) loss: 1.466098  
(Iteration 2801 / 4900) loss: 1.456798  
(Iteration 2901 / 4900) loss: 1.381156  
(Epoch 6 / 10) train acc: 0.500000; val\_acc: 0.485000  
(Iteration 3001 / 4900) loss: 1.336462  
(Iteration 3101 / 4900) loss: 1.495612  
(Iteration 3201 / 4900) loss: 1.394181  
(Iteration 3301 / 4900) loss: 1.455102  
(Iteration 3401 / 4900) loss: 1.431457  
(Epoch 7 / 10) train acc: 0.500000; val\_acc: 0.495000  
(Iteration 3501 / 4900) loss: 1.500162  
(Iteration 3601 / 4900) loss: 1.501405  
(Iteration 3701 / 4900) loss: 1.433480  
(Iteration 3801 / 4900) loss: 1.377827  
(Iteration 3901 / 4900) loss: 1.404421  
(Epoch 8 / 10) train acc: 0.504000; val\_acc: 0.512000  
(Iteration 4001 / 4900) loss: 1.331002  
(Iteration 4101 / 4900) loss: 1.519259  
(Iteration 4201 / 4900) loss: 1.513337  
(Iteration 4301 / 4900) loss: 1.338200  
(Iteration 4401 / 4900) loss: 1.517789  
(Epoch 9 / 10) train acc: 0.551000; val\_acc: 0.506000  
(Iteration 4501 / 4900) loss: 1.476329  
(Iteration 4601 / 4900) loss: 1.262637  
(Iteration 4701 / 4900) loss: 1.371927

(Iteration 4801 / 4900) loss: 1.257381  
(Epoch 10 / 10) train acc: 0.488000; val\_acc: 0.506000  
Best model not updated! Current best\_val\_acc = 0.521  
reg: 0.0036001964025159706, lr: 0.03, lr\_decay: 0.95  
(Iteration 1 / 4900) loss: 2.302735  
(Epoch 0 / 10) train acc: 0.084000; val\_acc: 0.117000  
(Iteration 101 / 4900) loss: 2.302748  
(Iteration 201 / 4900) loss: 2.298720  
(Iteration 301 / 4900) loss: 2.272819  
(Iteration 401 / 4900) loss: 2.171877  
(Epoch 1 / 10) train acc: 0.282000; val\_acc: 0.283000  
(Iteration 501 / 4900) loss: 1.996141  
(Iteration 601 / 4900) loss: 1.909205  
(Iteration 701 / 4900) loss: 1.656267  
(Iteration 801 / 4900) loss: 1.816395  
(Iteration 901 / 4900) loss: 1.472407  
(Epoch 2 / 10) train acc: 0.410000; val\_acc: 0.412000  
(Iteration 1001 / 4900) loss: 1.565494  
(Iteration 1101 / 4900) loss: 1.586828  
(Iteration 1201 / 4900) loss: 1.501856  
(Iteration 1301 / 4900) loss: 1.454508  
(Iteration 1401 / 4900) loss: 1.559640  
(Epoch 3 / 10) train acc: 0.464000; val\_acc: 0.477000  
(Iteration 1501 / 4900) loss: 1.581014  
(Iteration 1601 / 4900) loss: 1.399074  
(Iteration 1701 / 4900) loss: 1.448534  
(Iteration 1801 / 4900) loss: 1.441901  
(Iteration 1901 / 4900) loss: 1.408067  
(Epoch 4 / 10) train acc: 0.503000; val\_acc: 0.494000  
(Iteration 2001 / 4900) loss: 1.443239  
(Iteration 2101 / 4900) loss: 1.253840  
(Iteration 2201 / 4900) loss: 1.342251  
(Iteration 2301 / 4900) loss: 1.192912  
(Iteration 2401 / 4900) loss: 1.317035  
(Epoch 5 / 10) train acc: 0.514000; val\_acc: 0.510000  
(Iteration 2501 / 4900) loss: 1.491513  
(Iteration 2601 / 4900) loss: 1.248944  
(Iteration 2701 / 4900) loss: 1.386778  
(Iteration 2801 / 4900) loss: 1.374056  
(Iteration 2901 / 4900) loss: 1.414769  
(Epoch 6 / 10) train acc: 0.511000; val\_acc: 0.509000  
(Iteration 3001 / 4900) loss: 1.372700  
(Iteration 3101 / 4900) loss: 1.239072  
(Iteration 3201 / 4900) loss: 1.289897  
(Iteration 3301 / 4900) loss: 1.477414  
(Iteration 3401 / 4900) loss: 1.240550  
(Epoch 7 / 10) train acc: 0.529000; val\_acc: 0.514000  
(Iteration 3501 / 4900) loss: 1.453837

(Iteration 3601 / 4900) loss: 1.295629  
(Iteration 3701 / 4900) loss: 1.379697  
(Iteration 3801 / 4900) loss: 1.298346  
(Iteration 3901 / 4900) loss: 1.284672  
(Epoch 8 / 10) train acc: 0.555000; val\_acc: 0.519000  
(Iteration 4001 / 4900) loss: 1.276887  
(Iteration 4101 / 4900) loss: 1.335793  
(Iteration 4201 / 4900) loss: 1.336391  
(Iteration 4301 / 4900) loss: 1.548589  
(Iteration 4401 / 4900) loss: 1.253067  
(Epoch 9 / 10) train acc: 0.553000; val\_acc: 0.521000  
(Iteration 4501 / 4900) loss: 1.302257  
(Iteration 4601 / 4900) loss: 1.106692  
(Iteration 4701 / 4900) loss: 1.388893  
(Iteration 4801 / 4900) loss: 1.397395  
(Epoch 10 / 10) train acc: 0.511000; val\_acc: 0.529000  
Best model updated! Current best\_val\_acc = 0.529  
reg: 0.003088445987370089, lr: 0.03, lr\_decay: 0.95  
(Iteration 1 / 4900) loss: 2.302720  
(Epoch 0 / 10) train acc: 0.092000; val\_acc: 0.102000  
(Iteration 101 / 4900) loss: 2.301671  
(Iteration 201 / 4900) loss: 2.297774  
(Iteration 301 / 4900) loss: 2.275135  
(Iteration 401 / 4900) loss: 2.181215  
(Epoch 1 / 10) train acc: 0.257000; val\_acc: 0.271000  
(Iteration 501 / 4900) loss: 1.974060  
(Iteration 601 / 4900) loss: 1.976174  
(Iteration 701 / 4900) loss: 1.786811  
(Iteration 801 / 4900) loss: 1.676008  
(Iteration 901 / 4900) loss: 1.650471  
(Epoch 2 / 10) train acc: 0.435000; val\_acc: 0.414000  
(Iteration 1001 / 4900) loss: 1.677524  
(Iteration 1101 / 4900) loss: 1.496707  
(Iteration 1201 / 4900) loss: 1.452835  
(Iteration 1301 / 4900) loss: 1.578841  
(Iteration 1401 / 4900) loss: 1.402819  
(Epoch 3 / 10) train acc: 0.481000; val\_acc: 0.467000  
(Iteration 1501 / 4900) loss: 1.373629  
(Iteration 1601 / 4900) loss: 1.469462  
(Iteration 1701 / 4900) loss: 1.492323  
(Iteration 1801 / 4900) loss: 1.536912  
(Iteration 1901 / 4900) loss: 1.238781  
(Epoch 4 / 10) train acc: 0.492000; val\_acc: 0.496000  
(Iteration 2001 / 4900) loss: 1.353818  
(Iteration 2101 / 4900) loss: 1.505854  
(Iteration 2201 / 4900) loss: 1.189055  
(Iteration 2301 / 4900) loss: 1.513983  
(Iteration 2401 / 4900) loss: 1.636939



```

(Epoch 5 / 10) train acc: 0.528000; val_acc: 0.508000
(Iteration 2501 / 4900) loss: 1.427827
(Iteration 2601 / 4900) loss: 1.557327
(Iteration 2701 / 4900) loss: 1.313562
(Iteration 2801 / 4900) loss: 1.359917
(Iteration 2901 / 4900) loss: 1.464851
(Epoch 6 / 10) train acc: 0.529000; val_acc: 0.527000
(Iteration 3001 / 4900) loss: 1.365478
(Iteration 3101 / 4900) loss: 1.385279
(Iteration 3201 / 4900) loss: 1.312410
(Iteration 3301 / 4900) loss: 1.275741
(Iteration 3401 / 4900) loss: 1.389778
(Epoch 7 / 10) train acc: 0.511000; val_acc: 0.510000
(Iteration 3501 / 4900) loss: 1.642136
(Iteration 3601 / 4900) loss: 1.459408
(Iteration 3701 / 4900) loss: 1.396858
(Iteration 3801 / 4900) loss: 1.400433
(Iteration 3901 / 4900) loss: 1.164601
(Epoch 8 / 10) train acc: 0.516000; val_acc: 0.524000
(Iteration 4001 / 4900) loss: 1.434185
(Iteration 4101 / 4900) loss: 1.475159
(Iteration 4201 / 4900) loss: 1.288996
(Iteration 4301 / 4900) loss: 1.216085
(Iteration 4401 / 4900) loss: 1.344583
(Epoch 9 / 10) train acc: 0.546000; val_acc: 0.535000
(Iteration 4501 / 4900) loss: 1.400607
(Iteration 4601 / 4900) loss: 1.393992
(Iteration 4701 / 4900) loss: 1.472142
(Iteration 4801 / 4900) loss: 1.249673
(Epoch 10 / 10) train acc: 0.536000; val_acc: 0.533000
Best model updated! Current best_val_acc = 0.535
reg: 0.007061937068543055, lr: 0.03, lr_decay: 0.95
(Iteration 1 / 4900) loss: 2.302889
(Epoch 0 / 10) train acc: 0.088000; val_acc: 0.092000
(Iteration 101 / 4900) loss: 2.301880
(Iteration 201 / 4900) loss: 2.299842
(Iteration 301 / 4900) loss: 2.264829
(Iteration 401 / 4900) loss: 2.134453
(Epoch 1 / 10) train acc: 0.260000; val_acc: 0.272000
(Iteration 501 / 4900) loss: 2.009630
(Iteration 601 / 4900) loss: 1.872249
(Iteration 701 / 4900) loss: 1.835755
(Iteration 801 / 4900) loss: 1.705229
(Iteration 901 / 4900) loss: 1.601470
(Epoch 2 / 10) train acc: 0.410000; val_acc: 0.416000
(Iteration 1001 / 4900) loss: 1.497378
(Iteration 1101 / 4900) loss: 1.631437
(Iteration 1201 / 4900) loss: 1.657021

```

```

(Iteration 1301 / 4900) loss: 1.517785
(Iteration 1401 / 4900) loss: 1.534743
(Epoch 3 / 10) train acc: 0.506000; val_acc: 0.471000
(Iteration 1501 / 4900) loss: 1.579522
(Iteration 1601 / 4900) loss: 1.294889
(Iteration 1701 / 4900) loss: 1.569871
(Iteration 1801 / 4900) loss: 1.361690
(Iteration 1901 / 4900) loss: 1.504259
(Epoch 4 / 10) train acc: 0.497000; val_acc: 0.503000
(Iteration 2001 / 4900) loss: 1.409191
(Iteration 2101 / 4900) loss: 1.472215
(Iteration 2201 / 4900) loss: 1.461426
(Iteration 2301 / 4900) loss: 1.422191
(Iteration 2401 / 4900) loss: 1.475245
(Epoch 5 / 10) train acc: 0.518000; val_acc: 0.509000
(Iteration 2501 / 4900) loss: 1.441840
(Iteration 2601 / 4900) loss: 1.491759
(Iteration 2701 / 4900) loss: 1.605277
(Iteration 2801 / 4900) loss: 1.550305
(Iteration 2901 / 4900) loss: 1.378523
(Epoch 6 / 10) train acc: 0.506000; val_acc: 0.511000
(Iteration 3001 / 4900) loss: 1.538893
(Iteration 3101 / 4900) loss: 1.404542
(Iteration 3201 / 4900) loss: 1.623561
(Iteration 3301 / 4900) loss: 1.451256
(Iteration 3401 / 4900) loss: 1.688904
(Epoch 7 / 10) train acc: 0.530000; val_acc: 0.510000
(Iteration 3501 / 4900) loss: 1.402773
(Iteration 3601 / 4900) loss: 1.431356
(Iteration 3701 / 4900) loss: 1.396989
(Iteration 3801 / 4900) loss: 1.476211
(Iteration 3901 / 4900) loss: 1.502548
(Epoch 8 / 10) train acc: 0.526000; val_acc: 0.513000
(Iteration 4001 / 4900) loss: 1.399324
(Iteration 4101 / 4900) loss: 1.336638
(Iteration 4201 / 4900) loss: 1.325075
(Iteration 4301 / 4900) loss: 1.374993
(Iteration 4401 / 4900) loss: 1.352844
(Epoch 9 / 10) train acc: 0.517000; val_acc: 0.521000
(Iteration 4501 / 4900) loss: 1.265600
(Iteration 4601 / 4900) loss: 1.329405
(Iteration 4701 / 4900) loss: 1.541241
(Iteration 4801 / 4900) loss: 1.419663
(Epoch 10 / 10) train acc: 0.539000; val_acc: 0.518000
Best model not updated! Current best_val_acc = 0.535
reg: 0.009945080801889286, lr: 0.03, lr_decay: 0.9
(Iteration 1 / 4900) loss: 2.302980
(Epoch 0 / 10) train acc: 0.109000; val_acc: 0.115000

```

(Iteration 101 / 4900) loss: 2.303423  
(Iteration 201 / 4900) loss: 2.300150  
(Iteration 301 / 4900) loss: 2.282959  
(Iteration 401 / 4900) loss: 2.151936  
(Epoch 1 / 10) train acc: 0.309000; val\_acc: 0.301000  
(Iteration 501 / 4900) loss: 2.005799  
(Iteration 601 / 4900) loss: 1.872646  
(Iteration 701 / 4900) loss: 1.859065  
(Iteration 801 / 4900) loss: 1.806350  
(Iteration 901 / 4900) loss: 1.818051  
(Epoch 2 / 10) train acc: 0.419000; val\_acc: 0.413000  
(Iteration 1001 / 4900) loss: 1.669483  
(Iteration 1101 / 4900) loss: 1.530152  
(Iteration 1201 / 4900) loss: 1.584172  
(Iteration 1301 / 4900) loss: 1.546017  
(Iteration 1401 / 4900) loss: 1.557112  
(Epoch 3 / 10) train acc: 0.476000; val\_acc: 0.457000  
(Iteration 1501 / 4900) loss: 1.515272  
(Iteration 1601 / 4900) loss: 1.589124  
(Iteration 1701 / 4900) loss: 1.658198  
(Iteration 1801 / 4900) loss: 1.546890  
(Iteration 1901 / 4900) loss: 1.647046  
(Epoch 4 / 10) train acc: 0.500000; val\_acc: 0.492000  
(Iteration 2001 / 4900) loss: 1.544470  
(Iteration 2101 / 4900) loss: 1.569837  
(Iteration 2201 / 4900) loss: 1.471969  
(Iteration 2301 / 4900) loss: 1.441463  
(Iteration 2401 / 4900) loss: 1.531839  
(Epoch 5 / 10) train acc: 0.546000; val\_acc: 0.509000  
(Iteration 2501 / 4900) loss: 1.387261  
(Iteration 2601 / 4900) loss: 1.625086  
(Iteration 2701 / 4900) loss: 1.550702  
(Iteration 2801 / 4900) loss: 1.349510  
(Iteration 2901 / 4900) loss: 1.556547  
(Epoch 6 / 10) train acc: 0.525000; val\_acc: 0.513000  
(Iteration 3001 / 4900) loss: 1.334838  
(Iteration 3101 / 4900) loss: 1.575001  
(Iteration 3201 / 4900) loss: 1.543585  
(Iteration 3301 / 4900) loss: 1.625396  
(Iteration 3401 / 4900) loss: 1.457691  
(Epoch 7 / 10) train acc: 0.524000; val\_acc: 0.517000  
(Iteration 3501 / 4900) loss: 1.340714  
(Iteration 3601 / 4900) loss: 1.189967  
(Iteration 3701 / 4900) loss: 1.707100  
(Iteration 3801 / 4900) loss: 1.545314  
(Iteration 3901 / 4900) loss: 1.552841  
(Epoch 8 / 10) train acc: 0.556000; val\_acc: 0.517000  
(Iteration 4001 / 4900) loss: 1.403642

(Iteration 4101 / 4900) loss: 1.154744  
(Iteration 4201 / 4900) loss: 1.415446  
(Iteration 4301 / 4900) loss: 1.434426  
(Iteration 4401 / 4900) loss: 1.492329  
(Epoch 9 / 10) train acc: 0.521000; val\_acc: 0.512000  
(Iteration 4501 / 4900) loss: 1.364023  
(Iteration 4601 / 4900) loss: 1.495556  
(Iteration 4701 / 4900) loss: 1.380622  
(Iteration 4801 / 4900) loss: 1.473795  
(Epoch 10 / 10) train acc: 0.547000; val\_acc: 0.513000  
Best model not updated! Current best\_val\_acc = 0.535  
reg: 0.0017331682795015857, lr: 0.03, lr\_decay: 0.9  
(Iteration 1 / 4900) loss: 2.302685  
(Epoch 0 / 10) train acc: 0.098000; val\_acc: 0.078000  
(Iteration 101 / 4900) loss: 2.302285  
(Iteration 201 / 4900) loss: 2.298251  
(Iteration 301 / 4900) loss: 2.256724  
(Iteration 401 / 4900) loss: 2.120513  
(Epoch 1 / 10) train acc: 0.265000; val\_acc: 0.295000  
(Iteration 501 / 4900) loss: 1.994900  
(Iteration 601 / 4900) loss: 1.925031  
(Iteration 701 / 4900) loss: 1.774769  
(Iteration 801 / 4900) loss: 1.737080  
(Iteration 901 / 4900) loss: 1.651231  
(Epoch 2 / 10) train acc: 0.443000; val\_acc: 0.420000  
(Iteration 1001 / 4900) loss: 1.538141  
(Iteration 1101 / 4900) loss: 1.596848  
(Iteration 1201 / 4900) loss: 1.614572  
(Iteration 1301 / 4900) loss: 1.426510  
(Iteration 1401 / 4900) loss: 1.538031  
(Epoch 3 / 10) train acc: 0.511000; val\_acc: 0.469000  
(Iteration 1501 / 4900) loss: 1.427059  
(Iteration 1601 / 4900) loss: 1.355369  
(Iteration 1701 / 4900) loss: 1.481911  
(Iteration 1801 / 4900) loss: 1.494077  
(Iteration 1901 / 4900) loss: 1.495286  
(Epoch 4 / 10) train acc: 0.479000; val\_acc: 0.501000  
(Iteration 2001 / 4900) loss: 1.408780  
(Iteration 2101 / 4900) loss: 1.627954  
(Iteration 2201 / 4900) loss: 1.427974  
(Iteration 2301 / 4900) loss: 1.283680  
(Iteration 2401 / 4900) loss: 1.373896  
(Epoch 5 / 10) train acc: 0.534000; val\_acc: 0.518000  
(Iteration 2501 / 4900) loss: 1.229550  
(Iteration 2601 / 4900) loss: 1.493610  
(Iteration 2701 / 4900) loss: 1.355030  
(Iteration 2801 / 4900) loss: 1.284374  
(Iteration 2901 / 4900) loss: 1.276799

```

(Epoch 6 / 10) train acc: 0.516000; val_acc: 0.508000
(Iteration 3001 / 4900) loss: 1.563690
(Iteration 3101 / 4900) loss: 1.414483
(Iteration 3201 / 4900) loss: 1.219454
(Iteration 3301 / 4900) loss: 1.493506
(Iteration 3401 / 4900) loss: 1.357521
(Epoch 7 / 10) train acc: 0.525000; val_acc: 0.511000
(Iteration 3501 / 4900) loss: 1.294848
(Iteration 3601 / 4900) loss: 1.392900
(Iteration 3701 / 4900) loss: 1.402768
(Iteration 3801 / 4900) loss: 1.402790
(Iteration 3901 / 4900) loss: 1.372908
(Epoch 8 / 10) train acc: 0.531000; val_acc: 0.512000
(Iteration 4001 / 4900) loss: 1.320943
(Iteration 4101 / 4900) loss: 1.280615
(Iteration 4201 / 4900) loss: 1.314262
(Iteration 4301 / 4900) loss: 1.450851
(Iteration 4401 / 4900) loss: 1.322835
(Epoch 9 / 10) train acc: 0.551000; val_acc: 0.522000
(Iteration 4501 / 4900) loss: 1.167285
(Iteration 4601 / 4900) loss: 1.210650
(Iteration 4701 / 4900) loss: 1.158096
(Iteration 4801 / 4900) loss: 1.408575
(Epoch 10 / 10) train acc: 0.520000; val_acc: 0.519000
Best model not updated! Current best_val_acc = 0.535
reg: 0.008897566589609421, lr: 0.03, lr_decay: 0.9
(Iteration 1 / 4900) loss: 2.302945
(Epoch 0 / 10) train acc: 0.092000; val_acc: 0.117000
(Iteration 101 / 4900) loss: 2.300631
(Iteration 201 / 4900) loss: 2.300294
(Iteration 301 / 4900) loss: 2.274199
(Iteration 401 / 4900) loss: 2.160620
(Epoch 1 / 10) train acc: 0.276000; val_acc: 0.261000
(Iteration 501 / 4900) loss: 1.966335
(Iteration 601 / 4900) loss: 1.908198
(Iteration 701 / 4900) loss: 1.810882
(Iteration 801 / 4900) loss: 1.791182
(Iteration 901 / 4900) loss: 1.610780
(Epoch 2 / 10) train acc: 0.429000; val_acc: 0.414000
(Iteration 1001 / 4900) loss: 1.464936
(Iteration 1101 / 4900) loss: 1.559175
(Iteration 1201 / 4900) loss: 1.725194
(Iteration 1301 / 4900) loss: 1.499605
(Iteration 1401 / 4900) loss: 1.726554
(Epoch 3 / 10) train acc: 0.491000; val_acc: 0.470000
(Iteration 1501 / 4900) loss: 1.598460
(Iteration 1601 / 4900) loss: 1.555505
(Iteration 1701 / 4900) loss: 1.478003

```

```

(Iteration 1801 / 4900) loss: 1.604852
(Iteration 1901 / 4900) loss: 1.374973
(Epoch 4 / 10) train acc: 0.486000; val_acc: 0.497000
(Iteration 2001 / 4900) loss: 1.615172
(Iteration 2101 / 4900) loss: 1.443966
(Iteration 2201 / 4900) loss: 1.425907
(Iteration 2301 / 4900) loss: 1.525849
(Iteration 2401 / 4900) loss: 1.441216
(Epoch 5 / 10) train acc: 0.508000; val_acc: 0.504000
(Iteration 2501 / 4900) loss: 1.612227
(Iteration 2601 / 4900) loss: 1.533262
(Iteration 2701 / 4900) loss: 1.424977
(Iteration 2801 / 4900) loss: 1.463687
(Iteration 2901 / 4900) loss: 1.440932
(Epoch 6 / 10) train acc: 0.534000; val_acc: 0.499000
(Iteration 3001 / 4900) loss: 1.494282
(Iteration 3101 / 4900) loss: 1.466120
(Iteration 3201 / 4900) loss: 1.459016
(Iteration 3301 / 4900) loss: 1.456424
(Iteration 3401 / 4900) loss: 1.511847
(Epoch 7 / 10) train acc: 0.511000; val_acc: 0.513000
(Iteration 3501 / 4900) loss: 1.354479
(Iteration 3601 / 4900) loss: 1.524966
(Iteration 3701 / 4900) loss: 1.610592
(Iteration 3801 / 4900) loss: 1.557277
(Iteration 3901 / 4900) loss: 1.439574
(Epoch 8 / 10) train acc: 0.542000; val_acc: 0.510000
(Iteration 4001 / 4900) loss: 1.316151
(Iteration 4101 / 4900) loss: 1.529116
(Iteration 4201 / 4900) loss: 1.504594
(Iteration 4301 / 4900) loss: 1.526410
(Iteration 4401 / 4900) loss: 1.473089
(Epoch 9 / 10) train acc: 0.509000; val_acc: 0.517000
(Iteration 4501 / 4900) loss: 1.580902
(Iteration 4601 / 4900) loss: 1.294162
(Iteration 4701 / 4900) loss: 1.549044
(Iteration 4801 / 4900) loss: 1.368667
(Epoch 10 / 10) train acc: 0.550000; val_acc: 0.519000
Best model not updated! Current best_val_acc = 0.535
reg: 0.0042670201517997774, lr: 0.04, lr_decay: 0.95
(Iteration 1 / 4900) loss: 2.302747
(Epoch 0 / 10) train acc: 0.121000; val_acc: 0.117000
(Iteration 101 / 4900) loss: 2.301479
(Iteration 201 / 4900) loss: 2.283296
(Iteration 301 / 4900) loss: 2.166840
(Iteration 401 / 4900) loss: 1.903839
(Epoch 1 / 10) train acc: 0.347000; val_acc: 0.365000
(Iteration 501 / 4900) loss: 1.833393

```

(Iteration 601 / 4900) loss: 1.746269  
(Iteration 701 / 4900) loss: 1.605726  
(Iteration 801 / 4900) loss: 1.581127  
(Iteration 901 / 4900) loss: 1.424700  
(Epoch 2 / 10) train acc: 0.467000; val\_acc: 0.473000  
(Iteration 1001 / 4900) loss: 1.443430  
(Iteration 1101 / 4900) loss: 1.379664  
(Iteration 1201 / 4900) loss: 1.410274  
(Iteration 1301 / 4900) loss: 1.531078  
(Iteration 1401 / 4900) loss: 1.442967  
(Epoch 3 / 10) train acc: 0.474000; val\_acc: 0.502000  
(Iteration 1501 / 4900) loss: 1.582030  
(Iteration 1601 / 4900) loss: 1.280430  
(Iteration 1701 / 4900) loss: 1.441667  
(Iteration 1801 / 4900) loss: 1.334207  
(Iteration 1901 / 4900) loss: 1.497039  
(Epoch 4 / 10) train acc: 0.509000; val\_acc: 0.511000  
(Iteration 2001 / 4900) loss: 1.204012  
(Iteration 2101 / 4900) loss: 1.446197  
(Iteration 2201 / 4900) loss: 1.391292  
(Iteration 2301 / 4900) loss: 1.222993  
(Iteration 2401 / 4900) loss: 1.556796  
(Epoch 5 / 10) train acc: 0.542000; val\_acc: 0.523000  
(Iteration 2501 / 4900) loss: 1.245398  
(Iteration 2601 / 4900) loss: 1.455973  
(Iteration 2701 / 4900) loss: 1.249450  
(Iteration 2801 / 4900) loss: 1.412764  
(Iteration 2901 / 4900) loss: 1.469615  
(Epoch 6 / 10) train acc: 0.544000; val\_acc: 0.516000  
(Iteration 3001 / 4900) loss: 1.419421  
(Iteration 3101 / 4900) loss: 1.287655  
(Iteration 3201 / 4900) loss: 1.353718  
(Iteration 3301 / 4900) loss: 1.407999  
(Iteration 3401 / 4900) loss: 1.260587  
(Epoch 7 / 10) train acc: 0.515000; val\_acc: 0.529000  
(Iteration 3501 / 4900) loss: 1.350333  
(Iteration 3601 / 4900) loss: 1.354071  
(Iteration 3701 / 4900) loss: 1.355664  
(Iteration 3801 / 4900) loss: 1.362047  
(Iteration 3901 / 4900) loss: 1.340277  
(Epoch 8 / 10) train acc: 0.530000; val\_acc: 0.528000  
(Iteration 4001 / 4900) loss: 1.367506  
(Iteration 4101 / 4900) loss: 1.366121  
(Iteration 4201 / 4900) loss: 1.437349  
(Iteration 4301 / 4900) loss: 1.298983  
(Iteration 4401 / 4900) loss: 1.246331  
(Epoch 9 / 10) train acc: 0.528000; val\_acc: 0.538000  
(Iteration 4501 / 4900) loss: 1.337094

(Iteration 4601 / 4900) loss: 1.424596  
(Iteration 4701 / 4900) loss: 1.350804  
(Iteration 4801 / 4900) loss: 1.519393  
(Epoch 10 / 10) train acc: 0.537000; val\_acc: 0.539000  
Best model updated! Current best\_val\_acc = 0.539  
reg: 0.004754994024508956, lr: 0.04, lr\_decay: 0.95  
(Iteration 1 / 4900) loss: 2.302759  
(Epoch 0 / 10) train acc: 0.107000; val\_acc: 0.107000  
(Iteration 101 / 4900) loss: 2.302529  
(Iteration 201 / 4900) loss: 2.277933  
(Iteration 301 / 4900) loss: 2.118076  
(Iteration 401 / 4900) loss: 1.979691  
(Epoch 1 / 10) train acc: 0.326000; val\_acc: 0.338000  
(Iteration 501 / 4900) loss: 1.776086  
(Iteration 601 / 4900) loss: 1.731050  
(Iteration 701 / 4900) loss: 1.653489  
(Iteration 801 / 4900) loss: 1.671821  
(Iteration 901 / 4900) loss: 1.602608  
(Epoch 2 / 10) train acc: 0.473000; val\_acc: 0.467000  
(Iteration 1001 / 4900) loss: 1.455358  
(Iteration 1101 / 4900) loss: 1.446708  
(Iteration 1201 / 4900) loss: 1.554551  
(Iteration 1301 / 4900) loss: 1.625541  
(Iteration 1401 / 4900) loss: 1.402308  
(Epoch 3 / 10) train acc: 0.521000; val\_acc: 0.505000  
(Iteration 1501 / 4900) loss: 1.354211  
(Iteration 1601 / 4900) loss: 1.447121  
(Iteration 1701 / 4900) loss: 1.520160  
(Iteration 1801 / 4900) loss: 1.428114  
(Iteration 1901 / 4900) loss: 1.340998  
(Epoch 4 / 10) train acc: 0.539000; val\_acc: 0.515000  
(Iteration 2001 / 4900) loss: 1.391384  
(Iteration 2101 / 4900) loss: 1.325953  
(Iteration 2201 / 4900) loss: 1.325119  
(Iteration 2301 / 4900) loss: 1.420860  
(Iteration 2401 / 4900) loss: 1.404445  
(Epoch 5 / 10) train acc: 0.552000; val\_acc: 0.507000  
(Iteration 2501 / 4900) loss: 1.461697  
(Iteration 2601 / 4900) loss: 1.210367  
(Iteration 2701 / 4900) loss: 1.438408  
(Iteration 2801 / 4900) loss: 1.350638  
(Iteration 2901 / 4900) loss: 1.458127  
(Epoch 6 / 10) train acc: 0.541000; val\_acc: 0.529000  
(Iteration 3001 / 4900) loss: 1.398753  
(Iteration 3101 / 4900) loss: 1.359091  
(Iteration 3201 / 4900) loss: 1.451976  
(Iteration 3301 / 4900) loss: 1.443536  
(Iteration 3401 / 4900) loss: 1.272158



(Epoch 7 / 10) train acc: 0.549000; val\_acc: 0.526000  
(Iteration 3501 / 4900) loss: 1.256424  
(Iteration 3601 / 4900) loss: 1.521642  
(Iteration 3701 / 4900) loss: 1.288232  
(Iteration 3801 / 4900) loss: 1.330084  
(Iteration 3901 / 4900) loss: 1.223817  
(Epoch 8 / 10) train acc: 0.534000; val\_acc: 0.522000  
(Iteration 4001 / 4900) loss: 1.319028  
(Iteration 4101 / 4900) loss: 1.204268  
(Iteration 4201 / 4900) loss: 1.380608  
(Iteration 4301 / 4900) loss: 1.501211  
(Iteration 4401 / 4900) loss: 1.355918  
(Epoch 9 / 10) train acc: 0.529000; val\_acc: 0.527000  
(Iteration 4501 / 4900) loss: 1.212220  
(Iteration 4601 / 4900) loss: 1.312661  
(Iteration 4701 / 4900) loss: 1.277908  
(Iteration 4801 / 4900) loss: 1.288392  
(Epoch 10 / 10) train acc: 0.555000; val\_acc: 0.535000  
Best model not updated! Current best\_val\_acc = 0.539  
reg: 0.005583951659759184, lr: 0.04, lr\_decay: 0.95  
(Iteration 1 / 4900) loss: 2.302802  
(Epoch 0 / 10) train acc: 0.107000; val\_acc: 0.113000  
(Iteration 101 / 4900) loss: 2.302764  
(Iteration 201 / 4900) loss: 2.285118  
(Iteration 301 / 4900) loss: 2.148094  
(Iteration 401 / 4900) loss: 1.964022  
(Epoch 1 / 10) train acc: 0.342000; val\_acc: 0.343000  
(Iteration 501 / 4900) loss: 1.908868  
(Iteration 601 / 4900) loss: 1.721957  
(Iteration 701 / 4900) loss: 1.591381  
(Iteration 801 / 4900) loss: 1.487172  
(Iteration 901 / 4900) loss: 1.685429  
(Epoch 2 / 10) train acc: 0.473000; val\_acc: 0.469000  
(Iteration 1001 / 4900) loss: 1.475226  
(Iteration 1101 / 4900) loss: 1.551123  
(Iteration 1201 / 4900) loss: 1.549173  
(Iteration 1301 / 4900) loss: 1.506160  
(Iteration 1401 / 4900) loss: 1.539767  
(Epoch 3 / 10) train acc: 0.511000; val\_acc: 0.506000  
(Iteration 1501 / 4900) loss: 1.541444  
(Iteration 1601 / 4900) loss: 1.543346  
(Iteration 1701 / 4900) loss: 1.465909  
(Iteration 1801 / 4900) loss: 1.450490  
(Iteration 1901 / 4900) loss: 1.426127  
(Epoch 4 / 10) train acc: 0.531000; val\_acc: 0.518000  
(Iteration 2001 / 4900) loss: 1.532793  
(Iteration 2101 / 4900) loss: 1.435814  
(Iteration 2201 / 4900) loss: 1.278831

```

(Iteration 2301 / 4900) loss: 1.327719
(Iteration 2401 / 4900) loss: 1.341556
(Epoch 5 / 10) train acc: 0.514000; val_acc: 0.528000
(Iteration 2501 / 4900) loss: 1.281495
(Iteration 2601 / 4900) loss: 1.445321
(Iteration 2701 / 4900) loss: 1.426527
(Iteration 2801 / 4900) loss: 1.368680
(Iteration 2901 / 4900) loss: 1.273435
(Epoch 6 / 10) train acc: 0.554000; val_acc: 0.525000
(Iteration 3001 / 4900) loss: 1.522119
(Iteration 3101 / 4900) loss: 1.342912
(Iteration 3201 / 4900) loss: 1.324764
(Iteration 3301 / 4900) loss: 1.508130
(Iteration 3401 / 4900) loss: 1.339029
(Epoch 7 / 10) train acc: 0.544000; val_acc: 0.524000
(Iteration 3501 / 4900) loss: 1.415011
(Iteration 3601 / 4900) loss: 1.281462
(Iteration 3701 / 4900) loss: 1.392237
(Iteration 3801 / 4900) loss: 1.389093
(Iteration 3901 / 4900) loss: 1.498730
(Epoch 8 / 10) train acc: 0.554000; val_acc: 0.529000
(Iteration 4001 / 4900) loss: 1.408427
(Iteration 4101 / 4900) loss: 1.360371
(Iteration 4201 / 4900) loss: 1.327052
(Iteration 4301 / 4900) loss: 1.198700
(Iteration 4401 / 4900) loss: 1.403170
(Epoch 9 / 10) train acc: 0.566000; val_acc: 0.536000
(Iteration 4501 / 4900) loss: 1.490429
(Iteration 4601 / 4900) loss: 1.416360
(Iteration 4701 / 4900) loss: 1.248418
(Iteration 4801 / 4900) loss: 1.417437
(Epoch 10 / 10) train acc: 0.565000; val_acc: 0.543000
Best model updated! Current best_val_acc = 0.543
reg: 0.003737959513844151, lr: 0.04, lr_decay: 0.9
(Iteration 1 / 4900) loss: 2.302724
(Epoch 0 / 10) train acc: 0.112000; val_acc: 0.127000
(Iteration 101 / 4900) loss: 2.301120
(Iteration 201 / 4900) loss: 2.280862
(Iteration 301 / 4900) loss: 2.118139
(Iteration 401 / 4900) loss: 1.927550
(Epoch 1 / 10) train acc: 0.322000; val_acc: 0.330000
(Iteration 501 / 4900) loss: 1.900152
(Iteration 601 / 4900) loss: 1.738981
(Iteration 701 / 4900) loss: 1.735317
(Iteration 801 / 4900) loss: 1.622550
(Iteration 901 / 4900) loss: 1.631611
(Epoch 2 / 10) train acc: 0.490000; val_acc: 0.457000
(Iteration 1001 / 4900) loss: 1.502742

```

(Iteration 1101 / 4900) loss: 1.430946  
(Iteration 1201 / 4900) loss: 1.487548  
(Iteration 1301 / 4900) loss: 1.441445  
(Iteration 1401 / 4900) loss: 1.355419  
(Epoch 3 / 10) train acc: 0.506000; val\_acc: 0.501000  
(Iteration 1501 / 4900) loss: 1.431786  
(Iteration 1601 / 4900) loss: 1.416725  
(Iteration 1701 / 4900) loss: 1.510405  
(Iteration 1801 / 4900) loss: 1.308217  
(Iteration 1901 / 4900) loss: 1.430316  
(Epoch 4 / 10) train acc: 0.521000; val\_acc: 0.508000  
(Iteration 2001 / 4900) loss: 1.274294  
(Iteration 2101 / 4900) loss: 1.372514  
(Iteration 2201 / 4900) loss: 1.334828  
(Iteration 2301 / 4900) loss: 1.411366  
(Iteration 2401 / 4900) loss: 1.329219  
(Epoch 5 / 10) train acc: 0.530000; val\_acc: 0.527000  
(Iteration 2501 / 4900) loss: 1.350155  
(Iteration 2601 / 4900) loss: 1.374582  
(Iteration 2701 / 4900) loss: 1.325410  
(Iteration 2801 / 4900) loss: 1.500449  
(Iteration 2901 / 4900) loss: 1.211171  
(Epoch 6 / 10) train acc: 0.525000; val\_acc: 0.520000  
(Iteration 3001 / 4900) loss: 1.513531  
(Iteration 3101 / 4900) loss: 1.394831  
(Iteration 3201 / 4900) loss: 1.177537  
(Iteration 3301 / 4900) loss: 1.483656  
(Iteration 3401 / 4900) loss: 1.436726  
(Epoch 7 / 10) train acc: 0.515000; val\_acc: 0.537000  
(Iteration 3501 / 4900) loss: 1.371700  
(Iteration 3601 / 4900) loss: 1.316506  
(Iteration 3701 / 4900) loss: 1.422300  
(Iteration 3801 / 4900) loss: 1.398886  
(Iteration 3901 / 4900) loss: 1.312056  
(Epoch 8 / 10) train acc: 0.534000; val\_acc: 0.519000  
(Iteration 4001 / 4900) loss: 1.335733  
(Iteration 4101 / 4900) loss: 1.526128  
(Iteration 4201 / 4900) loss: 1.221636  
(Iteration 4301 / 4900) loss: 1.384552  
(Iteration 4401 / 4900) loss: 1.525612  
(Epoch 9 / 10) train acc: 0.527000; val\_acc: 0.527000  
(Iteration 4501 / 4900) loss: 1.395461  
(Iteration 4601 / 4900) loss: 1.403099  
(Iteration 4701 / 4900) loss: 1.315966  
(Iteration 4801 / 4900) loss: 1.523312  
(Epoch 10 / 10) train acc: 0.523000; val\_acc: 0.536000  
Best model not updated! Current best\_val\_acc = 0.543  
reg: 0.004342772083976483, lr: 0.04, lr\_decay: 0.9

(Iteration 1 / 4900) loss: 2.302734  
(Epoch 0 / 10) train acc: 0.104000; val\_acc: 0.081000  
(Iteration 101 / 4900) loss: 2.302025  
(Iteration 201 / 4900) loss: 2.287346  
(Iteration 301 / 4900) loss: 2.178829  
(Iteration 401 / 4900) loss: 1.992444  
(Epoch 1 / 10) train acc: 0.342000; val\_acc: 0.362000  
(Iteration 501 / 4900) loss: 1.756205  
(Iteration 601 / 4900) loss: 1.617102  
(Iteration 701 / 4900) loss: 1.614696  
(Iteration 801 / 4900) loss: 1.556049  
(Iteration 901 / 4900) loss: 1.414611  
(Epoch 2 / 10) train acc: 0.488000; val\_acc: 0.460000  
(Iteration 1001 / 4900) loss: 1.392721  
(Iteration 1101 / 4900) loss: 1.315381  
(Iteration 1201 / 4900) loss: 1.591214  
(Iteration 1301 / 4900) loss: 1.517157  
(Iteration 1401 / 4900) loss: 1.488181  
(Epoch 3 / 10) train acc: 0.537000; val\_acc: 0.499000  
(Iteration 1501 / 4900) loss: 1.483716  
(Iteration 1601 / 4900) loss: 1.235566  
(Iteration 1701 / 4900) loss: 1.358595  
(Iteration 1801 / 4900) loss: 1.537630  
(Iteration 1901 / 4900) loss: 1.395319  
(Epoch 4 / 10) train acc: 0.516000; val\_acc: 0.506000  
(Iteration 2001 / 4900) loss: 1.435938  
(Iteration 2101 / 4900) loss: 1.337325  
(Iteration 2201 / 4900) loss: 1.517518  
(Iteration 2301 / 4900) loss: 1.502472  
(Iteration 2401 / 4900) loss: 1.488751  
(Epoch 5 / 10) train acc: 0.514000; val\_acc: 0.506000  
(Iteration 2501 / 4900) loss: 1.409311  
(Iteration 2601 / 4900) loss: 1.391232  
(Iteration 2701 / 4900) loss: 1.452858  
(Iteration 2801 / 4900) loss: 1.451839  
(Iteration 2901 / 4900) loss: 1.283403  
(Epoch 6 / 10) train acc: 0.562000; val\_acc: 0.528000  
(Iteration 3001 / 4900) loss: 1.471788  
(Iteration 3101 / 4900) loss: 1.545582  
(Iteration 3201 / 4900) loss: 1.227683  
(Iteration 3301 / 4900) loss: 1.279847  
(Iteration 3401 / 4900) loss: 1.432803  
(Epoch 7 / 10) train acc: 0.528000; val\_acc: 0.522000  
(Iteration 3501 / 4900) loss: 1.278582  
(Iteration 3601 / 4900) loss: 1.420105  
(Iteration 3701 / 4900) loss: 1.497608  
(Iteration 3801 / 4900) loss: 1.359340  
(Iteration 3901 / 4900) loss: 1.393429

```

(Epoch 8 / 10) train acc: 0.517000; val_acc: 0.527000
(Iteration 4001 / 4900) loss: 1.428442
(Iteration 4101 / 4900) loss: 1.418994
(Iteration 4201 / 4900) loss: 1.370892
(Iteration 4301 / 4900) loss: 1.322235
(Iteration 4401 / 4900) loss: 1.097523
(Epoch 9 / 10) train acc: 0.575000; val_acc: 0.532000
(Iteration 4501 / 4900) loss: 1.410809
(Iteration 4601 / 4900) loss: 1.329689
(Iteration 4701 / 4900) loss: 1.379197
(Iteration 4801 / 4900) loss: 1.381653
(Epoch 10 / 10) train acc: 0.550000; val_acc: 0.526000
Best model not updated! Current best_val_acc = 0.543
reg: 0.001135546903360758, lr: 0.04, lr_decay: 0.9
(Iteration 1 / 4900) loss: 2.302630
(Epoch 0 / 10) train acc: 0.093000; val_acc: 0.107000
(Iteration 101 / 4900) loss: 2.300567
(Iteration 201 / 4900) loss: 2.291856
(Iteration 301 / 4900) loss: 2.169500
(Iteration 401 / 4900) loss: 1.827537
(Epoch 1 / 10) train acc: 0.329000; val_acc: 0.341000
(Iteration 501 / 4900) loss: 1.796357
(Iteration 601 / 4900) loss: 1.745209
(Iteration 701 / 4900) loss: 1.578174
(Iteration 801 / 4900) loss: 1.538519
(Iteration 901 / 4900) loss: 1.495037
(Epoch 2 / 10) train acc: 0.507000; val_acc: 0.456000
(Iteration 1001 / 4900) loss: 1.534344
(Iteration 1101 / 4900) loss: 1.422314
(Iteration 1201 / 4900) loss: 1.252532
(Iteration 1301 / 4900) loss: 1.468575
(Iteration 1401 / 4900) loss: 1.389968
(Epoch 3 / 10) train acc: 0.524000; val_acc: 0.491000
(Iteration 1501 / 4900) loss: 1.357789
(Iteration 1601 / 4900) loss: 1.415780
(Iteration 1701 / 4900) loss: 1.280913
(Iteration 1801 / 4900) loss: 1.419053
(Iteration 1901 / 4900) loss: 1.386028
(Epoch 4 / 10) train acc: 0.536000; val_acc: 0.500000
(Iteration 2001 / 4900) loss: 1.564752
(Iteration 2101 / 4900) loss: 1.441150
(Iteration 2201 / 4900) loss: 1.566663
(Iteration 2301 / 4900) loss: 1.432050
(Iteration 2401 / 4900) loss: 1.447753
(Epoch 5 / 10) train acc: 0.527000; val_acc: 0.513000
(Iteration 2501 / 4900) loss: 1.403063
(Iteration 2601 / 4900) loss: 1.358056
(Iteration 2701 / 4900) loss: 1.450400

```

(Iteration 2801 / 4900) loss: 1.309527  
(Iteration 2901 / 4900) loss: 1.251523  
(Epoch 6 / 10) train acc: 0.522000; val\_acc: 0.511000  
(Iteration 3001 / 4900) loss: 1.570505  
(Iteration 3101 / 4900) loss: 1.385391  
(Iteration 3201 / 4900) loss: 1.360045  
(Iteration 3301 / 4900) loss: 1.568540  
(Iteration 3401 / 4900) loss: 1.394607  
(Epoch 7 / 10) train acc: 0.534000; val\_acc: 0.521000  
(Iteration 3501 / 4900) loss: 1.425360  
(Iteration 3601 / 4900) loss: 1.277113  
(Iteration 3701 / 4900) loss: 1.080870  
(Iteration 3801 / 4900) loss: 1.397057  
(Iteration 3901 / 4900) loss: 1.338287  
(Epoch 8 / 10) train acc: 0.571000; val\_acc: 0.531000  
(Iteration 4001 / 4900) loss: 1.232568  
(Iteration 4101 / 4900) loss: 1.308517  
(Iteration 4201 / 4900) loss: 1.207522  
(Iteration 4301 / 4900) loss: 1.279025  
(Iteration 4401 / 4900) loss: 1.191086  
(Epoch 9 / 10) train acc: 0.574000; val\_acc: 0.528000  
(Iteration 4501 / 4900) loss: 1.248692  
(Iteration 4601 / 4900) loss: 1.279684  
(Iteration 4701 / 4900) loss: 1.109553  
(Iteration 4801 / 4900) loss: 1.251159  
(Epoch 10 / 10) train acc: 0.552000; val\_acc: 0.536000  
Best model not updated! Current best\_val\_acc = 0.543  
reg: 0.0014622493216023536, lr: 0.05, lr\_decay: 0.95  
(Iteration 1 / 4900) loss: 2.302640  
(Epoch 0 / 10) train acc: 0.109000; val\_acc: 0.105000  
(Iteration 101 / 4900) loss: 2.299347  
(Iteration 201 / 4900) loss: 2.252869  
(Iteration 301 / 4900) loss: 2.050606  
(Iteration 401 / 4900) loss: 1.741156  
(Epoch 1 / 10) train acc: 0.393000; val\_acc: 0.391000  
(Iteration 501 / 4900) loss: 1.637575  
(Iteration 601 / 4900) loss: 1.561043  
(Iteration 701 / 4900) loss: 1.540753  
(Iteration 801 / 4900) loss: 1.587663  
(Iteration 901 / 4900) loss: 1.608477  
(Epoch 2 / 10) train acc: 0.478000; val\_acc: 0.490000  
(Iteration 1001 / 4900) loss: 1.500582  
(Iteration 1101 / 4900) loss: 1.483937  
(Iteration 1201 / 4900) loss: 1.438319  
(Iteration 1301 / 4900) loss: 1.391242  
(Iteration 1401 / 4900) loss: 1.345385  
(Epoch 3 / 10) train acc: 0.514000; val\_acc: 0.515000  
(Iteration 1501 / 4900) loss: 1.457959

(Iteration 1601 / 4900) loss: 1.441968  
(Iteration 1701 / 4900) loss: 1.311658  
(Iteration 1801 / 4900) loss: 1.333653  
(Iteration 1901 / 4900) loss: 1.335108  
(Epoch 4 / 10) train acc: 0.551000; val\_acc: 0.523000  
(Iteration 2001 / 4900) loss: 1.432192  
(Iteration 2101 / 4900) loss: 1.194434  
(Iteration 2201 / 4900) loss: 1.512674  
(Iteration 2301 / 4900) loss: 1.316328  
(Iteration 2401 / 4900) loss: 1.384581  
(Epoch 5 / 10) train acc: 0.533000; val\_acc: 0.521000  
(Iteration 2501 / 4900) loss: 1.305808  
(Iteration 2601 / 4900) loss: 1.399481  
(Iteration 2701 / 4900) loss: 1.270162  
(Iteration 2801 / 4900) loss: 1.276870  
(Iteration 2901 / 4900) loss: 1.399230  
(Epoch 6 / 10) train acc: 0.528000; val\_acc: 0.540000  
(Iteration 3001 / 4900) loss: 1.333853  
(Iteration 3101 / 4900) loss: 1.340327  
(Iteration 3201 / 4900) loss: 1.316031  
(Iteration 3301 / 4900) loss: 1.288520  
(Iteration 3401 / 4900) loss: 1.337040  
(Epoch 7 / 10) train acc: 0.546000; val\_acc: 0.544000  
(Iteration 3501 / 4900) loss: 1.211602  
(Iteration 3601 / 4900) loss: 1.294909  
(Iteration 3701 / 4900) loss: 1.080995  
(Iteration 3801 / 4900) loss: 1.238356  
(Iteration 3901 / 4900) loss: 1.162288  
(Epoch 8 / 10) train acc: 0.566000; val\_acc: 0.538000  
(Iteration 4001 / 4900) loss: 1.121382  
(Iteration 4101 / 4900) loss: 1.396107  
(Iteration 4201 / 4900) loss: 1.190199  
(Iteration 4301 / 4900) loss: 1.135050  
(Iteration 4401 / 4900) loss: 1.114441  
(Epoch 9 / 10) train acc: 0.591000; val\_acc: 0.537000  
(Iteration 4501 / 4900) loss: 1.276771  
(Iteration 4601 / 4900) loss: 1.153642  
(Iteration 4701 / 4900) loss: 1.285658  
(Iteration 4801 / 4900) loss: 1.089603  
(Epoch 10 / 10) train acc: 0.557000; val\_acc: 0.550000  
Best model updated! Current best\_val\_acc = 0.55  
reg: 0.0058840074606068375, lr: 0.05, lr\_decay: 0.95  
(Iteration 1 / 4900) loss: 2.302829  
(Epoch 0 / 10) train acc: 0.109000; val\_acc: 0.110000  
(Iteration 101 / 4900) loss: 2.300753  
(Iteration 201 / 4900) loss: 2.247047  
(Iteration 301 / 4900) loss: 1.978799  
(Iteration 401 / 4900) loss: 1.891076

(Epoch 1 / 10) train acc: 0.403000; val\_acc: 0.382000  
(Iteration 501 / 4900) loss: 1.656623  
(Iteration 601 / 4900) loss: 1.744760  
(Iteration 701 / 4900) loss: 1.477666  
(Iteration 801 / 4900) loss: 1.545750  
(Iteration 901 / 4900) loss: 1.522882  
(Epoch 2 / 10) train acc: 0.489000; val\_acc: 0.476000  
(Iteration 1001 / 4900) loss: 1.487454  
(Iteration 1101 / 4900) loss: 1.760146  
(Iteration 1201 / 4900) loss: 1.483809  
(Iteration 1301 / 4900) loss: 1.283048  
(Iteration 1401 / 4900) loss: 1.399890  
(Epoch 3 / 10) train acc: 0.531000; val\_acc: 0.511000  
(Iteration 1501 / 4900) loss: 1.468313  
(Iteration 1601 / 4900) loss: 1.521420  
(Iteration 1701 / 4900) loss: 1.246556  
(Iteration 1801 / 4900) loss: 1.355741  
(Iteration 1901 / 4900) loss: 1.527269  
(Epoch 4 / 10) train acc: 0.528000; val\_acc: 0.514000  
(Iteration 2001 / 4900) loss: 1.269606  
(Iteration 2101 / 4900) loss: 1.458064  
(Iteration 2201 / 4900) loss: 1.426698  
(Iteration 2301 / 4900) loss: 1.512628  
(Iteration 2401 / 4900) loss: 1.346540  
(Epoch 5 / 10) train acc: 0.542000; val\_acc: 0.522000  
(Iteration 2501 / 4900) loss: 1.315628  
(Iteration 2601 / 4900) loss: 1.403844  
(Iteration 2701 / 4900) loss: 1.420026  
(Iteration 2801 / 4900) loss: 1.336880  
(Iteration 2901 / 4900) loss: 1.546546  
(Epoch 6 / 10) train acc: 0.542000; val\_acc: 0.525000  
(Iteration 3001 / 4900) loss: 1.327140  
(Iteration 3101 / 4900) loss: 1.215464  
(Iteration 3201 / 4900) loss: 1.307765  
(Iteration 3301 / 4900) loss: 1.455142  
(Iteration 3401 / 4900) loss: 1.176396  
(Epoch 7 / 10) train acc: 0.553000; val\_acc: 0.527000  
(Iteration 3501 / 4900) loss: 1.202558  
(Iteration 3601 / 4900) loss: 1.411264  
(Iteration 3701 / 4900) loss: 1.283535  
(Iteration 3801 / 4900) loss: 1.355352  
(Iteration 3901 / 4900) loss: 1.400762  
(Epoch 8 / 10) train acc: 0.566000; val\_acc: 0.529000  
(Iteration 4001 / 4900) loss: 1.166888  
(Iteration 4101 / 4900) loss: 1.296801  
(Iteration 4201 / 4900) loss: 1.260314  
(Iteration 4301 / 4900) loss: 1.410738  
(Iteration 4401 / 4900) loss: 1.325958



(Epoch 9 / 10) train acc: 0.548000; val\_acc: 0.535000  
(Iteration 4501 / 4900) loss: 1.359226  
(Iteration 4601 / 4900) loss: 1.487783  
(Iteration 4701 / 4900) loss: 1.280453  
(Iteration 4801 / 4900) loss: 1.454405  
(Epoch 10 / 10) train acc: 0.563000; val\_acc: 0.534000  
Best model not updated! Current best\_val\_acc = 0.55  
reg: 0.006201180328356196, lr: 0.05, lr\_decay: 0.95  
(Iteration 1 / 4900) loss: 2.302837  
(Epoch 0 / 10) train acc: 0.109000; val\_acc: 0.113000  
(Iteration 101 / 4900) loss: 2.300061  
(Iteration 201 / 4900) loss: 2.248859  
(Iteration 301 / 4900) loss: 1.939016  
(Iteration 401 / 4900) loss: 1.742056  
(Epoch 1 / 10) train acc: 0.431000; val\_acc: 0.378000  
(Iteration 501 / 4900) loss: 1.762128  
(Iteration 601 / 4900) loss: 1.618898  
(Iteration 701 / 4900) loss: 1.640279  
(Iteration 801 / 4900) loss: 1.533401  
(Iteration 901 / 4900) loss: 1.589469  
(Epoch 2 / 10) train acc: 0.499000; val\_acc: 0.496000  
(Iteration 1001 / 4900) loss: 1.399639  
(Iteration 1101 / 4900) loss: 1.565930  
(Iteration 1201 / 4900) loss: 1.447606  
(Iteration 1301 / 4900) loss: 1.363689  
(Iteration 1401 / 4900) loss: 1.409356  
(Epoch 3 / 10) train acc: 0.513000; val\_acc: 0.519000  
(Iteration 1501 / 4900) loss: 1.477206  
(Iteration 1601 / 4900) loss: 1.697472  
(Iteration 1701 / 4900) loss: 1.395383  
(Iteration 1801 / 4900) loss: 1.448575  
(Iteration 1901 / 4900) loss: 1.410484  
(Epoch 4 / 10) train acc: 0.554000; val\_acc: 0.518000  
(Iteration 2001 / 4900) loss: 1.314541  
(Iteration 2101 / 4900) loss: 1.406195  
(Iteration 2201 / 4900) loss: 1.334702  
(Iteration 2301 / 4900) loss: 1.440236  
(Iteration 2401 / 4900) loss: 1.373688  
(Epoch 5 / 10) train acc: 0.529000; val\_acc: 0.524000  
(Iteration 2501 / 4900) loss: 1.395942  
(Iteration 2601 / 4900) loss: 1.359225  
(Iteration 2701 / 4900) loss: 1.292595  
(Iteration 2801 / 4900) loss: 1.463891  
(Iteration 2901 / 4900) loss: 1.361303  
(Epoch 6 / 10) train acc: 0.544000; val\_acc: 0.521000  
(Iteration 3001 / 4900) loss: 1.569166  
(Iteration 3101 / 4900) loss: 1.309004  
(Iteration 3201 / 4900) loss: 1.541461

(Iteration 3301 / 4900) loss: 1.403744  
(Iteration 3401 / 4900) loss: 1.258613  
(Epoch 7 / 10) train acc: 0.550000; val\_acc: 0.531000  
(Iteration 3501 / 4900) loss: 1.394688  
(Iteration 3601 / 4900) loss: 1.269674  
(Iteration 3701 / 4900) loss: 1.537852  
(Iteration 3801 / 4900) loss: 1.353359  
(Iteration 3901 / 4900) loss: 1.292975  
(Epoch 8 / 10) train acc: 0.570000; val\_acc: 0.531000  
(Iteration 4001 / 4900) loss: 1.533944  
(Iteration 4101 / 4900) loss: 1.458882  
(Iteration 4201 / 4900) loss: 1.360439  
(Iteration 4301 / 4900) loss: 1.376974  
(Iteration 4401 / 4900) loss: 1.349002  
(Epoch 9 / 10) train acc: 0.558000; val\_acc: 0.535000  
(Iteration 4501 / 4900) loss: 1.462045  
(Iteration 4601 / 4900) loss: 1.416158  
(Iteration 4701 / 4900) loss: 1.355129  
(Iteration 4801 / 4900) loss: 1.461851  
(Epoch 10 / 10) train acc: 0.555000; val\_acc: 0.542000  
Best model not updated! Current best\_val\_acc = 0.55  
reg: 9.004636087590723e-05, lr: 0.05, lr\_decay: 0.9  
(Iteration 1 / 4900) loss: 2.302589  
(Epoch 0 / 10) train acc: 0.122000; val\_acc: 0.113000  
(Iteration 101 / 4900) loss: 2.297850  
(Iteration 201 / 4900) loss: 2.251503  
(Iteration 301 / 4900) loss: 1.992687  
(Iteration 401 / 4900) loss: 1.730822  
(Epoch 1 / 10) train acc: 0.425000; val\_acc: 0.404000  
(Iteration 501 / 4900) loss: 1.681764  
(Iteration 601 / 4900) loss: 1.537411  
(Iteration 701 / 4900) loss: 1.743662  
(Iteration 801 / 4900) loss: 1.374220  
(Iteration 901 / 4900) loss: 1.464760  
(Epoch 2 / 10) train acc: 0.472000; val\_acc: 0.484000  
(Iteration 1001 / 4900) loss: 1.395565  
(Iteration 1101 / 4900) loss: 1.372940  
(Iteration 1201 / 4900) loss: 1.462465  
(Iteration 1301 / 4900) loss: 1.284131  
(Iteration 1401 / 4900) loss: 1.346532  
(Epoch 3 / 10) train acc: 0.525000; val\_acc: 0.504000  
(Iteration 1501 / 4900) loss: 1.376112  
(Iteration 1601 / 4900) loss: 1.342611  
(Iteration 1701 / 4900) loss: 1.319477  
(Iteration 1801 / 4900) loss: 1.427429  
(Iteration 1901 / 4900) loss: 1.473876  
(Epoch 4 / 10) train acc: 0.515000; val\_acc: 0.519000  
(Iteration 2001 / 4900) loss: 1.382248

(Iteration 2101 / 4900) loss: 1.105966  
(Iteration 2201 / 4900) loss: 1.364949  
(Iteration 2301 / 4900) loss: 1.302300  
(Iteration 2401 / 4900) loss: 1.156425  
(Epoch 5 / 10) train acc: 0.547000; val\_acc: 0.520000  
(Iteration 2501 / 4900) loss: 1.377891  
(Iteration 2601 / 4900) loss: 1.293091  
(Iteration 2701 / 4900) loss: 1.302942  
(Iteration 2801 / 4900) loss: 1.425380  
(Iteration 2901 / 4900) loss: 1.232560  
(Epoch 6 / 10) train acc: 0.555000; val\_acc: 0.530000  
(Iteration 3001 / 4900) loss: 1.259354  
(Iteration 3101 / 4900) loss: 1.296527  
(Iteration 3201 / 4900) loss: 1.285576  
(Iteration 3301 / 4900) loss: 1.181589  
(Iteration 3401 / 4900) loss: 1.252187  
(Epoch 7 / 10) train acc: 0.539000; val\_acc: 0.543000  
(Iteration 3501 / 4900) loss: 1.338936  
(Iteration 3601 / 4900) loss: 1.121984  
(Iteration 3701 / 4900) loss: 1.188004  
(Iteration 3801 / 4900) loss: 1.492186  
(Iteration 3901 / 4900) loss: 1.374556  
(Epoch 8 / 10) train acc: 0.546000; val\_acc: 0.536000  
(Iteration 4001 / 4900) loss: 1.102000  
(Iteration 4101 / 4900) loss: 1.159629  
(Iteration 4201 / 4900) loss: 1.311102  
(Iteration 4301 / 4900) loss: 1.451952  
(Iteration 4401 / 4900) loss: 1.113446  
(Epoch 9 / 10) train acc: 0.565000; val\_acc: 0.542000  
(Iteration 4501 / 4900) loss: 1.085811  
(Iteration 4601 / 4900) loss: 1.364109  
(Iteration 4701 / 4900) loss: 1.316195  
(Iteration 4801 / 4900) loss: 1.156950  
(Epoch 10 / 10) train acc: 0.565000; val\_acc: 0.544000  
Best model not updated! Current best\_val\_acc = 0.55  
reg: 0.0023892758381157155, lr: 0.05, lr\_decay: 0.9  
(Iteration 1 / 4900) loss: 2.302633  
(Epoch 0 / 10) train acc: 0.124000; val\_acc: 0.092000  
(Iteration 101 / 4900) loss: 2.301574  
(Iteration 201 / 4900) loss: 2.248760  
(Iteration 301 / 4900) loss: 2.002132  
(Iteration 401 / 4900) loss: 1.828884  
(Epoch 1 / 10) train acc: 0.384000; val\_acc: 0.388000  
(Iteration 501 / 4900) loss: 1.881721  
(Iteration 601 / 4900) loss: 1.484431  
(Iteration 701 / 4900) loss: 1.560315  
(Iteration 801 / 4900) loss: 1.521358  
(Iteration 901 / 4900) loss: 1.334481


(Epoch 2 / 10) train acc: 0.499000; val\_acc: 0.491000  
(Iteration 1001 / 4900) loss: 1.361315  
(Iteration 1101 / 4900) loss: 1.390466  
(Iteration 1201 / 4900) loss: 1.454913  
(Iteration 1301 / 4900) loss: 1.465441  
(Iteration 1401 / 4900) loss: 1.626172  
(Epoch 3 / 10) train acc: 0.554000; val\_acc: 0.513000  
(Iteration 1501 / 4900) loss: 1.425924  
(Iteration 1601 / 4900) loss: 1.268329  
(Iteration 1701 / 4900) loss: 1.465043  
(Iteration 1801 / 4900) loss: 1.314765  
(Iteration 1901 / 4900) loss: 1.287956  
(Epoch 4 / 10) train acc: 0.528000; val\_acc: 0.501000  
(Iteration 2001 / 4900) loss: 1.357973  
(Iteration 2101 / 4900) loss: 1.324467  
(Iteration 2201 / 4900) loss: 1.261577  
(Iteration 2301 / 4900) loss: 1.519400  
(Iteration 2401 / 4900) loss: 1.291968  
(Epoch 5 / 10) train acc: 0.551000; val\_acc: 0.524000  
(Iteration 2501 / 4900) loss: 1.240498  
(Iteration 2601 / 4900) loss: 1.343357  
(Iteration 2701 / 4900) loss: 1.476549  
(Iteration 2801 / 4900) loss: 1.212636  
(Iteration 2901 / 4900) loss: 1.405382  
(Epoch 6 / 10) train acc: 0.559000; val\_acc: 0.531000  
(Iteration 3001 / 4900) loss: 1.493392  
(Iteration 3101 / 4900) loss: 1.436087  
(Iteration 3201 / 4900) loss: 1.294527  
(Iteration 3301 / 4900) loss: 1.317041  
(Iteration 3401 / 4900) loss: 1.111045  
(Epoch 7 / 10) train acc: 0.580000; val\_acc: 0.541000  
(Iteration 3501 / 4900) loss: 1.198578  
(Iteration 3601 / 4900) loss: 1.263932  
(Iteration 3701 / 4900) loss: 1.223090  
(Iteration 3801 / 4900) loss: 1.292356  
(Iteration 3901 / 4900) loss: 1.448635  
(Epoch 8 / 10) train acc: 0.564000; val\_acc: 0.538000  
(Iteration 4001 / 4900) loss: 1.330117  
(Iteration 4101 / 4900) loss: 1.365933  
(Iteration 4201 / 4900) loss: 1.239045  
(Iteration 4301 / 4900) loss: 1.549408  
(Iteration 4401 / 4900) loss: 1.227450  
(Epoch 9 / 10) train acc: 0.558000; val\_acc: 0.542000  
(Iteration 4501 / 4900) loss: 1.189017  
(Iteration 4601 / 4900) loss: 1.397579  
(Iteration 4701 / 4900) loss: 1.313204  
(Iteration 4801 / 4900) loss: 1.265475  
(Epoch 10 / 10) train acc: 0.576000; val\_acc: 0.540000

Best model not updated! Current best\_val\_acc = 0.55  
reg: 0.00975405419068744, lr: 0.05, lr\_decay: 0.9  
(Iteration 1 / 4900) loss: 2.302983  
(Epoch 0 / 10) train acc: 0.092000; val\_acc: 0.106000  
(Iteration 101 / 4900) loss: 2.302462  
(Iteration 201 / 4900) loss: 2.248751  
(Iteration 301 / 4900) loss: 1.991316  
(Iteration 401 / 4900) loss: 1.816203  
(Epoch 1 / 10) train acc: 0.360000; val\_acc: 0.379000  
(Iteration 501 / 4900) loss: 1.818780  
(Iteration 601 / 4900) loss: 1.678632  
(Iteration 701 / 4900) loss: 1.536077  
(Iteration 801 / 4900) loss: 1.553191  
(Iteration 901 / 4900) loss: 1.633481  
(Epoch 2 / 10) train acc: 0.481000; val\_acc: 0.480000  
(Iteration 1001 / 4900) loss: 1.675639  
(Iteration 1101 / 4900) loss: 1.706263  
(Iteration 1201 / 4900) loss: 1.383079  
(Iteration 1301 / 4900) loss: 1.650871  
(Iteration 1401 / 4900) loss: 1.484168  
(Epoch 3 / 10) train acc: 0.510000; val\_acc: 0.512000  
(Iteration 1501 / 4900) loss: 1.378727  
(Iteration 1601 / 4900) loss: 1.337651  
(Iteration 1701 / 4900) loss: 1.270965  
(Iteration 1801 / 4900) loss: 1.310999  
(Iteration 1901 / 4900) loss: 1.582698  
(Epoch 4 / 10) train acc: 0.508000; val\_acc: 0.518000  
(Iteration 2001 / 4900) loss: 1.532426  
(Iteration 2101 / 4900) loss: 1.436942  
(Iteration 2201 / 4900) loss: 1.382967  
(Iteration 2301 / 4900) loss: 1.622004  
(Iteration 2401 / 4900) loss: 1.500725  
(Epoch 5 / 10) train acc: 0.519000; val\_acc: 0.526000  
(Iteration 2501 / 4900) loss: 1.305647  
(Iteration 2601 / 4900) loss: 1.344022  
(Iteration 2701 / 4900) loss: 1.373037  
(Iteration 2801 / 4900) loss: 1.549987  
(Iteration 2901 / 4900) loss: 1.471379  
(Epoch 6 / 10) train acc: 0.525000; val\_acc: 0.526000  
(Iteration 3001 / 4900) loss: 1.474687  
(Iteration 3101 / 4900) loss: 1.490745  
(Iteration 3201 / 4900) loss: 1.384546  
(Iteration 3301 / 4900) loss: 1.481700  
(Iteration 3401 / 4900) loss: 1.316688  
(Epoch 7 / 10) train acc: 0.538000; val\_acc: 0.522000  
(Iteration 3501 / 4900) loss: 1.674150  
(Iteration 3601 / 4900) loss: 1.435071  
(Iteration 3701 / 4900) loss: 1.384248

```
(Iteration 3801 / 4900) loss: 1.423214
(Iteration 3901 / 4900) loss: 1.320293
(Epoch 8 / 10) train acc: 0.565000; val_acc: 0.525000
(Iteration 4001 / 4900) loss: 1.516869
(Iteration 4101 / 4900) loss: 1.521298
(Iteration 4201 / 4900) loss: 1.548898
(Iteration 4301 / 4900) loss: 1.472299
(Iteration 4401 / 4900) loss: 1.485015
(Epoch 9 / 10) train acc: 0.539000; val_acc: 0.520000
(Iteration 4501 / 4900) loss: 1.616251
(Iteration 4601 / 4900) loss: 1.471228
(Iteration 4701 / 4900) loss: 1.377117
(Iteration 4801 / 4900) loss: 1.592709
(Epoch 10 / 10) train acc: 0.540000; val_acc: 0.536000
Best model not updated! Current best_val_acc = 0.55
```

```
[18]: plt.subplot(2, 1, 1)
plt.title('Training loss')
plt.plot(best_solver.loss_history, 'o')
plt.xlabel('Iteration')

plt.subplot(2, 1, 2)
plt.title('Accuracy')
plt.plot(best_solver.train_acc_history, '-o', label='train')
plt.plot(best_solver.val_acc_history, '-o', label='val')
plt.plot([0.5] * len(best_solver.val_acc_history), 'k--')
plt.xlabel('Epoch')
plt.legend(loc='lower right')
plt.gcf().set_size_inches(15, 12)
plt.show()
```



[19]: *# Run your best neural net classifier on the test set. You should be able  
# to get more than 55% accuracy.*

```
y_test_pred = np.argmax(best_net.loss(data['X_test']), axis=1)
test_acc = (y_test_pred == data['y_test']).mean()
print(test_acc)
```

0.553

```
[20]: examples_per_class = 8
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse',
↳ 'ship', 'truck']
for cls, cls_name in enumerate(classes):
    idxs = np.where((y_test != cls) & (y_test_pred == cls))[0]
    idxs = np.random.choice(idxs, examples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt.subplot(examples_per_class, len(classes), i * len(classes) + cls +
↳ 1)
        plt.imshow(X_test[idx].astype('uint8'))
```

```
plt.axis('off')
if i == 0:
    plt.title(cls_name)
plt.show()
```

